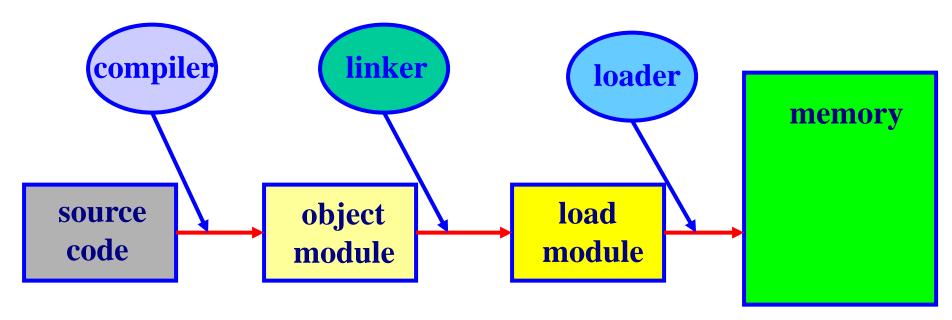
Part II Processes and Threads Process Basics

Program testing can be used to show the presence of bugs, but never to show their absence

From Compilation to Execution

- A compiler compiles source files to . o files.
- A linker links .o files and other libraries together, producing a binary executable (e.g., a.out).
- A loader loads a binary executable into memory for execution.



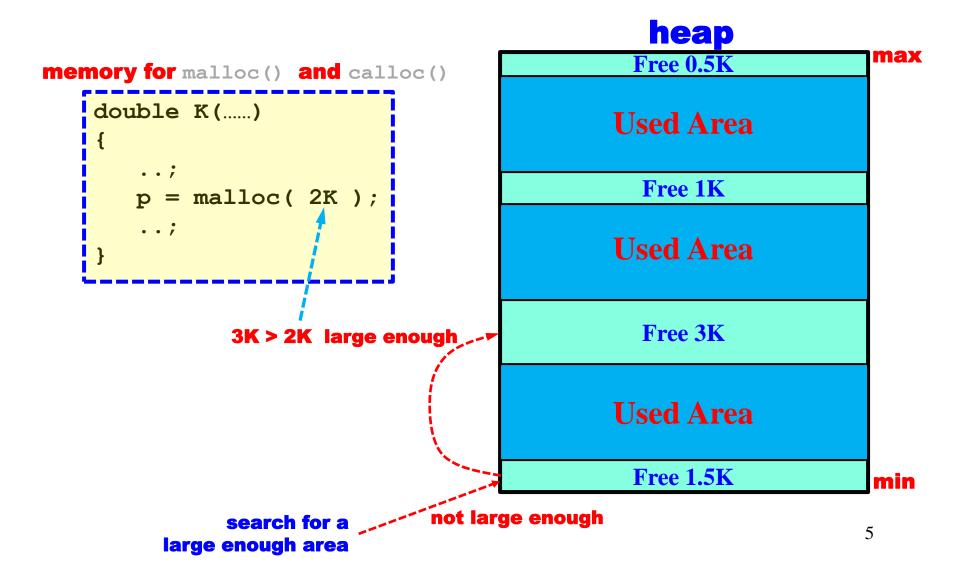
What Is a Process?

- When the OS runs a program (i.e., a binary executable), this program is loaded into memory and the control is transferred to this program's first instruction. Then, the program runs.
- A process is a program in execution.
- A process is more than a program, because a process has a program counter, stack, data section, code section, etc. (i.e., the runtime stuffs)
- Moreover, multiple processes may be associated with one program (e.g., running the same program, say a.out, multiple times at the same time).

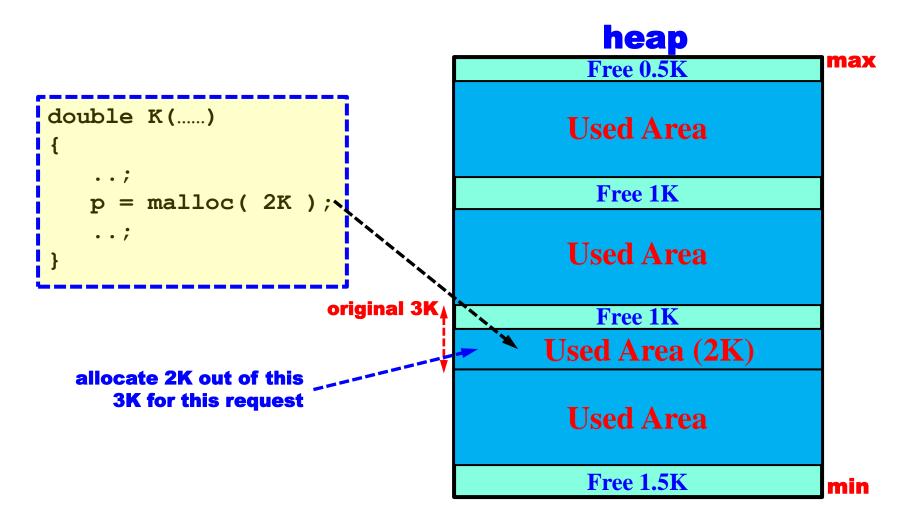
Memory Use of a Program: 1/10

```
int
      h, k;
                                      Global Variables
 double x;
                                         They are declared
 long A(int y, int w)
                                         outside of any function
                                     --Local Variables
    int a, b;
    float t;
                                         These variables are
                                         declared in a function
    a = 12345;
                                         and can only be used
    t = (w+a)/y;
                                         in that function.
    b = sqrt(t)/log(t);
    return (long) b;
                                         Outside of that function
fixed during execution
                                         these variables cannot
                                         be seen and cannot be
int e, f;
                                         used
 int B(.....)
                                       Code
                                         Executable instructions
    return 100;
             After compilation, the compiler knows the total size of global variables
             After linking, the linker knows the total size of the code
```

Memory Use of a Program: 2/10



Memory Use of a Program: 3/10



The memory area for malloc() and calloc() is called the heap

Memory Use of a Program: 4/10

```
int D(.. ..)
   int u, v;
double C(...)
   double w:
   long t;
long B(...)
   char b[20];
   C(...);
void A(.. ..)
   short m, n;
   B(...);
```

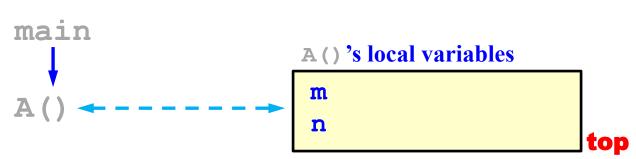
NOTE: When a function is called, all variables declared in this function must be available for this function to execute properly. After this function returns these local variables are no more needed.

- When main () calls A (), m and n must be available
- When A() calls B(), b[20] must be available
- When B() calls C(), w and t must be available
- When C () returns, w and t are not needed
- When B() returns, b[20] is not needed
- When A () returns, m and n are not needed

Call-Return order follows a first-in-last-out (FILO) order The allocation order of local variable is also FILO.

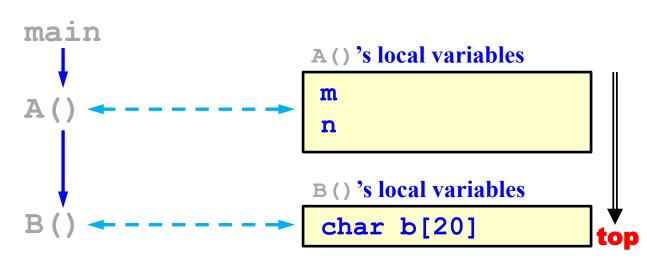
Memory Use of a Program: 5/10

```
int D(...)
  int u, v;
double C(...)
  double w;
  long t;
long B(...)
  char b[20];
  C(...);
void A(.. ..)
  short m, n;
  B(...);
```



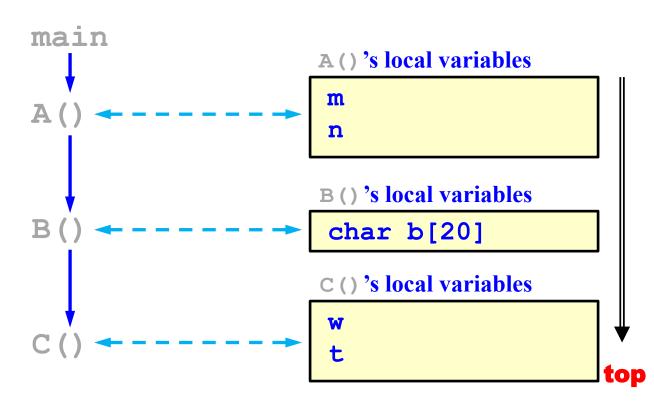
Memory Use of a Program: 6/10

```
int D(...)
   int u,v;
double C(...)
  double w;
  long t;
long B(.. ..)
  char b[20];
  C(...);
void A(.. ..)
  short m, n;
  B(...);
```



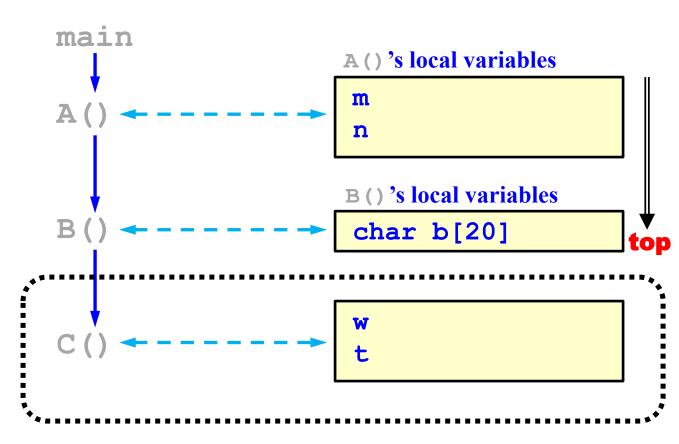
Memory Use of a Program: 7/10

```
int D(...)
   int u,v;
double C(.. ..)
   double w;
   long t;
long B(.. ..)
   char b[20];
   C(...);
void A(.. ..)
   short m, n;
   B(...);
```



Memory Use of a Program: 8/10

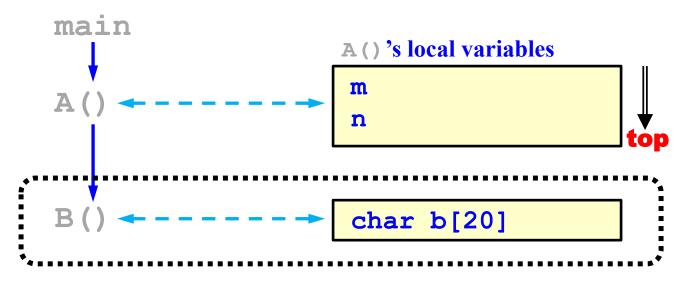
```
int D(...)
  int u,v;
double C(...)
  double w;
  long t;
long B(.. ..)
  char b[20];
  C(...);
void A(.. ..)
  short m, n;
  B(...);
```



C() returns, its local variables are not needed.

Memory Use of a Program: 9/10

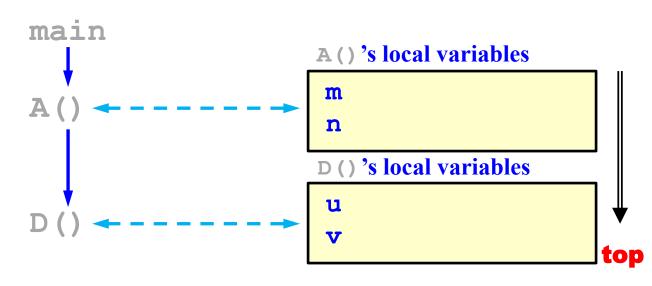
```
int D(...)
  int u,v;
double C(...)
  double w;
  long t;
long B(...)
  char b[20];
  C(...);
void A(.. ..)
  short m, n;
  B(...);
```



B() returns, its local variables are not needed.

Memory Use of a Program: 10/10

```
int D(...)
  int u,v;
double C(...)
  double w;
  long t;
long B(...)
  char b[20];
  C(...);
void A(.. ..)
  short m, n;
  B(...);
```

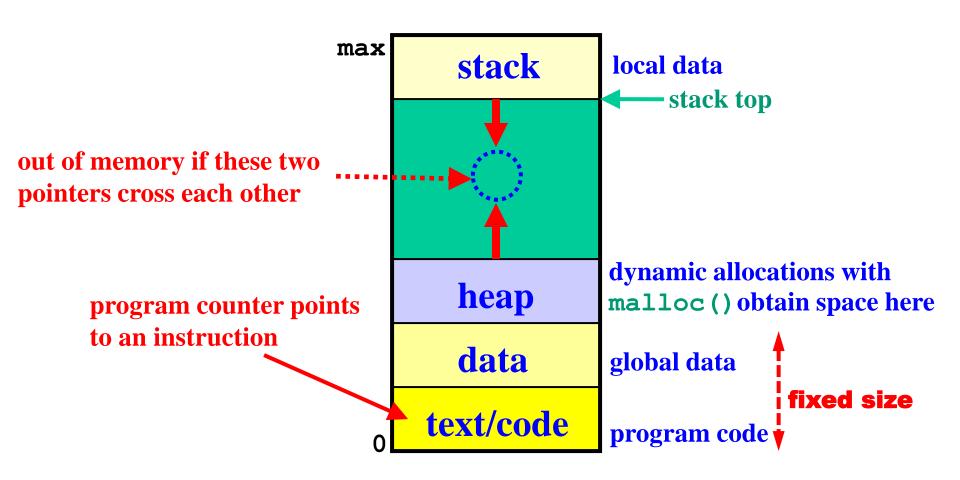


What if A() calls D(), D()'s local variables must be allocated, perhaps overwriting the area previous used by B().

What Have We Learned?

- The call-return sequence is first-in-last-out or last-in-first-out!
- This is a stack!
- To allocate memory for the local variables of functions, we need a STACK.
- In summary, a heap is needed for malloc(), calloc(), etc. (dynamic allocation), and a stack to be used for local variables declared in functions (first-in-last-out).

Process Space

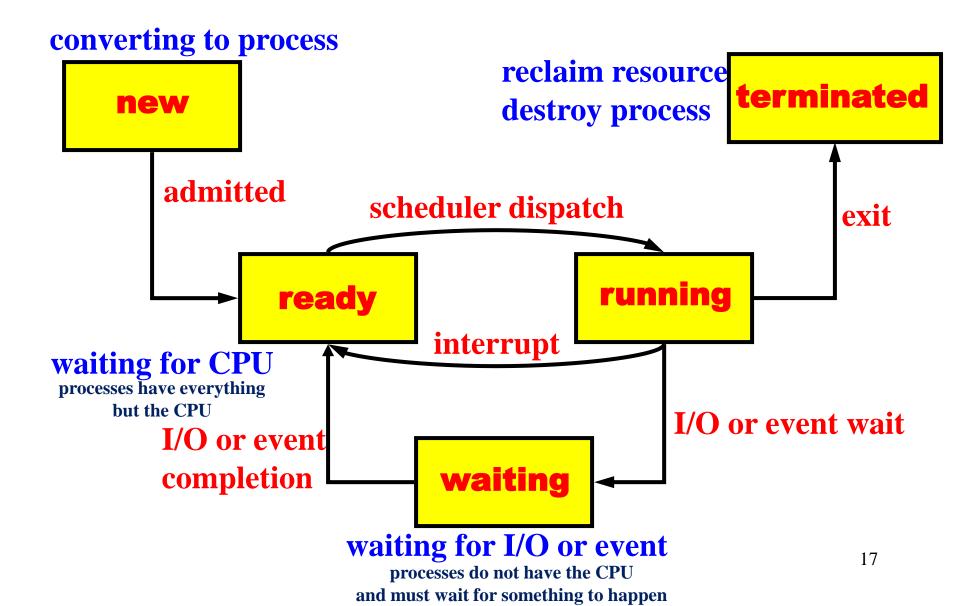


Process States

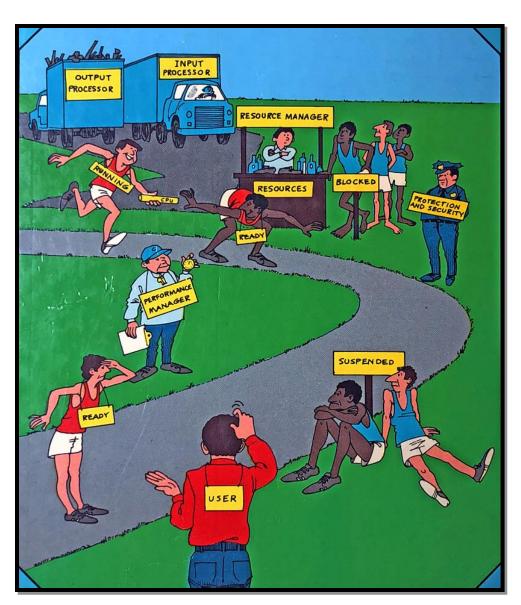
At any moment, a process can be in one of the five states: new, running, waiting, ready and terminated.

- New: The process is being created
- Running: The process is executing on a CPU
- Waiting: The process is waiting for some event to occur (e.g., waiting for I/O completion)
- Ready: The process has everything but the CPU. It is waiting to be assigned to a processor.
- Terminated: The process has finished execution.

Process State Diagram

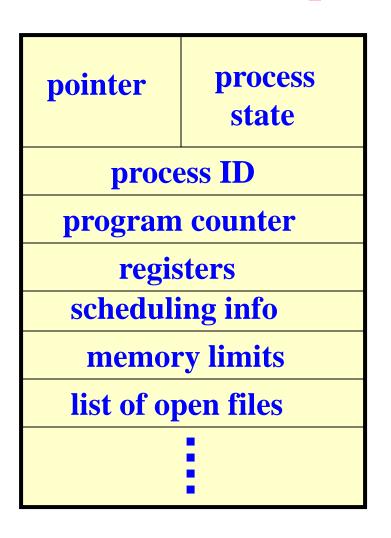


Process State Diagram



Lubomir Bic and Alan C. Shaw,
The Logical Design of
Operating Systems,
Second Edition,
Prentice Hall, 1988

Process Representation in OS



- Each process is assigned a unique number, the process ID.
- Process info are stored in a table, the process control block (PCB).
- These PCBs are chained into several data structures. For example, all processes in the ready state are in the ready queue.

Process Scheduling: 1/2

- Since the number of processes may be larger than the number of available CPUs, the OS must maintain maximum CPU utilization (i.e., all CPU's being as busy as possible).
- To determine which process can do what, processes are chained into a number of scheduling queues.
- For example, in addition to the ready queue, each event may have its own scheduling queue (i.e., waiting queue).

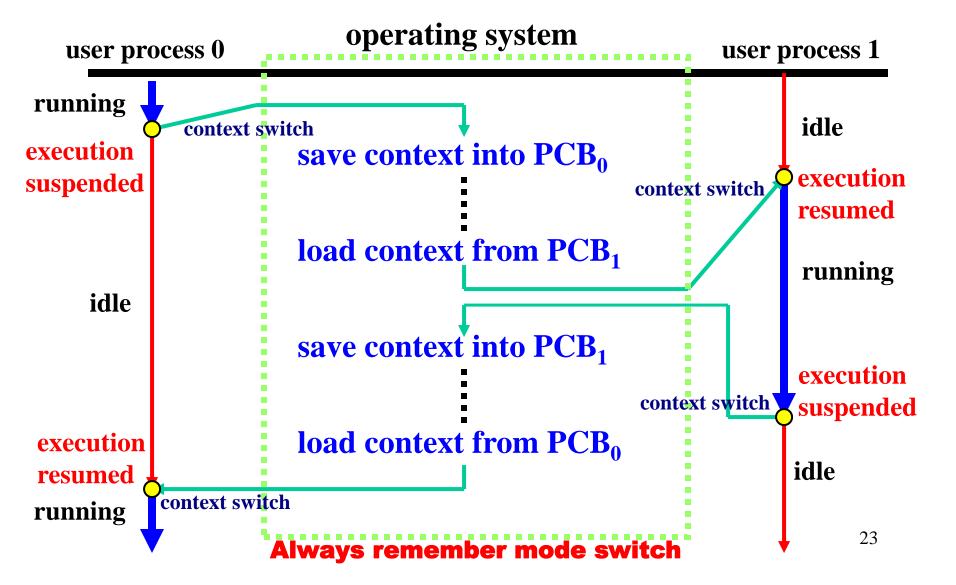
Process Scheduling: 2/2

- The ready queue, which may be organized into several sub-queues, has all processes ready to run.
- The OS has a CPU scheduler.
- When a CPU is free, the CPU scheduler looks at the ready queue, picks a process, and resumes its execution.
- The way of picking a process from the ready queue is referred to as scheduling policy.
- Scheduling policy is unimportant to this course because we cannot make any assumption about it.

Context Switch: 1/2

- What is a process context? The **context** of a process includes process ID, process state, the values of CPU registers, the program counter, and other memory/file management information (i.e., execution environment).
- What is a context switch? After the CPU scheduler selects a process and before allocates a CPU to it, the CPU scheduler must
 - > save the **context** of the currently running process,
 - > put it back to the ready queue or waiting state,
 - > load the context of the selected process, and
 - **>** go to the saved program counter.

Context Switch: 2/2



Operations on Processes

- There are three commonly seen operations:
 - *Process Creation: Create a new process. The newly created is the child of the original. Unix uses fork () to create new processes.
 - **Process Termination:** Terminate the execution of a process. Unix uses exit().
 - **Process Join:** Wait for the completion of a child process. Unix uses wait().
- fork(), exit() and wait() are system calls.

Some Required Header Files

- Before you use processes, include header files sys/types.h and unistd.h.
- sys/types.h has all system data types, and unistd.h declares standard symbolic constants and types.

```
#include <sys/types.h>
#include <unistd.h>
```

The fork() System Call

- The purpose of fork () is to create a child process. The creating and created processes are the parent and child, respectively.
- fork () does not require any argument!
- If the call to fork () is successful, Unix creates an identical but separate address space for the child process to run.
- Both processes start running with the instruction following the fork () system call.

fork() Return Values

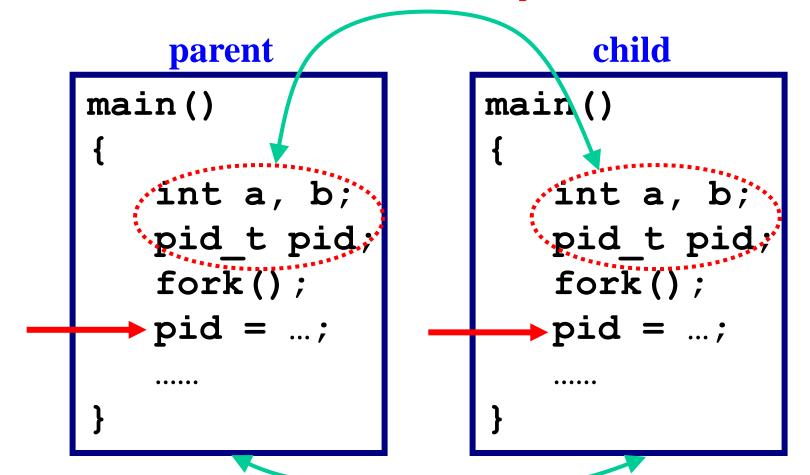
- A negative value means the creation of a child process was unsuccessful.
- A zero means the process is a child.
- Otherwise, fork() returns the process ID of the child process. The ID is of type pid_t.
- Function getpid() returns the process ID of the caller.
- Function getppid() returns the parent's process ID. If the calling process has no parent, getppid() returns 1.

Before the Execution of fork()

parent main() int a, b; pid_t pid → fork();

After the Execution of fork()

in different address spaces



Be Aware of printf(): 1/3

- Processes running from the same terminal window share the same stdout.
- This can cause problems when using printf() due to interleaved execution.
- When printing multiple items using printf() from different processes execution may be switched in the middle.

Be Aware of printf(): 2/3

- Process 1 and Process 2 are shown below.
- After printing A by Process 1, Process 2 may print something and vice versa.
- Both processes print to the same stdout!
- The output could be a mixed one from Process 1 and Process 2.

```
Process 1

A = 10; B = 20;
printf("A=%d B=%d\n",A,B);

X = 100; Y = 200;
printf("X=%d Y=%d\n",X,Y);
```

Be Aware of printf(): 3/3

- Process 1 and Process 2 are shown below.
- Remember interleaved execution.

Process 1	Process 2						
print A							
	print X						
print B							
	print Y						

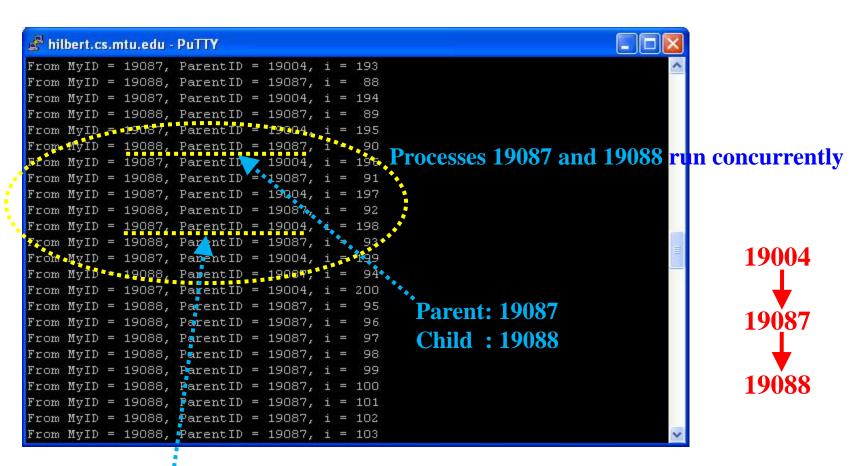
_	1			2	_		1		→	_		2		—	1	<u>_</u>
ſ	A	=	X	=	1	0		В	=	1	0	0	Y	=	2	0
ĺ	2	0	0													
-	—	2	—	-												

we assume that the process is switched out when doing conversions from int to ASCII (%d)

Example 1: 1/2

```
fork-1.c
#include <stdio.h>
#include <string.h>
#include <sys/types.h>
#include <unistd.h>
void main(void)
  pid t MyID, ParentID;
  int i;
  char buf[100];
   fork();
                        // create a child process
  MyID = getpid(); // get my process ID
  ParentID = getppid(); // get my parent's process ID
   for (i = 1; i \le 200; i++) {
      sprintf(buf, "From MyID=%ld, ParentID=%ld, i=%3d\n",
                  MyID, ParentID, i);
     write(1, buf, strlen(buf)); // REMEMBER: why we
                                  // don't use printf?
                                                       33
                  this is stdout
```

Example 1: 2/2



Parent 19087's parent is 19004, the shell that executes fork-1

fork(): A Typical Use

```
main (void)
  pid t pid;
  pid = fork();
  if (pid < 0)
    printf("Oops!");
  else if (pid == 0)
    child();
  else // pid > 0
    parent();
```

```
void child(void)
  int i;
  for (i=1; i<=10; i++)
    printf(" Child:%d\n", i);
  printf("Child done\n");
void parent(void)
  int i;
  for (i=1; i<=10; i++)
    printf("Parent:%d\n", i);
  printf("Parent done\n");
```

Before the Execution of fork()

parent

```
main (void)
 pid = fork();
  if (pid == 0)
    child();
  else
    parent();
void child(void)
void parent(void)
```

After the Execution of fork() 1/2

parent

```
main (void)
               pid=123
  pid = fork();
 if (pid == 0)
               n two different address spaces
else
     child();
  else
    parent();
void child(void)
{ ..... }
void parent(void)
{ ..... }
```

child

```
main (void)
             pid = 0
  pid = fork();
 if (pid == .0)
    child();
    parent();
void child(void)
{ ...... }
void parent(void)
```

After the Execution of fork() 2/2

parent

```
main (void)
             pid=123
  pid = fork();
  if (pid == 0)
    child();
  else
   -parent();
void child(void)
void parent(void)
{ ...... }
```

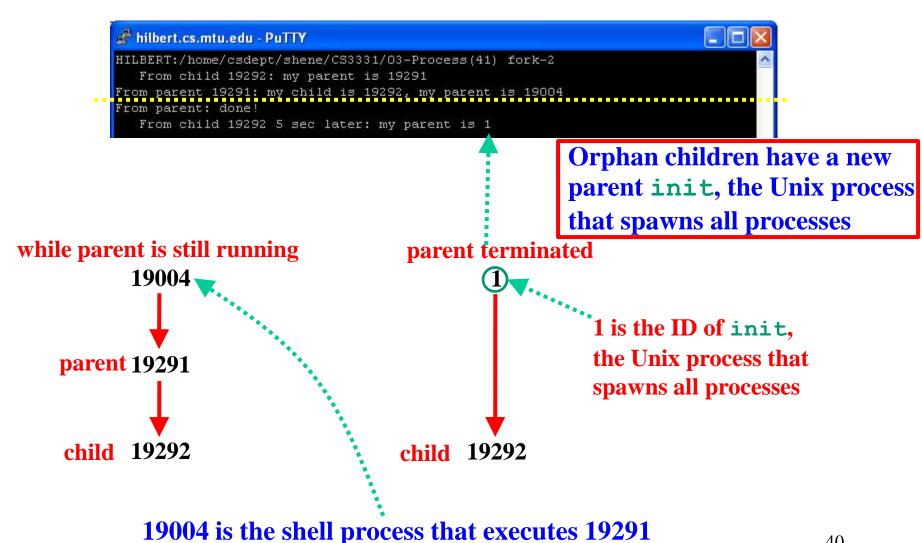
child

```
main (void)
              pid = 0
  pid = fork();
  if (pid == 0)
    -child();
  else
    parent();
void child(void)
{ ...... }
void parent(void)
{ ..... }
```

Example 2: 1/2

```
#include .....
                     // child exits first
                                                     fork-2.c
void main(void)
                     force the child to print after the parent terminates
  pid t pid;
   pid = fork();
   if (pid == 0)
                             // child here
      printf("
                .from child %ld: my parent is %ld\n",
                            getpid(), getppid());
     sleep(5);
      printf("
                 From child %ld 5 sec later: parent %ld\n",
                            getpid(), getppid());
   else {
                             // parent here
    sleep(2);
      printf("From parent %ld: child %ld, parent %ld\n",
                 getpid(), pid, getppid());
      printf("From parent: done!\n");
                                                            39
```

Example 2: 2/2



Example 3: 1/2

```
fork-3.c
#include ..... // separate address spaces
                                  this is equivalent to
void main(void)
                                  pid = fork();
                                  if (pid == 0) ...
  pid t pid;
   char out[100];
   int i = 10, j = 20;
  if ((pid = fork()) == 0) {    // child here
      i = 1000; j = 2000; // child changes values
      sprintf(out, " From child: i=%d, j=%d\n", i, j);
      write(1, out, strlen(out));
                                 // parent here
   else {
     sleep(3);
      sprintf(dat, "From parent: i=%d, j=%d\n", i, j);
      write(1, out, strlen(out));
                 force the parent to print after the child terminated
```

Example 3: 2/2

```
child changes i and j to 1000 and 2000
from 10 and 20, respectively

parent's i and j are not affected because parent and child have independent and separate address space
```

The wait() System Call

- The wait() system call blocks the caller until one of its child processes exits or a signal is received.
- wait() takes a pointer to an integer variable and returns the process ID of the completed process. If no child process is running, wait() returns -1.
- Some flags that indicate the completion status of the child process are passed back with the integer pointer.

How to Use wait()?

Wait for an unspecified child process:

```
wait(&status);
```

Wait for a number, say n, of unspecified child processes:

```
for (i = 0; i < n; i++)
    wait(&status);</pre>
```

Wait for a specific child process whose ID is known:

```
while (pid != wait(&status))
;
```

wait() System Call Example

```
void main(void)
   pid t pid, pid child;
   int status;
   if ((pid = fork()) == 0) // child here
      child();
                              // parent here
   else {
      parent();
      pid child = wait(&status);
```

Zombie Processes: 1/4

- A zombie or defunct process is a process that terminates its execution before its parent executes the wait() system call.
- A zombie process only has an entry in the system process table and does not consume much system resource.
- A zombie process will be removed completely when its parent terminates or when its parent executes the wait() system call waiting for this child.

Zombie Processes: 2/4

- Between the termination of the child and the termination of its parent (or the parent executes the wait() system call), the child is a zombie or defunct process.
- A zombie process is in the terminated state.
- The kill command cannot kill zombie processes because they had died (terminated)!
- To find out whether you have zombie processes, use ps -af.

Zombie Processes: 3/4

Suppose we have the following code:

```
if (fork() > 0)
    if (fork() > 0)
        for (fork() > 0)
        while (1 == 1) ;
```

- The parent creates three child processes and loops forever (i.e., never dies). Each child process does nothing and terminates immediately.
- Therefore, we have three zombie processes!

Zombie Processes: 4/4

We have three zombie processes, marked as defunct.

```
COLOSSUS:/home/campus13/shene(30) ps -af
          PID PPID C STIME TTY
UID
                                          TIME CMD
         1313 871 99 01:06 pts/1
                                      00:00:42 zom
shene
         1314 1313 0 01:06 pts/1
                                     00:00:00 [zom] <defunct>
shene
         1315 1313 0 01:06 pts/1
                                      00:00:00 [zom] <defunct>
shene
                                      00:00:00 [zom] <defunct>
         1316 1313 0 01:06 pts/1
shene
         1375 1201 0 01:06 pts/2
                                      UU:UU:UU ps -ai
shene
COLOSSUS:/home/campus13/shene(31)
```

You need to kill the parent to kill its child zombies.

```
COLOSSUS:/home/campus13/shene(30) ps -af
                            PID PPID C STIME TTY
                                                            TIME CMD
                 UID
                           1313
                                  871 99 01:06 pts/1
                 shene
                                                        00:00:42 zom
    process 1313 shene
                           1314 1313 0 01:06 pts/1
                                                        00:00:00 [zom] <defunct>
                                1313 0 01:06 pts/1
                                                        00:00:00 [zom] <defunct>
                 shene
                                                                                 kill the parent 1313
   created 1314,
                           1316 1313 0 01:06 pts/1
                                                        00:00:00 [zom] <defunct>
  1315 and 1316, shene
                                1201 0 01:06 pts/2
                                                        00:00:00 ps -af
       the three COLOSSUS:/home/campus13/shene(31) kill -KILL 1313
                 COLOSSUS:/home/campus13/shene(32) ps -af
zombie processes
                 UID
                                PPID C STIME TTY
                            PID
                                                            TIME CMD
                           1455
                                1201 0 01:08 pts/2
                                                        00:00:00 ps -af
                 shene
                 COLOSSUS:/home/campus13/shene(33)
```

Daemon Processes: 1/2

- A daemon (DEE-mun or DAY-mun) process is a background process that is detached from a terminal.
- They are used to answer requests and/or to provide services. Commonly seen ones are inetd, httpd, nfsd, sshd, lpd, syslogd ...
- Daemon processes are usually created when the system started by the init process, the mother of all processes with PID 1.
- It could also be created by a process that immediately exits after creation, making the created process adapted by init.

Daemon Processes: 2/2

- Daemon processes usually have PPID 1 (the init process) and have names like *d.
 this is the PPID column
- Not all process names end with d are daemons.
- You may use ps -ef | awk \\$3 == 1' to list all processes whose parent is init:

```
root PID 2124
                1 0 2018 ?
                                    00:00:00 /usr/sbin/lvmetad -f
              1 0 2018 ?
         2159
                                    00:00:03 /usr/lib/systemd/systemd-udevd
root
              1 0 2018 ?
        4158
                                    00:00:26 /sbin/auditd
root
                                    00:03:18 /usr/sbin/irqbalance - foreground
              1 0 2018 ?
        4185
root
              1 0 2018 ?
        4187
                                    00:00:04 /sbin/rpcbind -w
rpc
                                    00:00:24 /usr/libexec/rtkit-daemon
                 1 0 2018 ?
rtkit
        4188
              1 0 2018 ?
         4221
                                    00:00:00 /usr/sbin/smartd/-n -q\never
root
                 1 0 2018 ?
         4951
                                    00:00:01 rhnsd
root
         4961
                 1 0 2018 ?
                                    00:00:22 /usr/sbin/ypbind
root
```

The exec() System Calls

- A newly created process may run a different program rather than that of the parent.
- This is done using the exec system calls. We will only discuss execup():

```
int execvp(char *file, char *argv[]);
```

- *file is a char array that contains the name of an executable file. Depending on your system settings, you may need the . / prefix for files in the current directory.
- *argv[] is the argument passed to your main program
- *argv[0] is a pointer to a string that contains the program name
- *argv[1], argv[2], ... are pointers to strings that contain the arguments

execvp(): An Example 1/2

```
argv[]
#include <stdio.h>
#include <unistd.h>
void main(void)
 char prog[] = { "cp" };
 char in[] = { "this.c" };
 char out[] = { "that.c" }; ......
 char *argv[4];
 int status;
 pid t pid;
 argv[0] = prog; argv[1] = in;
 argv[2] = out; argv[3] = '\0';
          // see next slide
```

execvp(): An Example 2/2

```
if ((pid = fork()) < 0) {
                                      argv[]
    printf("fork() failed\n");
    exit(1);
 else if (pid == 0)
    if (execvp(prog, argv) < 0) {</pre>
       printf("execup() failed\n"); t h i s
       exit(1);
 else
    wait(&status);
                   execute program cp
```

A Mini-Shell: 1/3

```
void parse(char *line, char **argv)
   while (*line != '\0') { // not EOLN
      while (*line == ' ' || *line == '\t' || *line == '\n')
         *line++ = '\0'; // replace white spaces with 0
      *argv++ = line; // save the argument position
      while (*line != '\0' && *line ! =' '
                           && *line!='\t' && *line != '\n')
         line++;
                          // skip the argument until ...
   *argv = ' \ 0';
                          // mark the end of argument list
line[]
                    t h a t
                   line[]
                                        \0 t h a t
                           t h i s
                                                         55
                             arqv
```

A Mini-Shell: 2/3

```
void execute(char **argv)
{
   pid t pid;
   int status;
   if ((pid = fork()) < 0) { // fork a child process
      printf("*** ERROR: forking child process failed\n");
      exit(1);
   else if (pid == 0) {     // for the child process:
      if (execup (*argv, argv) < 0) { // execute the command
         printf("*** ERROR: exec failed\n");
         exit(1);
   else {
                             // for the parent:
      while (wait(&status) != pid) // wait for completion
```

A Mini-Shell: 3/3

```
void main(void)
{
   char line[1024]; // the input line
   char *argv[64]; // the command line argument
  while (1) { // repeat until done ....
     printf("Shell -> "); // display a prompt
      gets(line); // read in the command line
      printf("\n");
     parse(line, argv); // parse the line
      if (strcmp(argv[0], "exit") == 0) // is it an "exit"?
        exit(0); // exit if it is
      execute(argv); // otherwise, execute the command
        Don't forget that gets () is risky! Use fgets () instead.
```

What is Shared Memory?

- The parent and child processes are run in independent and separate address spaces. All processes, parent and children included, do not share anything.
- A shared memory segment is a piece of memory that can be allocated and attached to an address space. Only those processes that have this memory segment attached can have access to it.
- But, race conditions can occur!

Procedure for Using Shared Memory

- Find a **key**. Unix uses this key for identifying shared memory segments.
- Use shmget () to allocate a shared memory.
- Use shmat() to attach a shared memory to an address space.
- Use shmdt () to detach a shared memory from an address space.
- Use shmctl() to deallocate a shared memory.

Keys: 1/2

To use shared memory, include the following:

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>
```

- A key is a value of type key_t. There are three ways to generate a key:
 - **❖**Do it yourself
 - Use function ftok ()
 - *Ask the system to provide a private key.

Keys: 2/2

Do it yourself:

```
key_t SomeKey;
SomeKey = 1234;
```

Use ftok () to generate one for you:

```
key_t = ftok(char *path, int ID);
*path is a path name(e.g., "./")
*ID is a small integer(e.g., 'a')
*Function ftok() returns a key of type key_t:
SomeKey = ftok("./", 'x');
```

- Keys are global entities. If other processes know your key, they can access your shared memory.
- Ask the system to provide a private key with IPC PRIVATE.

Ask for a Shared Memory: 1/4

• Include the following:

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>
```

Use shmget () to request a shared memory:

```
shm_id = shmget(
   key_t key, /* identity key */
   int size, /* memory size */
   int flag); /* creation or use */
```

- shmget() returns a shared memory ID.
- The flag means permission, which is either IPC_CREAT | 0666 for creation or 0666 for use only, where 0xxx is a Unix permission.62

Ask for a Shared Memory: 2/4

The following creates a shared memory of size struct Data with a private key IPC_PRIVATE. This is a creation (IPC_CREAT) with read and write permission (0666).

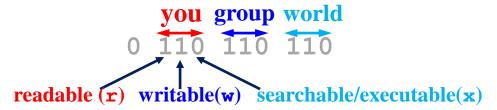
Ask for a Shared Memory: 3/4

The following creates a shared memory with a key based on the current directory:

```
struct Data { int a; double b; char x;};
int ShmID;
key t Key;
Key = ftok("./", 'h');
ShmID = shmget(
         Key, /* a key */
          sizeof(struct Data),
         IPC CREAT | 0666);
```

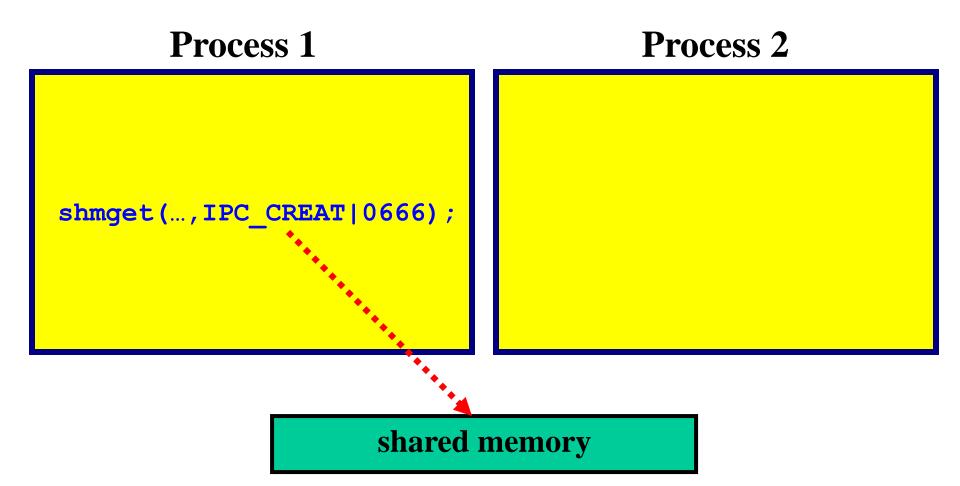
Ask for a Shared Memory: 4/4

- When asking for a shared memory, the process that creates it uses IPC_CREAT | 0666 and processes that access a created one use 0666.
- 0666 is an octal number with a bit pattern



- If the return value is < 0 (Unix convention), the request was unsuccessful, and no shared memory is allocated.
- Create a shared memory before its use!

After the Execution of shmget()



A shared memory is allocated; but is not part of the address space

Attaching a Shared Memory: 1/4

Use shmat () to attach an existing shared memory to an address space:

```
shm_ptr = shmat(
   int shm_id, /* ID from shmget() */
   char *ptr, /* use NULL here */
   int flag); /* use 0 here */
```

- shm_id is the shared memory ID returned by shmget().
- Use NULL and 0 for the second and third arguments, respectively.
- shmat() returns a void pointer to the memory. If unsuccessful, it returns a negative integer.

Attaching a Shared Memory: 2/4

```
struct Data { int a; double b; char x;};
int ShmID;
key t Key;
struct Data *p;
Key = ftok("./", 'h');
ShmID = shmget(Key, sizeof(struct Data),
            IPC CREAT | 0666);
p = (struct Data *) shmat(ShmID, NULL, 0);
if ((int) p < 0) {
   printf("shmat() failed\n"); exit(1);
p->a = 1; p->b = 5.0; p->x = '.';
```

Attaching a Shared Memory: 3/4

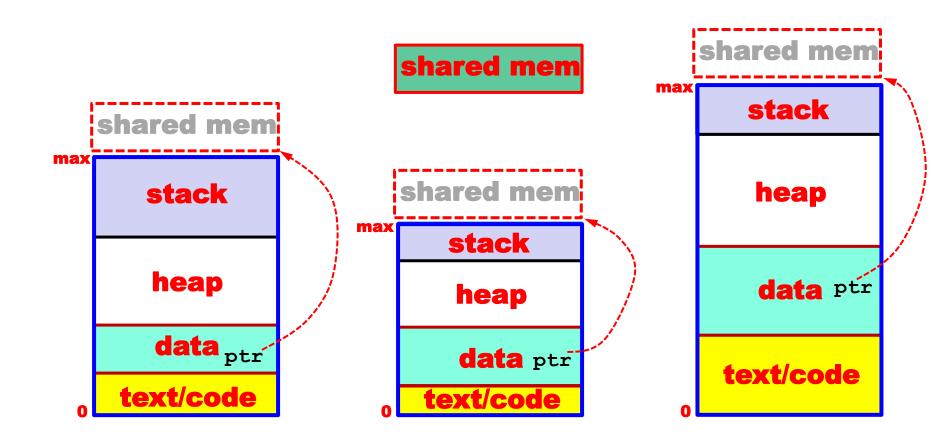
Process 1 Process 2 shmget(...,IPC CREAT|0666); shmget(...,0666); ptr = shmat(.....); ptr = shmat(.....); ptr shared memory Now processes can access the shared memory

Are the pointers returned to Process 1

and Process 2 the same?

Attaching a Shared Memory: 4/4

Are the pointers returned to the attaching processes the same?



Detaching/Removing Shared Memory

To detach a shared memory, use

```
shmdt(shm_ptr);
shm_ptr is the pointer returned by shmat().
```

- After a shared memory is detached, it is still there. You can attach and use it again.
- To remove a shared memory, use

```
shmctl(shm_ID, IPC_RMID, NULL);
shm_ID is the shared memory ID returned by
shmget(). After a shared memory is removed,
it no longer exists.
```

Communicating with a Child: 1/2

```
void main(int argc, char *argv[])
  int ShmID, *ShmPTR, status;
 pid t pid;
  ShmID = shmget(IPC PRIVATE, 4*sizeof(int), IPC CREAT | 0666);
  ShmPTR = (int *) shmat(ShmID, NULL, 0);
  ShmPTR[0] = getpid(); ShmPTR[1] = atoi(argv[1]);
  ShmPTR[2] = atoi(argv[2]); ShmPTR[3] = atoi(argv[3]);
  if ((pid = fork()) == 0) {
    Child(ShmPTR);
    exit(0);
                               parent's process ID here
  wait(&status);
  shmdt((void *) ShmPTR); shmctl(ShmID, IPC RMID, NULL);
  exit(0);
                                                      72
```

Communicating with a Child: 2/2

• Why are shmget() and shmat() not needed in the child process?

Communicating Among Separate Processes: 1/5

Define the structure of a shared memory segment as follows:

```
#define NOT_READY (-1)
#define FILLED (0)
#define TAKEN (1)

struct Memory {
  int status;
  int data[4];
};
```

Communicating Among Separate Processes: 2/5

```
The "Server"
                              Prepare for a shared memory
int main(int argc, char *argv[])
   key t
                 ShmKEY;
   int
                  ShmID, i;
   struct Memory *ShmPTR;
   ShmKEY = ftok("./", 'x');
   ShmID = shmget(ShmKEY, sizeof(struct Memory),
                  IPC CREAT | 0666);
   ShmPTR = (struct Memory *) shmat(ShmID, NULL, 0);
```

Communicating Among Separate Processes: 3/5

```
shared memory not ready
ShmPTR->status = NOT READY;
                              filling in data
ShmPTR->data[0] = getpid();
for (i = 1; i < 4; i++)
   ShmPTR->data[i] = atoi(argv[i]);
ShmPTR->status = FILLED;
while (ShmPTR->status != TAKEN)
   sleep(1); /* sleep for 1 second */
printf("My buddy is %ld\n", ShmPR->data[0]);
shmdt((void *) ShmPTR);
shmctl(ShmID, IPC RMID, NULL);
exit(0);
                            wait until the data is taken
```

Communicating Among Separate Processes: 4/5

```
int main(void)
                                        The "Client"
  key t
                  ShmKEY;
                                   prepare for shared memory
   int
                  ShmID;
   struct Memory
                 *ShmPTR;
   ShmKEY=ftok("./", 'x');
   ShmID = shmget(ShmKEY, sizeof(struct Memory), 0666);
   ShmPTR = (struct Memory *) shmat(ShmID, NULL, 0);
   while (ShmPTR->status != FILLED)
  printf("%d %d %d %d\n", ShmPTR->data[0],
      ShmPTR->data[1], ShmPTR->data[2], ShmPTR->data[3]);
   ShmPTR->data[0] = getpid();
   ShmPTR->status = TAKEN;
   shmdt((void *) ShmPTR);
   exit(0);
                                                     77
```

Communicating Among Separate Processes: 5/5

- The "server" must run first to prepare a shared memory.
- Run the server in one window, and run the client in another a little later.
- Or, run the server as a background process. Then, run the client in the foreground later:

```
server 1 3 5 &
client
```

- This version uses busy waiting.
- One may use Unix semaphores for mutual exclusion.

Important Notes: 1/4

- If you do not remove your shared memory segments (e.g., program crashes before the execution of shmctl()), they will be in the system forever until the system is shut down. This will degrade the system performance.
- Use the ipcs command to check if you have shared memory segments left in the system.
- Use the ipcrm command to remove your shared memory segments.

Important Notes: 2/4

- Checking existing shared memory segments in the system, use ipcs -m, where m means shared memory.
- The following is a snapshot on wopr:

```
♣ shene@wopr:~

[shene@wopr ~] $ ipcs -m
 ---- Shared Memory Seaments
           shmid
                                             bytes
                                                         nattch
                      owner
                                                                    status
                                  perms
0x78181367 1573912576 machoudh
                                  666
                                             12
                                  666
                                             204
0x7817433c 1336737793 hyunjik
0x78181363 1575583746 machoudh
                                  666
                                             12
0x7818132a 1577582595 machoudh
                                  666
                                             12
0x781813da 1579515908 machoudh
                                  666
0x6b179e35 1612414981 mswillia
                                  666
0x6b18b8b0 1909686278 machoudh
                                  666
                                             40
0x7918299c 1910013959 machoudh
                                  666
                                             92
[shene@wopr ~]$
```

Important Notes: 3/4

- Checking existing shared memory segments in the system, use ipcs -m, where m means shared memory.
- The following is a snapshot on colossus:

shared memory segment marked to be destroyed

```
COLOSSUS:/home/campus13/shene(26) ipcs -m
 ---- Shared Memory Segments -----
                                                     nattch
kev
          shmid
                     owner
                                perms
                                           bytes
                                                                status
                               600
0x00000000 0
                     zabbix
                                           576
                                                                dest
0x00000000 1
                     zabbix
                               600
                                           4745056
                                                     3
                                                                dest
0x00000000 4
                     gdm
                               777
                                           16384
                                                                dest
                                                     2
0x00000000 7
                                777
                                           3145728
                                                                dest
                     adm
COLOSSUS:/home/campus13/shene(27)
```

Important Notes: 4/4

To remove a shared memory, use the ipcrm command as follows:

```
❖ipcrm -M shm-key❖ipcrm -m shm-ID
```

You have to be the owner (or super user) to remove a shared memory.

The End