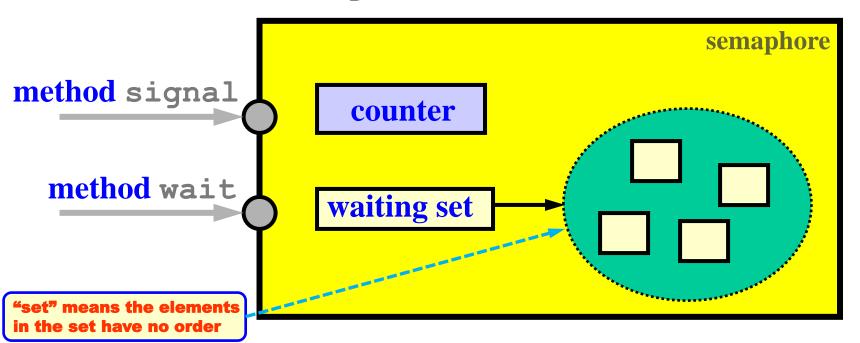
Part III Synchronization Semaphores

The bearing of a child takes nine months, no matter how many women are assigned.

Semaphores

A semaphore is an object that consists of a private counter, a private waiting set of threads, and two public methods (e.g., member functions): signal and wait.



Semaphore Method: wait

```
void wait(sem S)
{
    S.count--;
    if (S.count < 0) {
        add the caller to the waiting set;
        block();
    }
}</pre>
```

- After decreasing the counter by 1, if the new value becomes negative, then
 - add the caller to the waiting set, and
 - *block the caller.

Semaphore Method: signal

```
void signal(sem S)
{
    S.count++;
    if (S.count <= 0) {
        remove a thread T from the waiting set;
        resume(T);
    }
}</pre>
```

- After increasing the counter by 1, if the new value is not positive (e.g., non-negative), then
 - *remove a thread T from the waiting set,
 - *resume the execution of thread T, and return

Practice Example: 1/8

```
void wait(sem S)
{
    S.count--;
    if (S.count < 0) {
        add the caller to the waiting set;
        block();
    }
}</pre>
void signal(sem S)
{
    S.count++;
    if (S.count <= 0) {
        remove a thread T from the waiting set;
        resume(T);
    }
}
```

- Suppose we have four threads, A, B, C and D.
- We also have a semaphore S with initial value 2.
- What would happen if A, B, C and D calls wait (S) in this order?

Practice Example: 2/8

```
void wait(sem S)
{
    S.count--;
    if (S.count < 0) {
        add the caller to the waiting set;
        block();
    }
}</pre>
void signal(sem S)
{
    S.count++;
    if (S.count <= 0) {
        remove a thread T from the waiting set;
        resume(T);
    }
}
```

A	В	C	D	semaphore counter	semaphore waiting set
				2	Ø
wait(S)					
		wait(S)			
			wait(S)		

Practice Example: 3/8

```
void wait(sem S)
{
    S.count--;
    if (S.count < 0) {
        add the caller to the waiting set;
        block();
    }
}</pre>
void signal(sem S)
{
    S.count++;
    if (S.count <= 0) {
        remove a thread T from the waiting set;
        resume(T);
    }
}
```

A	В	С	D	semaphore counter	semaphore waiting set
				2	Ø
wait(S)				1	Ø
		wait(S)			
V			wait(S)		

Practice Example: 4/8

```
void wait(sem S)
{
    S.count--;
    if (S.count < 0) {
        add the caller to the waiting set;
        block();
    }
}</pre>
void signal(sem S)
{
    S.count++;
    if (S.count <= 0) {
        remove a thread T from the waiting set;
        resume(T);
    }
}
```

A	В	C	D	semaphore counter	semaphore waiting set
				2	Ø
wait(S)				1	Ø
	wait(S)			0	Ø
		wait(S)			
V	*		wait(S)		

Practice Example: 5/8

```
void wait(sem S)
{
    S.count--;
    if (S.count < 0) {
        add the caller to the waiting set;
        block();
    }
}</pre>
void signal(sem S)
{
    S.count++;
    if (S.count <= 0) {
        remove a thread T from the waiting set;
        resume(T);
    }
}
```

A	В	С	D	semaphore counter	semaphore waiting set
				2	Ø
wait(S)				1	Ø
	wait(S)			0	Ø
		wait(S) (blocked)		-1	C
¥	*		wait(S)		

Practice Example: 6/8

```
void wait(sem S)
{
    S.count--;
    if (S.count < 0) {
        add the caller to the waiting set;
        block();
    }
}</pre>
void signal(sem S)
{
    S.count++;
    if (S.count <= 0) {
        remove a thread T from the waiting set;
        resume(T);
    }
}
```

A	В	C	D	semaphore counter	semaphore waiting set
				2	Ø
wait(S)				1	Ø
	wait(S)			0	Ø
		wait(S) (blocked)		-1	C
•	•		wait(S) (blocked)	-2	C and D

10

Practice Example: 7/8

```
void wait(sem S)
{
    S.count--;
    if (S.count < 0) {
        add the caller to the waiting set;
        block();
    }
}</pre>
void signal(sem S)
{
    S.count++;
    if (S.count <= 0) {
        remove a thread T from the waiting set;
        resume(T);
    }
}
```

What would happen if A calls signal (S) and then B calls wait (S)?

Practice Example: 8/8

```
void wait(sem S)
{
    S.count--;
    if (S.count < 0) {
        add the caller to the waiting set;
        block();
    }
}</pre>
void signal(sem S)
{
    S.count++;
    if (S.count <= 0) {
        remove a thread T from the waiting set;
        resume(T);
    }
}
```

A	В	С	D	semaphore counter	semaphore waiting set
		wait(S)	wait(S)	-2	C and D
signal(S)			D relea	nsed -1	C
V	wait(S) (blocked)			-2	B and C

- Which one of C and D is released when A calls signal (S)?
- We don't know, because there is no ordering in the waiting set.
- Let us say D is lucky and released.-
- B is blocked.

Important Note: 1/4

```
S.count--;
    if (S.count<0) {
        add to set;
        block();
    }
    S.count++;
    if (S.count<=0) {
        remove T;
        resume(T);
    }
}</pre>
```

- If S.count < 0, abs (S.count) is the number of waiting threads.
- This is because threads are added to (resp., removed from) the waiting set only if the counter value is < 0 (resp., <= 0).

Important Note: 2/4

```
S.count--;
    if (S.count<0) {
        add to set;
        block();
    }
    S.count++;
    if (S.count<=0) {
        remove T;
        resume(T);
    }
}</pre>
```

- The waiting set can be implemented with a queue if FIFO order is desired.
- However, the correctness of a program should not depend on a particular implementation (e.g., ordering) of the waiting set.

Important Note: 3/4

```
S.count--;
    if (S.count<0) {
        add to set;
        block();
    }
    S.count++;
    if (S.count<=0) {
        remove T;
        resume(T);
    }
}</pre>
```

- The caller may be blocked in the call to wait().
- The caller is never blocked in the call to signal(). If S.count > 0, signal() returns and the caller continues. Otherwise, a waiting thread is released, and the caller continues. In this case, two threads continue.

The Most Important Note: 4/4

```
S.count--;
    if (S.count<0) {
        add to set;
        block();
    }
    S.count++;
    if (S.count<=0) {
        remove T;
        resume(T);
    }
}</pre>
```

- wait() and signal() with respect to the same semaphore must be executed atomically. Otherwise, race conditions may occur.
- Homework: use execution sequences to show race conditions if wait() and/or signal() is not executed atomically. Also show that mutual exclusion cannot be guaranteed.

Typical Uses of Semaphores

- There are three typical uses of semaphores:
 - *****mutual exclusion:

Mutex (i.e., Mutual Exclusion) locks

***count-down lock:**

Keep in mind that a semaphore has a private counter that can count.

*notification:

Wait for an event to occur and indicate an event has occurred.

Use 1: Mutual Exclusion (Lock)

```
initialization is important
semaphore(S = 1)
           count = 0; // shared variable
int
     Thread 1
                                   Thread 2
while (1) {
                            while (1) {
   // do something entry
                                   do something
   S.wait();
                                S.wait()
                                S.signal();
   S.signal();
   // do something
                       exit
                                // do something
```

- What if the initial value of S is zero?
- S is a binary semaphore (count being 0/FALSE or 1/TRUE)

Discussion: 1/2

- Note that the execution of S.wait() and S.signal() must be atomic. When multiple calls to wait() and signal() at the same time, the execution of these calls is sequential.
- Therefore, these two S.wait() calls cannot be executed at the same time. One will happen before the other.

Discussion: 2/2

- Mutual Exclusion: Because only 1 thread can call wait() at a time, if thread 1 gets into the critical section, the semaphore counter of S is 0. The next thread that calls wait() blocks. Thus, mutual exclusion holds!
- Progress: Suppose the critical section is empty and threads 1 and 2 are waiting to enter. Then, the system selects one thread to call wait(). The execution of wait() only needs finite instructions. Both take finite time!
- Outsider Issue? Bounded Waiting?

Use 2: Count-Down Counter

```
semaphore
      Thread 1
                                   Thread 2
while (1) {
                            while (1) {
   // do something
                                // do something
   S.wait();
                                S.wait();
         at most 3 threads can be here!!!
   S.signal();
                                S.signal();
   // do something
                                // do something
```

• After three threads passing through wait(), this section is locked until a thread calls signal().

Use 3: Notification

- Thread 1 uses S2.signal() to notify thread 2, indicating "I am done. Please go ahead."
- Thread 2 uses S1.signal() to notify thread 1, indicating "I am done. Please go ahead."
- The output is 1 2 1 2 1 2

Discussion: 1/10

```
semaphore S1 = 1, S2 = 0;

Thread 1
while (1) {
    // do something
    S1.wait();
    cout << "1";
    S2.signal();
    // do something
}
</pre>
S2.wait();
    cout << "2";
    S1.signal();
    // do something
}</pre>
```

• If T_2 reaches S2.wait() **before** T_1 reaches S2.signal().

$\mathbf{T_1}$	$\mathbf{T_2}$	S1	S2	Comment
		1	0	Initial Values
S1.wait()		0	0	T ₁ continues
į	S2.wait()	0	-1	T ₂ blocks
cout<<"1"		0	-1	
S2.signal()'		0	0	T ₁ signals T ₂
S1.wait()	cout<<"2"	-1	0	T ₂ is released & prints
	S1.signal()	0	0	T ₂ signals T ₁
cout<<"1"	S2.wait()	0	-1	T ₂ waits

Discussion: 2/10

```
semaphore S1 = 1, S2 = 0;

Thread 1
while (1) {
    // do something
    S1.wait();
    cout << "1";
    S2.signal();
    // do something
}
</pre>
S1.signal();
    // do something
}
```

If T₂ reaches S2.wait() after T₁ reaches S2.signal().

$\mathbf{T_1}$	T_2	S1	S2	Comment
		1	0	Initial Values
S1.wait()		0	0	
cout<<"1"		0	0	
S2.signal()		0	1	T ₁ signals
	S2.wait()	0	0	T ₂ continues
S1.wait()	cout<<"2"	-1	0	T ₁ blocks
	S1.signal()	0	0	T ₂ signals T ₁
cout<<"1"	S2.wait()	0	-1	T ₂ waits

Discussion: 3/10

```
semaphore S1 = 0, S2 = 1;
Thread 1
while (1) {
    // do something
    S1.wait();
    cout << "1";
    S2.signal();
    // do something
}</pre>
S1.signal();
    // do something
}
```

What if the initial values are changed to S1 = 0 and S2 = 1?

Discussion: 4/10

```
semaphore S1 = 0, S2 = 1;

Thread 1
while (1) {
    // do something
    S1.wait();
    cout << "1";
    S2.signal();
    // do something
}
</pre>
S2.wait();
    cout << "2";
    S1.signal();
    // do something
}</pre>
```

- What if the initial values are changed to S1 = 0 and S2 = 1?
- Because the code for Thread 1 and Thread 2 are symmetric, the output is 2 1 2 1 2 1!

Discussion: 5/10

```
semaphore S1 = 0, S2 = 0;
Thread 1
while (1) {
    // do something
    S1.wait();
    cout << "1";
    S2.signal();
    // do something
}</pre>
S1.signal();
    // do something
}
```

What if the initial values are changed to S1 = 0 and S2 = 0?

Discussion: 6/10

```
semaphore S1 = 0, S2 = 0;
Thread 1
while (1) {
    // do something
    S1.wait();
    cout << "1";
    S2.signal();
    // do something
}</pre>
S1.signal();
    // do something
}
```

• What if the initial values are changed to S1 = 0 and S2 = 0?

Deadlock!

- ➤ Both threads are blocked by their wait() calls, waiting for the other thread to call signal().
- Each of these two threads are blocked by an event that can only be caused to happen by the other (waiting) thread.

Discussion: 7/10

```
Thread 1
while (1) {
    // do something
    S1.wait();
    cout << "1";
    S2.signal();
    // do something
}</pre>

Thread 2
while (1) {
    // do something
    S2.wait();
    cout << "2";
    S1.signal();
    // do something
}
```

What if the initial values are changed to S1 = 1 and S2 = 1?

Discussion: 8/10

```
Thread 1
while (1) {
    // do something
    S1.wait();
    cout << "1";
    S2.signal();
    // do something
}
</pre>

Thread 2
while (1) {
    // do something
    S2.wait();
    cout << "2";
    S1.signal();
    // do something
}
</pre>
```

- What if the initial values are changed to S1 = 1 and S2 = 1?
 - ➤ Because semaphores S1 and S2 have initial value 1, both threads can pass their wait() calls in an unpredictable order.
 - ➤ As a result, the order of printing 1 or 2 is also not predictable.
 - The output from Thread 1 and Thread 2 is kind of random.

Discussion: 9/10

```
semaphore S1 = 1, S2 = 1;

Thread 1
while (1) {
    // do something
    S1.wait();
    cout << "1";
    S2.signal();
    // do something
}
</pre>

S1.signal();
    // do something
}
```

- What if the initial values are changed to S1 = 1 and S2 = 1?
 - ➤ It is easy to get two consecutive 1's (or 2's)
 - ➤ **REASON:** If Thread 1 calls S2.signal() before Thread 2 reaches its S2.wait(), the counter of semaphore S2 is 2.
 - ➤ Therefore, Thread 2 can print two consecutive 2's in its output.
 - ➤ Because the code is symmetric, Thread 1 can print two consecutive 1's in its output.

Discussion: 10/10

T_1	T_2	s1	S2	Comment
		1	1	
S1.wait()		0	1	T ₁ gets through
cout<<"1"		0	1	
S2.signal() 0	utput =1 2 2 1	- ··· 0	2	T ₁ signals T ₂
	S2.wait()	0	1	T ₂ does not wait
	cout<<"2"	0	1	
	S1.signal()	1	1	T ₂ signals T ₁
	S2.wait()	1	0	T ₂ does not wait
	cout<<"2"	1	0	
	S1.signal()	2	0	T ₂ signals T ₁
	S2.wait()	2	-1	T ₂ must wait
S1.wait()		1	-1	T ₁ does not wait
cout<<"1"		1	-1	
S2.signal()		1	0	T ₁ signals T ₂
	T ₂ released			

Food for Thought

- Is it possible that Thread 1 and Thread 2 can print 3 consecutive 1's or 2's such as ... 1 2 2 2 1...?
- Suppose Thread i can only print the value of i. Modify the sample code by adding semaphores and threads to do the following:

```
> Print 1 2 2 1 2 2 1 2 2 1 ...
```

- > Print 1 2 3 1 2 3 1 2 3 ...
- > Print 1 2 3 2 1 2 3 2 1 ...
- > Print 1 2 3 2 3 1 2 3 2 3 1 ...

Dining Philosophers

Dining Philosophers: Its Origin


The Dining Philosophers problem, originally called Dining Quintuple – later dubbed "The Dining Philosophers" by Tony Hoare – was a problem in the final exam of a course taught by Edsger W Dijkstra.

FMD 1000-5

EWD123, "Cooperating Sequential Processes" was written in 1965 and served as lecture notes for my course in the fall semester of that year. (The problem of The Dining Quintuple -later dubbed "The Dining Philosophers" by Tony Hoare - was the examination problem at the end of that semester.) The rate with which the EWD-numbers increased was in those days not a measure of my productivity: I assigned numbers when I started on manuscripts and many of them were not completed.

Dining Philosophers

- Five philosophers are in a thinking - eating cycle.
- When a philosopher gets hungry, he sits down, picks up his left and then his right chopsticks, and eats.
- A philosopher can eat only if he has both chopsticks.
- After eating, he puts down both chopsticks and thinks.
- This cycle continues.



Dining Philosopher: Ideas


- Each philosopher is a thread.
- Chopsticks are shared items (by two neighboring philosophers) and must be protected.
- Each chopstick must be used in a mutually exclusive way (i.e., used in the critical section of that chopstick).
- Each chopstick has a semaphore with initial value 1 for mutual exclusion.
- A philosopher calls wait() to pick up a chopstick and calls signal() to release it.

Dining Philosophers: Code

```
semaphore C[5] = 1;
              left chop critical section
philosopher i
                              wait for my left chop
while (1) {
                                , right chop critical section
    // thinking
                               ... wait for my right chop
    C[i].wait();
       (i+1)%5].wait();
                               release my right chop
    C[(i+1)%5].signal()
                           release my left chop
    C[i].signal();
    // finishes eating
       Does this solution work?
```

Dining Philosophers: Deadlock!

- If all five philosophers sit down and pick up their left chopsticks at the same time, this causes a circular waiting and the program deadlocks.
- An easy way to remove this deadlock is to introduce a weirdo who picks up his right chopstick first!



Dining Philosophers: A Better Idea: 1/2


- Let the philosophers be 0, 1, 2, 3 and 4, and the weirdo is 4.
- The "normal" philosophers (i.e., 0-3) always pick up their left chopsticks followed by their right ones.
- The weirdo (i.e., 4) always picks up his right chopstick followed by his left.
- This solution is also referred to as the leftyrighty solution.

Dining Philosophers: A Better Idea: 2/2

```
semaphore C[5] = 1;
philosopher i (0, 1, 2, 3)
                          Philosopher 4: the weirdo
                          while (1) {
while (1) {
                            // thinking
   // thinking
                            C[(i+1) %5].wait();
  C[i].wait();
  C[(i+1)%5] wait();
                            C[i].wait();
                           /// eating
   // eating
  C[(i+1)%5].signal()
                            C[i].signal();
                            C[(i+1)%5].signal();
  C[i].signal();
   // finishes eating;
                           .// finishes eating
            lock left chop
                            lock right chop
```

Discussion: 1/2

- Suppose philosopher 4 is the weirdo.
- Suppose there is a deadlock.
- All the normal philosophers have their left and wait for their right.
- Because the weirdo picks up his right chopstick followed by his left, his right neighbor (i.e., Philosopher 1) cannot have his left.
- Or, if Philosopher 1 has his left, then the weirdo cannot have his right chopstick.
- This cannot happen, and the circular waiting pattern cannot happen!




Discussion: 2/2

- The following are some important questions:
 - ***** We choose philosopher 4 to be the weirdo. Does this choice matter?
 - ***** What if we have more than one weirdos?
 - *How many weirdos can be so that this solution is still deadlock-free?
 - *Can four or fewer philosophers cause a deadlock?
- *This solution is not **symmetric** because not all threads run the same code. Furthermore, you need to write two versions of the philosopher code. More codes usually mean more troubles.

Observation

- All philosophers are normal.
- If all four philosophers sit down; there will be an empty seat.
- The right-most chopstick is free, and the right-most philosopher has a chance to eat.
- What if the right-most philosopher cannot eat? **Exercise!**
- Circular waiting is broken.
- If only four philosophers are allowed to sit down, no deadlock can occur!



Count-Down Lock Example

```
semaphore C[5]= 1;
semaphore Chair = 4
          get a chair
                               this count-down lock
while
                               only allows 4 to go!
    // thinking
   Chair.wait()
       C[i].wait();
       C[(i+1)%5].wait();
       |// eating
                                 ◆····· this is our old friend
       C[(i+1)%5].signal();
       C[i].signal();
   Chair.signal();
                           ·release my chair
```

Exercises: 1/2

- We discussed the weirdo and 4-chair versions.
 - ➤ Use execution sequences to show that some philosophers may have no chance to eat indefinitely (i.e., starvation).
 - > Prove that the 4-chair version is deadlock-free.
- Suppose chopsticks are numbered from 0 to 4. A philosopher always picks up the low number one followed by the high number one. Compare this solution with the weirdo (or lefty-righty) one. What do you find?

Exercises: 2/2

- Suppose all chopsticks are in a tray. A philosopher picks up a chopstick as his left followed by another as his right. Is this deadlockfree? How about starvation?
- Return to the original version. A philosopher sits down and flips a coin. If the result is a head, he picks up his left chopstick first. Otherwise, he picks his right first. Is this version deadlock-free?
- More solutions will be presented in later chapters/units.

Classical Problems

- Producer/Consumer (Bounded Buffer)
- ***** Readers-Writers
- ***** Roller-Coaster

The Producer/Consumer Problem


- The Producer/Consumer problem is also referred to as the Bounded Buffer problem.
- It was discussed in E. W. Dijkstra's seminal paper, actually his lecture notes for his course.

4.3. The Bounded Buffer.

I shall give a last simple example to illustrate the use of the general semaphore. In section 4.1 we have studied a producer and a consumer coupled via a buffer with unbounded capacity. This is a typically one-sided restriction: the producer can be arbitrarily far ahead of the consumer, on the other hand the consumer can never be ahaed of the producer. The relation becomes symmetric, if the two are coupled via a buffer of finite size, say N portions. We give the program without any further discussion; we ask the reader to convince himself of the complete symmetry. ("The consumer produces and the producer consumes empty positions in the buffer".) The value N, as the buffer, is supposed to be defined in the surrounding universe into which the following program should be embedded.

E.W.Dijkstra (1968), Co-operating Sequential Processes, in Programming Languages: NATO Advanced Study Institute: Lectures given at a three weeks Summer School held in Villard-le-Lans 1996, edited by F. Genuys (pp. 43-112), Academic Press, Inc.


The Producer/Consumer Problem




- Suppose we have a circular buffer of n slots.
- Pointer in (resp., out) points to the first empty (resp., filled) slot.
- Producer threads keep adding data into the buffer.
- Consumer threads keep retrieving data from the buffer.

50

Problem Analysis



- A producer deposits data into Buf[in] and a consumer retrieves info from Buf[out].
- in and out must be advanced.
- in is shared among producers.
- out is shared among consumers.
- If Buf is full, producers should be blocked.
- If Buf is empty, consumers should be blocked.



- A semaphore to protect the buffer.
- A semaphore to block producers if the buffer is full.
- A semaphore to block consumers if the buffer is empty.
 - a wait or locka notification or unlock

Solution

```
number of slots
semaphore NotFull=n;
                       NotEmpty=0, Mutex=1;
producer
                                consumer
while (1) {
                          while (1) {
  NotFull.wait().
                            NotEmpty.wait();
    Mutex.wait()
                              Mutex.wait();
      Buf[in] = x;
                                 x = Buf[out];
      in = (in+1)%n;
                                 out = (out+1) %n;
    Mutex.signal();
                             : Mutex.signal();
                            NotFull signal();
  NotEmpty.signal ()
              notifications
                                    critical section
```

Question

- What if the producer code is modified as follows?
- Answer: a deadlock may occur. Why?

```
while (1) {
    Mutex.wait();
    NotFull.wait();

Buf[in] = x;
    in = (in+1)%n;

NotEmpty.signal();
Mutex.signal();
}
```

The Readers/Writers Problem


- Two groups of threads, readers and writers, access a shared resource by the following rules:
 - *Readers can read simultaneously.
 - **Only one** writer can write at any time.
 - *When a writer is writing, no reader can read.
 - **❖**If there is any reader reading, all incoming writers must wait. Thus, readers have a higher priority.

Problem Analysis

- A semaphore is needed to block the first reader and writers if a writer is writing.
- When a writer arrives, it must know if there are readers reading. A reader count is required and must be protected by a lock.
- This reader-priority version has a problem: if readers keep coming in an overlapping way, waiting writers have no chance to write.


writer is waiting		reader <i>n</i>
writer arrives	 	
reader 3	 	
reader 2	 	
reader 1	 	

56



Readers


- When a reader arrives, it adds 1 to the counter.
- If it is the first reader, waits until no writer is
- Decreases the counter.
- If it is the last reader, tells the waiting readers or writers that no reader is reading.



Readers

- When a reader arrives, it adds 1 to the counter.
- If it is the first reader, waits until no writer is writing.
- Reads data.
- Decreases the counter.
- If it is the last reader, tells the waiting readers or writers that no reader is reading.


if a writer waits here, all subsequent writers are blocked here



writers are blocked here.

Readers


- When a reader arrives, it adds 1 to the counter.
- If it is the first reader, waits until no writer is writing.
- Reads data.
- Decreases the counter.
- If it is the last reader, tells the waiting readers or writers that no reader is reading.



Readers

- When a reader arrives, it adds 1 to the counter.
- If it is the first reader, waits until no writer is writing.
- Reads data.
- Decreases the counter.
- If it is the last reader, tells the waiting readers or writers that no reader is reading.

if a writer waits here,
all subsequent writers
are blocked here



Writers

- When a writer comes in, it waits until no reader is reading, and no writer is writing.
- Then, it writes data.
- Finally, tells the waiting readers or writers that no writer is writing.

there is a writer writing, or the 1st reader gets through

Solution

```
semaphore Mutex = 1, WrtMutex = 1;
int RdrCount;
reader
                             writer
while (1) {
                             while (1) {
 Mutex.wait();
    RdrCount++;
    if (RdrCount == 1)
                         blocks both readers and writers
     WrtMutex.wait();
                               WrtMutex.wait();
  Mutex.signal();
  // read data
                               // write data
 Mutex.wait();
    RdrCount--;
    if (RdrCount == 0)
      WrtMutex.signal();
                               WrtMutex.signal();
  Mutex.signal();
                                                62
```


The Roller-Coaster Problem: 1/11

- Suppose there are n passengers and one roller coaster car. The passengers repeatedly wait to ride in the car, which can hold maximum C passengers, where C < n.
- The car can go around the track only when it is full. After finishes a ride, each passenger gets off the car, and wanders around the amusement park before returning to the roller coaster for another ride.
- Due to safety concerns, the car only rides T times and then shut-down.

The Roller-Coaster Problem: 2/11


- The car always rides with exactly C passengers
- No passengers will jump off the car while the car is running
- No passengers will jump on the car while the car is running
- No passengers will request another ride before they get off the car.

The Roller-Coaster Problem: 3/11




- A passenger decides to have a ride and joins the queue.
- The queue is managed by a gate keeper.
- Passengers check in one-by-one.
- The last passenger tells the car that all passengers are on board.
- Then, they have a ride.
- After riding, passengers get off the car one-by-one.
- Passengers play for a while and come back for a ride.

The Roller-Coaster Problem: 4/11




- A passenger decides to have a ride and joins the queue.
- The queue is managed by a gate keeper.
- Passengers check in one-by-one.
- The last passenger tells the car that all passengers are on board.
- Then, they have a ride.
- After riding, passengers get off the car one-by-one.
- Passengers play for a while and come back for a ride.

The Roller-Coaster Problem: 5/11




The Roller-Coaster Problem: 6/11




- A passenger decides to have a ride and joins the queue.
- The queue is managed by a gate keeper.
- Passengers check in one-by-one.
- The last passenger tells the car that all passengers are on board.
- Then, they have a ride.
- After riding, passengers get off the car one-by-one.
- Passengers play for a while and come back for a ride.

The Roller-Coaster Problem: 7/11



- The car comes and lets the gate keeper know it is available so that the gate keeper could release passengers to check in.
- The car is blocked for loading.
- When the last passenger is in the car, s/he informs the car that all passengers are on board, the car starts a ride.
- After this, the car waits until all passengers are off. Then, go for another round.

The Roller-Coaster Problem: 8/11




The car comes and lets the gate keeper know it is available so that the gate keeper could release passengers to check in.

The car is blocked for loading.


- When the last passenger is in the car, s/he informs the car that all passengers are on board, the car starts a ride.
 - After this, the car waits until all passengers are off. Then, go for another round.

The Roller-Coaster Problem: 9/11



- The car comes and lets the gate keeper know it is available so that the gate keeper could release passengers to check in.
- The car is blocked for loading.
- When the last passenger is in the car, s/he informs the car that all passengers are on board, the car starts a ride.
- After this, the car waits until all passengers are off. Then, go for another round.

The Roller-Coaster Problem: 10/11



- The car comes and lets the gate keeper know it is available so that the gate keeper could release passengers to check in.
- The car is blocked for loading.
- When the last passenger is in the car, s/he informs the car that all passengers are on board, the car starts a ride.
 - After this, the car waits until all passengers are off. Then, go for another round.

The Roller-Coaster Problem: 11/11

```
int count = 0;
Semaphore Queue = Boarding = Riding = Unloading = 0;
Semaphore Check-In = 1;
                                count is shared but not protected. why?
Passenger
                          Car
```

```
Wait (Queue) ;
Wait(Check-In);
 if (++count==Maximum)
    Signal (Boarding) ;
...Signal(Check-In);
 Wait (Riding); 🔨
 Signal (Unloading);
```

Exercise:

```
This code unloads passengers
  one-by-one. Is this necessary?!
Can Unloading be removed?
```

```
for (i = 0; i < #rides; i++) {
   count = 0; // reset counter before boarding
   for (j = 1; j \le Maximum; j++)
    Signal (Queue) ; // car available
   Wait(Boarding);
   // all passengers in car
   // and riding
   for (j = 1; j \le Maximum; j++) {
     Signal(Riding);
      Wait(Unloading);
   // all passengers are off
```

one ride

A Quick Summary: 1/2

- We have learned a few tricks in this component: locks, count-down locks and notification.
- Very often a counter is needed to determine if certain condition is met (e.g., number of readers in the readers-writers problem, check-in and boarding in the roller-coaster problem).
- Sometimes threads may have to be "paired-up" like the get-off process we saw in the rollercoaster problem.
- Use these basic and frequently seen patterns to solve other problems.

A Quick Summary: 2/2

- Using many semaphores could mean more locking and unlocking activities and could be inefficient.
- Using only a few semaphores could produce very large critical sections, and a thread could stay in a critical section for a long time. Thus, other threads may have to wait very long to get in.
- Therefore, try your best to minimize the number of semaphores and reduce the length of locking time.

Patterns and Pass-the-Baton

What Is a Pattern?

- ☐ A pattern is simply a description/template for solving a problem that can be used in several situations.
- ☐ A pattern is **NOT** a complete solution to a problem. It is just a template and requires extra work to make it a solution to a specific problem.
- ☐ We will discuss a few patterns related to the use of semaphores.

Mutual Exclusion – Of Course!

- ☐ This is the easiest one for enforcing mutual exclusion so that race conditions will not occur.
- A semaphore is initialized to 1. Then, use the Wait() and Signal() methods to lock and unlock the semaphore, respectively.

```
Semaphore Lock(1);
Wait(Lock);
    // critical section
Signal(Lock);
```

```
Semaphore S(1);
int c = 0;

Wait(S); Wait(S);
c++;
Signal(S); Signal(S);
if (c >= 3) {\ if (c == 0) {\}
...
```

78

Enter-and-Test: 1/2

- □ In many applications, a thread may enter a critical section and test for a condition. If that condition is met, the thread does something₁. Otherwise, its does something₂.
- ☐ Frequently, one of the two **something**s may involve a wait.

```
Reader: Enter
Mutex.wait();
RdrCount++;
if (RdrCount == 1)
WrtMutex.wait();
Mutex.signal();
// read data
```

(i.e., RdrCount being 1),
then waits until it is released.
In this case, the first reader does
something (i.e., waiting) and
at the same time has the Mutex.
In this way, no other threads
can enter the critical section.

Enter-and-Test: 2/2

☐ Usually, a wait may be used in the entry part to wait for a particular condition to occur, and a signal is used upon exit to release a waiting thread.

```
Reader: Exit
// read data

Mutex.wait();
RdrCount--;
if (RdrCount == 0)
WrtMutex.signal();
Mutex.signal();
}
--- critical section
```

if the condition is met

(i.e., RdrCount being 0),

then tell someone, a reader

or a writer, to continue.

In this case, the last reader does
something.

Exit-Before-Wait: 1/2

- ☐ In many applications, a thread exits a critical section and then blocks itself.
- ☐ Usually, a thread updates some variables in a critical section, and then waits for a resource from another thread.

```
Roller-Coaster: Passenger
Wait(Queue);
Wait(Check-In);
if (++count==Maximum)
    Signal(Boarding);
Signal(Check-In);
Wait(Riding);
Signal(Unloading);
```

if the condition is met (i.e., count being the maximum), then notify some thread.

after exiting the critical section, wait for some event to happen.

Exit-Before-Wait: 2/2

- ☐ This signaling an event followed by waiting on another must be used with care.
- ☐ A context switch can happen between the signal and the wait.
- □ For example, a thread enters the critical section, signals s1 upon exit, and gets swapped out before reaches the wait. This could cause some problems. Why? So, be careful!

```
Wait(s1);
   // critical section
Signal(s1);
Wait(s2);
```

This thread may not wait immediately Is this OK? It all depends on the logic of your program.

Conditional Waiting/Signaling

- ☐ A thread waits or notifies another thread if a condition is satisfied.
- Make sure that no race conditions will occur while the condition is being tested.

Passing the Baton: 1/8

- ☐ If a thread is in its critical section, it holds the baton (i.e., the critical section).
- □ Upon exit, if there are threads waiting to enter the CS, the exiting thread passes the **baton** (*i.e.*, the critical section) to one of them directly.
- ☐ In this way, we save a signal-wait pair.
- ☐ If no thread is waiting, the **baton** is passed to the next thread that will try to enter the CS later.
- ☐ This is a technique that can make the use of semaphores more efficient.

Passing the Baton: 2/8

☐ The **Waiting** thread waits on Condition if **Event** is not there. The **Signaling** thread sets **Event** and releases a **Waiting** thread.

```
Semaphore Mutex(1);
        Semaphore Condition(0);
        Bool Event = FALSE;
        Waiting Thread
                                  Signaling Thread
        Wait(Mutex);
                                 Wait(Mutex);
        while (!Event) {
                                    Event = TRUE;
          Signal (Mutex) ;
                                    Signal (Condition)
          Wait (Condition)
        --- Wait (Mutex);
                                  Signal (Mutex) ;
try again!
```

Passing the Baton: 3/8

- Waiting does not acquire the CS. Instead, Signaling has the CS, does not release it, and gives the CS to Waiting (i.e., baton passed)
- **□ Signaling** must be sure that **Waiting** will not do any harm to the CS.

```
Waiting Thread
Wait(Mutex);
while (!Event) {-
    Signal (Mutex);
    Signal (Mutex);
    Wait(Condition);
    Wait(Mutex);
    Signal (Condition);
    Signal (Condition);
    Signal (Mutex);
}
```

Passing the Baton: 4/8

Threads waits for an event to occur before they can continue.

A **Waiting** thread checks to see if the event is there.

If the event is not there, it waits.

Variable Waiting is used to count the number of waiting thread.

Semaphore Condition is used to block threads if the event is not there.

```
Semaphore Mutex(1), Condition(0);
int
         Event = FALSE, Waiting = 0;
Waiting Thread
Wait(Mutex);
                    // lock Event/Waiting
                   // if Event not there
if (!Event) {
                  //
 Waiting++;
                          join the waiting
  Signal (Mutex);  // release the lock
 Wait(Condition);  // wait!
    // has the event & does something
    // upon exit, releases a waiting thread
if (Waiting > 1) { /// anyone waiting?
                  /// try to release
 Waiting--;
  Signal(Condition); // release a thread
Signal (Mutex) ;
                    // release the lock
```

Passing the Baton: 5/8

- A Signaling thread locks the semaphore Mutex and hence Waiting.
- Then, this Signaling thread makes the event to happen.
- It checks if there are threads waiting for this event.
- If there are waiting threads, release one of them.
- Finally, this Signaling thread releases the lock.

Passing the Baton: 6/8

```
Semaphore Mutex(1), Condition(0);
          Event = FALSE, Waiting = 0;
int
Waiting Thread
                       Signaling Thread
Wait(Mutex);
                       Wait(Mutex); →……
if (!Event) {
                       Event = TRUE;
  Waiting++;
  Signal (Mutex) ; ··
  Wait(Condition);
/// Process Event
                        if (Waiting > 0) {
if (Waiting > 1) ! {
  Waiting--;
                         Signal (Condition);
  Signal (Condition);
Signal (Mutex) ;
                      Signal (Mutex); ...
```

Only needed if Waiting = 0
The signaling thread has the CS and can pass it to the next

Passing the Baton: 7/8

```
Semaphore Mutex(1), Condition(0);
                   Event = FALSE, Waiting = 0;
        int
        Waiting Thread
                                Signaling Thread
        Wait(Mutex);
                                -Wait(Mutex);
        if (!Event) {
                                Event = TRUE;
          Waiting++; 👡
                               baton acquired
           Signal (Mutex)
                                a Mutex needed to protect Waiting
         Wait(Condition);
                                 if (Waiting > 0)
        if (Waiting >
          Waiting--;
           Signal (Condition);
                                 Signal (Condition);
            - baton passed
                                 else
        else
                                   Signal (Mutex) ; |
          Signal (Mutex);
baton released
                                                         90
```

Passing the Baton: 8/8

- **□ Passing the baton** technically transfers the ownership of a critical section from a thread to another thread.
- ☐ The thread that has the baton does not need a signal to release it. Instead, the CS is directly given to another that needs it. The receiving thread does not need a wait for the CS.
- ☐ In this way, mutual exclusion may be destroyed; but, we reduce the number of entering and exiting a mutex.

Pass-the-Baton Example

Passing the Baton: Example 1/23

- ☐ We shall use the reader-priority version of the readers/writers problem as a more complex example.
- **■** Note the following conditions:
 - **❖**If there is no writer writing, a reader can read.
 - **❖**If there is no readers reading and there are waiting writers, allow a writer to write (i.e., better!).
 - **❖**If there are readers reading **○R** a writer writing, no writer can write.
 - **❖**If there are waiting readers, a finishing writer should allow a reader to read (i.e., reader priority).
 - **❖**If there are waiting writers and no waiting reader, a finishing writer should allow a writer to write.

Passing the Baton: Example^{2/23}

- **■** We will need counters for counting waiting readers and writers and active readers and writer.
- ☐ A semaphore for protecting all counters.
- ☐ A semaphore for blocking readers.
- ☐ A semaphore for blocking writers.

Passing the Baton: Example 3/23

Developing the Entry Section for Readers: 1/6

```
Mutex is used to protect all counters
Wait(Mutex);
if (aWriters > 0)
   wReaders++;
   Signal (Mutex);
   Wait(Reader);
   Wait (Mutex)
aReaders++;
if (wReaders > 0)
   wReaders--;
   Signal (Reader) ;
Signal (Mutex);
```

Passing the Baton: Example 4/23

Developing the Entry Section for **Readers**: 2/6

```
Wait(Mutex);
if (aWriters > 0)
   wReaders++;
   Signal (Mutex);
   Wait(Reader);
   Wait (Mutex)
aReaders++;
if (wReaders > 0)
   wReaders--;
   Signal (Reader);
Signal (Mutex) ;
```

If there is a writer writing,
this reader must wait.
As a result, this reader must release
the Mutex and wait on Reader.

If a reader is released from Reader,
-this reader must reacquire the
Mutex in order to update
aReaders and wReaders,
and test wReaders.

Passing the Baton: Example 5/23

Developing the Entry Section for Readers: 3/6

```
Wait(Mutex);
if (aWriters > 0)
   wReaders++;
   Signal (Mutex);
   Wait(Reader)
   Wait (Mutex)
aReaders++;
if (wReaders
   wReaders--;
   Signal (Reader)
Signal (Mutex)
```

If there is a writer writing,
this reader must wait.
As a result, this reader must release
the Mutex and wait on Reader.

If a reader is released from Reader, this reader must reacquire the Mutex in order to update aReaders and test wReaders.

If a reader reaches here, then
--(1) it is released from Reader,
or
-- (2) there is no writer writing.
Then, this reader can read.

Passing the Baton: Example 6/23

Developing the Entry Section for **Readers: 4/6**

```
Wait(Mutex);
if (aWriters > 0)
   wReaders++;
   Signal (Mutex)
   Wait(Reader)
   Wait(Mutex);
aReaders++; 🚄
   (wReaders > 0) --{*
if
   wReaders--;
   Signal (Reader);
Signal (Mutex) ;
```

```
If a reader reaches here, then
(1) it is released from Reader,
or
(2) there is no writer writing.
Then, this reader can read.
```

```
This reader adds 1 to the active reader count.

If there are waiting readers, release one.

This released reader releases other waiting readers in a cascading way.
```

Passing the Baton: Example 7/23

Developing the Entry Section for Readers: 5/6

```
Wait(Mutex);
if (aWriters > 0)
   wReaders++;
   Signal (Mutex) ;
   Wait(Reader);
   Wait (Mutex)
aReaders+4
  (wRéaders > 0) {
  wReaders--;
   Signal (Reader);
Signal (Mutex) ;
```

```
This first reader that sees no writing
  writer can read immediately.
This reader starts to signal Reader
  to release other waiting readers.
The first reader executes Signal ()
  to Reader.
If a reader is released, this reader
  can read because readers can
  read simultaneously.
This reader signals Reader to
  release the next reader waiting
  on Reader.
Therefore, a sequence of signals
  releases the waiting readers on
  Reader one-by-one.
This is a cascading signal/release.
```

Passing the Baton: Example 8/23

Developing the Entry Section for **Readers: 6/6 (Summary)**

```
Wait(Mutex); ~ acquire
if (aWriters > 0) \{
    wReaders++;
    Signal (Mutex) ; release
    Wait(Reader);
    Wait(Mutex); <
            re-acquire
aReaders++;
    (wReaders > 0
    wReaders--;
   ·Signal (Reader)
Signal (Mutex) ; 🔺
```

The first reader, who sees
no writer writing and
some readers waiting,
releases the waiting readers.

The release of readers is in a cascading way (i.e., one after the other).

The first thread acquires the baton, and may pass it to each of the released thread.

Passing the Baton: Example 9/23

Developing the Entry Section for Readers: Passing the Baton

```
Wait(Mutex);
if (aWriters > 0) {
   wReaders++;
   Signal (Mutex) ;
   Wait (Reader);
  -Wait (Mutex) ;
aReaders++;
if (wReaders > 0). {
   wReaders--;
   Signal (Reader) ;
else
   Signal (Mutex) ;
```

The first reader, who gets the baton and sees no writer writing and -- some readers waiting, will release the waiting readers and pass its baton to the released reader.

Each released reader, after releasing the next reader, passes the baton -- to it. Then, it goes away without releasing the baton.

Only the last released returns the baton.

Note that the critical section (i.e., the baton) is locked throughout this cascading release, and no other thread can pass through the first Wait (Mutex).

Passing the Baton: Example 10/23

Developing the Exit Section for Readers

```
Wait(Mutex); acquire
aReaders--;
if (aReaders=0 & wWriters>0) {
    wWriters--;
    Signal(Mutex); release
    Signal(Writer);
}
else
    Signal(Mutex);
```

If this is the <u>last reader</u> and there are <u>waiting writers</u>, let one writer go.

Otherwise, simply go away.

This signal is not needed, because the last reader can pass the baton to the released writer!

Passing the Baton: Example 11/23

Developing the Exit Section for **Readers**: Passing the Baton

```
Wait(Mutex);
aReaders--;
if (aReaders=0 & wWriters>0) {
   wWriters--;
   -Signal(Mutex);
}
else
Signal(Mutex);
```

The baton is passed to the released writer.

If this is the last reader and there are waiting writers, let one writer go.

Otherwise, simply go away.

This release of the baton is not needed, because the last reader passes the baton to the released writer.

Does the following work?

```
Wait(Mutex);
aReaders--;
if (aReaders=0 & wWriters>0) {
   wWriters--;
   Signal(Writer);
}
Signal(Mutex);
```

Would this version work properly in terms of passing the baton? Why? Exercise!

Passing the Baton: Example 12/23

Developing the Entry Section for Writers

```
Wait(Mutex); -acquire
if (aReaders>0 | aWriters>0) {
    wWriters++;
    Signal(Mutex); release
    Wait(Writer);
    Wait(Mutex); re-acquire
}
aWriters++;
Signal(Mutex); re-release
```

A writer, who gets the baton and sees some readers reading OR a writer writing, will release the baton and wait.

Later, this writer reacquires
the baton to update
aWriters.

Otherwise (i.e., no readers reading AND no writer writing), this writer adds 1 to the number of active writers, releases the baton, and starts writing.

Passing the Baton: Example 13/23

Developing the Entry Section for Writers: Passing the Baton

```
Wait(Mutex); --acquire
if (aReaders>0 | aWriters>0) {
    wWriters++; baton
    Signal(Mutex); release
-- Wait(Writer); receives the baton
    -Wait(Mutex); from a Reader
}
aWriters++; release the received baton
Signal(Mutex);
```

The exiting reader acquires the baton, releases a writer w/o releasing the baton.

Hence, the released writer has the baton, increases the active writers count, releases the baton and writes.

baton

```
Exit Section of Readers: Passing the Baton
Wait (Mutex);
aReaders--;
if (aReaders=0 & wWriters>0) {
    wWriters--;
    Signal (Writer);
}
else
    Signal (Mutex);
```

Passing the Baton: Example 14/23

Developing the Exit Section for Writers

```
Wait (Mutex); <a>acquire</a>
aWriters--;
if (wReaders > 0)
   wReaders--;
   Signal (Mutex)
   Signal (Readers);
else if (wWriters >_
   wWriters--;
   Signal (Mutex);
   Signal (Writer)
   Signal (Mutex)
```

Lock the Mutex first.

If there are waiting readers, let one of them go.

If there is no waiting readers but there are waiting writers, let one of them go.

If there is no waiting readers and no waiting writers, then do nothing.

But, don't forget to release the Mutex.

Passing the Baton: Example 15/23

Developing the Exit Section for Writers: Passing the Baton

```
aWriters--;
if (wReaders > 0) {
   wReaders--;
  -Signal (Mutex)-
   Signal (Readers);
else if (wWriters > 0)
   wWriters--;
  -Signal (Mutex)-;-
   Signal (Writer)
else
   Signal (Mutex) ;
```

This writer acquires the baton.

This writer may pass the baton to a released reader or a released writer.

The baton is not passed if there is no waiting readers and there is no waiting writers.

In this case, this writer must — release the baton explicitly.

```
Wait(Mutex);
Wait(Mutex);
                       Passing the Baton
if (aWriters > 0) {
                                           if (aReaders > 0 | aWriters > 0) {
   wReaders++;
                                             |wWriters++;|
acquire
Signal (Mutex);
                                              Signal (Mutex);
                                pass
   Wait(Reader); ←
                                              Wait (Writer);
                 loops here to
                                          aWriters++;
aReaders++;
                 release all readers
if (wReaders > 0)
                                                  release by the
   wReaders -- ;/
                                                 finishing writer
   Signal(Reader);
                                    pass
                release by the
else
                last released reader
   Signal(Mutex);
                                          Signal(Mutex);
// READING
                                          // WRITING !
Wait(Mutex);
                                          Wait(Mutex);
aReaders--;
                                          aWriters--;
if (aReaders = 0.6 wWriters > 0)
                                           if (wReaders > 0) {
   wWriters--;
                                              wReaders--;
Signal (Writer) ;
                                              Signal (Reader) ;
                                          else if (wWriters > 0) {
                                              wWriters--;
                                              Signal(Writer);
                                          else // wReaders = wWriters = 0
else
   Signal (Mutex);
                                              Signal (Mutex);
```

Summary

```
Wait(Mutex);
                                          Wait(Mutex);
                       Passing the Baton
if (aWriters > 0) {
                                           if (aReaders > 0 | aWriters > 0) {
   wReaders++;
                                             |wWriters++;|
   Signal (Mutex);
                                              Signal (Mutex) ;
                                  pass
   Wait(Reader);
                                              Wait(Writer); __
                 loops here to
aReaders++;
                                           aWriters++;
                 <u>\release all readers</u>
if (wReaders > 0)
                                                  release by the
   wReaders -- ;/
                                                  finishing writer
   Signal (Reader);
                                                                          pass
                 release by the
else
               last released reader
   Signal(Mutex);
                                           Signal(Mutex);
                                          // WRITING
// READING
                                          Wait (Mutex); acquire
Wait(Mutex);
                                   acquire | aWriters--;
aReaders--;
if (aReaders = 0 & wWriters > 0)
                                          if (wReaders > 0)
   wWriters--;
                                              wReaders--;
   Signal(Writer);
                                              Signal (Reader);
                                           else if (wWriters > 10)
                                             wWriters--;
                                              Signal (Writer);
                                          writers // wReaders = wWriters = 0
else
   Signal (Mutex);
                                              Signal (Mutex);
```

Summary

Passing the Baton: Example 18/23

- One of the several advantages of this solution is that it can easily be modified to achieve other goals. Here is a writer-priority version.
- A writer-priority version should satisfy the following conditions:
 - 1. New readers are blocked if a writer is waiting, and
 - 2. A waiting reader is released if no writer is writing.

Study the conditions for releasing readers and writers.

Passing the Baton: Example 19/23

☐ To meet Condition (1), we must change the first (enter) if statement of the reader thread:

```
if (aWriters > 0 || wWriters > 0) { // yield to writers
   wReaders++;
   Signal(Mutex);
   Wait(Reader);
   // as long as there
   // is a writer waiting
   // yield to writers
   // joint the line,
   // release the baton,
   // and wait
}
```

When a reader arrives, if there are writers waiting or writing, this reader blocks, and hence yields to writers.

Passing the Baton: Example^{20/23}

☐ To meet Condition (2), we must strengthen the last (exit) if statement of the writer thread:

```
if (wReaders > 0 && wWriters == 0)
                                     // release readers
                     added component
                                         only if there
   wReaders--;
  Signal (Reader) ;
                                      // is no writers
else if (wWriters > 0) {
                                      // if no readers
                                      // but some writers
  wWriters--;
                                      // release one
  Signal(Writer);
else
                                      // no readers/writers
 Signal (Mutex);
```

'Release a reader if there is no writer waiting.

Otherwise (i.e., no waiting readers or some waiting writers), release a writer.

112

Finally (i.e., no waiting readers and no waiting writers), do nothing.

A waiting reader is released if no writer is writing.

Passing the Baton: Example^{21/23}

- ☐ A fair version should allow readers and writers take turns (i.e., no starvation).
- **■** We assume the semaphores being used is implemented so that every blocked thread will be eventually released (i.e., no starvation). This is difficult to achieve.
- A fair version must satisfy the following:
 - 1. When a writer finishes, all waiting readers get a turn
 - 2. When all current readers finish reading, one waiting writer can write.

Passing the Baton: Example^{22/23}

- **■** We need to change the last (exit) if of the writer.
- A new Bool variable Writing is need to indicate whether a writer is writing. Writing is set to TRUE when a writer starts writing and is set to FALSE when a reader starts reading.

```
if (wReaders > 0 & (wWriters = 0 | Writing)) {
                      NEW: no waiting or writing writer \( \)
    wReaders--;
                                                      This is set by the current
                                                       exiting writer.
    Signal (Reader); a reader can read
                                                      Because it exits, no writer
                                                        is writing.
else if (wWriters > 0 & (wReaders = 0 | ¬Writing);)
                             NEW: no waiting readers or no writing writer
   wWriters--;
    Signal (Writer); a writer can writer
                                            Why is the starvation
                                         free assumption needed?
else // wReaders = wWriters
    Signal (Mutex);
                                                                  114
```

Passing the Baton: Example^{23/23}

- ☐ This example, including its extensions and the passing the baton pattern, is due to Gregory R. Andrews. Refer to the following for more detailed discussions.
 - 1. Gregory R. Andrews, *Concurrent Programming: Principles and Practice*, Benjamin/Cummings, 1991.
 - 2. Gregory R. Andrews, A Method for Solving Synchronization Problems, *Science of Computer Programming*, Vol. 13 (1989/1991), pp. 1-21.

Semaphores with ThreadMentor

Semaphores with ThreadMentor

- ThreadMentor has a class Semaphore with two methods Wait() and Signal().
- Class Semaphore requires a nonnegative integer as an initial value.
- A name is optional.

```
Semaphore Sem("S",1);
Sem.Wait();
// critical section
Sem.Signal();
Semaphore *Sem;
Sem = new
  Semaphore("S",1);
Sem->Wait();
// critical section
Sem->Signal();
```

Dining Philosophers: 4 Chairs

```
Semaphore Chairs (4);
Mutex Chops[5]={1,..,1};
class phil::public Thread
  public:
    phil(int n, int it);
  private:
    int
         Number;
    int iter;
    void ThreadFunc()
};
Count-Down and Lock
```

```
Void phil::ThreadFunc()
 int i, Left=Number,
     Right=(Number+1)%5;
 Thread::ThreadFunc();
 for (i=0; i<iter; i++) {
  Chairs.Wait();
     Chops[Left].Lock();
     Chops[Right].Lock();
     // Eat
     Chops[Left].Unlock();
     Chops[Right].Unlock();
   Chairs.Signal();
```


The Smokers Problem: 1/6

- Three ingredients are needed to make a cigarette: tobacco, paper and matches.
- An agent has an infinite supply of all three.
- Each of the three smokers has an infinite supply of one ingredient only. That is, one of them has tobacco, the second has paper, and the third has matches.
- They share a table.

The Smokers Problem: 2/6

- The agent adds two randomly selected different ingredients on the table and notifies the needed smoker.
- A smoker waits until agent's notification. Then, takes the two needed ingredients, makes a cigarette, and smokes for a while.
- This process continues forever.
- How can we use semaphores to solve this problem?

The Smokers Problem: 3/6



The Smokers Problem: 4/6

- Semaphore Table protects the table.
- Three semaphores Sem[3] are used, one for each smoker:

Smoker #	Has	Needs	Sem
0	0	1 & 2	Sem[0]
1	1	2 & 0	Sem[1]
2	2	0 & 1	Sem[2]

The Smokers Problem: 5/6

```
class A::public Thread
               agent thread
  private:
    void ThreadFunc();
};
       smoker thread
class Smk::public Thread
  public:
    Smk(int n);
  private:
    void ThreadFunc();
    int No;
};
   clear the table
```

```
Smk::Smk(int n)
  No = n;
            waiting for ingredients
Void Smk::ThreadFunc()
  Thread::ThreadFunc();
  while (1) {
    Sem[No]->Wait();
    Table.Signal();
    // smoker a while
```

The Smokers Problem: 6/6

```
void A::ThreadFunc()
  Thread::ThreadFunc();
  int Ran;
                ingredients are ready
  while (1) {
    Ran = // random #
           // in [0,2]
    Sem[Ran]->Signal();
    Table.Wait();
waiting for the table
to be cleared
```

```
void main()
  Smk *Smoker[3];
      Agent;
  Agent.Begin();
  for (i=0;i<3;i++) {
   Smoker = new Smk(i);
   Smoker->Begin();
  Agent.Join();
```

The General Smokers Problem: 1/13

The original version of the Smokers problem was due to Suhas Patil in 1971. The agent and smokers have codes like the following:

```
Semaphore Table = 1, Sem[3] = {0,0,0};

Agent

while (1) {
    generate ingredients i and j
    Table.Wait();
    Sem[i].Signal();
    Sem[j].Signal();
    Sem[j].Signal();
}

    Semaphore Table = 1, Sem[3] = {0,0,0};
    while (1);
    Sem[m].Wait();
    Sem[n].Wait();
    Table.Signal();
    smoke for a while
}
```

The General Smokers Problem: 2/13

- However, deadlock can happen.
- Use an execution sequence to reveal a possible deadlock.

```
Semaphore Table = 1, Sem[3] = {0,0,0};

Agent

while (1) {
    generate ingredients i and j
    Table.Wait();
    Sem[i].Signal();
    Sem[j].Signal();
    Sem[j].Signal();
}

    Semaphore Table = 1, Sem[3] = {0,0,0};

    while (1) {
        Sem[m].Wait();
        Sem[n].Wait();
        Table.Signal();
        smoke for a while
}
```

The General Smokers Problem: 3/13

- Solving this general smokers problem with the semaphores discussed here is tedious, but doable.
- David L. Parnas published a deadlock-free solution and Nico Habermann proved that Parnas' solution was correct.
- It is interesting to point out that E. W. Dijkstra and Suhas Patil proposed the parallel Wait() approach in 1971.

David L. Parnas, On a Solution to the Cigarette Smoker's Problem (without conditional statements), *Communications of the ACM*, Vol. 18 (1975), No. 3, pp. 181-183.

Nico Habermann, On a Solution and a Generalization of the Cigarette Smokers' Problem, Computer Science Department, Carnegie-Mellon University, August 1972.

The General Smokers Problem: 4/13

- Dijkstra proposed to extend the Wait() and Signal() calls to use multiple semaphores.
- A process calls Wait_p (S₁, S₂,..., S_n) does not wait if and only if every semaphore S₁ has a positive counter value. Each semaphore counter value is decreased by 1 as usual.
- Signal_p (S₁, S₂,..., S_n) is simple: adding 1 to the counter of each semaphore S₁. Of course, some waiting process may be released due to the change of the values of semaphore counters.
- Dijkstra calls this type of Wait_p() parallel Wait().
 128

The General Smokers Problem: 5/13

- A general form of semaphores, called *semaphore* array, was proposed by Tilak Agerwala in 1977.
- Let $S_1, S_2, ..., S_n$ and $\underline{T}_1, \underline{T}_2, ..., \underline{T}_m$ be semaphores.
- Wait_e() and Signal_e() are defined as follows:

```
 \begin{aligned} & \text{Wait}_{e}(S_{1}, S_{2}, ..., S_{n}, \underline{T}_{1}, \underline{T}_{2}, ..., \underline{T}_{m}) \\ & \text{if (for all i in } [1, n] \colon S_{i} > 0 \text{ and for all j in } [1, m] \colon \underline{T}_{j} = 0) \\ & \text{for all i in } [1, n] \colon S_{i} - -; \\ & \text{else} \\ & \text{the caller process/thread is blocked} \\ & \text{Signal}_{e}(S_{1}, S_{2}, ..., S_{n}) \\ & \text{for all i in } [1, n] \colon S_{i} + +; \\ & \text{} \end{aligned}
```

The General Smokers Problem: 6/13

- Counters of Semaphores $S_1, S_2, ..., S_n$ are tested and decreased at the same time.
- Semaphores \underline{T}_{j} 's are semaphores T_{j} 's used for 0 testing.

The General Smokers Problem: 7/13

If the semaphores $\underline{T}_1, \underline{T}_2, ..., \underline{T}_m$ are not used, Agerwala's extension becomes Dijkstra's parallel Wait().

```
Wait<sub>e</sub>(S<sub>1</sub>,S<sub>2</sub>,...,S<sub>n</sub>)
{
    if (for all i in [1,n]: S<sub>i</sub> > 0)
        for all i in [1,n]: S<sub>i</sub>--;
    else
        the caller process/thread is blocked
}

Signal<sub>e</sub>(S<sub>1</sub>,S<sub>2</sub>,...,S<sub>n</sub>)
{
    for all i in [1,n]: S<sub>i</sub>++;
}
```

The General Smokers Problem: 8/13

- If semaphores S_i are not used, this extension becomes waiting for all the counters of $\underline{T}_1, \underline{T}_2, ..., \underline{T}_m$ to become zero.
- Any semaphore S can be use for 0 testing and is denoted as S.

```
\label{eq:wait} \begin{aligned} \text{Wait}_{\text{e}}\left(\underline{\textbf{T}}_{1},\underline{\textbf{T}}_{2},...,\underline{\textbf{T}}_{\text{m}}\right) \\ \text{if (for all j in } [1,m]: \ \underline{\textbf{T}}_{\text{j}} = 0) \\ \text{do nothing} \\ \text{else} \\ \text{the caller process/thread is blocked} \\ \end{aligned}
```

The General Smokers Problem: 9/13

- The philosophers problem can be solved very easily.
- Show that this solution is deadlock-free.

```
Philosophger i

Semaphore C[5] = { 1, 1, 1, 1, 1 };

while (1) {
    // thinking
    Wait<sub>e</sub>(C[(i+4)%5], C[(i+1)%5]);
    // eating
    Signal<sub>e</sub>(C[(i+4)%5], C[(i+1)%5]);
}
```

The General Smokers Problem: 10/13

- The general smokers problem can also be solved very easily.
- Show that this solution is deadlock-free.

```
Agent
Semaphore Table = 1;
Semaphore Ingred[3] = { 0, 0, 0 };
while (1) {
    Wait(Table);
    // generate two ingredients i and j
    Signal<sub>e</sub>(Ingred[i], Ingred[j]);
}
```

```
Smoker k who needs ingredients m and n

while (1) {
    Wait<sub>e</sub>(Ingred[m], Ingred[n]);
    Signal<sub>e</sub>(Table);
    // smoke for a while
}
```

The General Smokers Problem: 11/13

- **How do we use the** T **semaphores?** Here is an example.
- Process P_i has higher priority than process P_{i+1} (0 $\leq i \leq n-2$). There are n processes.
- The processes request access to the resource and are allocated in a mutually exclusive way based in the priorities.
- A request by a process is not honored until all higher priority requests are taken care of.

The General Smokers Problem: 12/13

- Process P_i sets its semaphore to 1, making a request.
- P_i waits for higher priority processes' semaphores $\underline{S}_1, \underline{S}_2, ..., \underline{S}_{i-1}$ to become 0, withdraws its request, and accesses the resource via semaphore R.

The General Smokers Problem: 13/13

- Unix/Linux supports semaphore arrays like Agerwala's semaphore extension.
- Use semget () to obtain a semaphore set using a key, which is similar to shared memory.
- Use semop () to operate on a semaphore set.
- For each semaphore element in a semaphore set, the following are permitted:
 - 1. There is no operation
 - 2. A value can be added or subtracted from the semaphore counter.
 - 3. A zero means waiting for 0.
- Refer to Unix/Linux manual for the details.

Food for Thought: 1/2

 Dijkstra suggested the following solution to the philosophers problem.

Global Items

Food for Thought: 2/2

 Dijkstra suggested the following solution to the philosophers problem.

Philosopher i

```
| Wait(Mutex);
                         // lock the Mutex to change state
                         // i is hungry
   State[i] = HUNGRY;
   CanEat(i);
                         // Can i eat? CanEat() is called in a Mutex
Signal (Mutex);
                         // if i can eat, it was signaled in CanEat()
Wait(toEat[i]);
   State[i] = EATING;
                               state changed to EATING
   // eat
Wait(Mutex);
                         // after eating, change state again
   State[i] = THINKING; // i is thinking
                         // allow left neighbor to eat
   CanEat((i+4)%5);
                         // allow right neighbor to eat
   CanEat((i+1)%5);
Signal (Mutex);
```

The End