Part III Synchronization Monitors

You cannot build (or understand) a modern operating system unless you know the principles of concurrent programming.

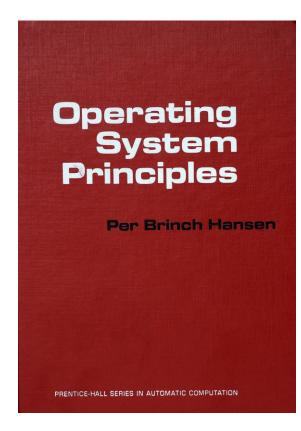
What Will Be Covered?

- Some historical remarks
- Monitor Basics
- What is a condition variable?
- Condition variable wait and signal
- Two Types of monitors: Hoare and Mesa
- Examples
- Hoare type vs. Mesa type and Semaphores vs. Monitors
- ThreadMentor Monitor Programming
- ThreadMentor Monitor Visualization

Some Historical Remarks: 1/2

- The concept of a monitor was invented by Per Brinch Hansen in early 1970s.
- Per Brinch Hansen used the concept of class in Simula 67 and defined a shared class as the beginning of today's monitor.
- C. A. R. Hoare refined Hansen's work to become today's monitor in 1974.
- Hansen is also considered as a pioneer of concurrent programming. His *Concurrent Pascal* language used monitors for synchronization.

Some Historical Remarks: 2/2



Per Brinch Hansen The Architecture Concurrent **Programs** Prentice-Hall Series in Automatic Computation

Hansen's landmark book Operating System Principles (Section 7.2) in 1973 Hansen's another landmark book The Architecture of Concurrent Programs (1977) in which the concept of monitor and his Concurrent Pascal were clearly defined and discussed. Operating Systems

C. Weissman Editor

Monitors: An

Operating System

Structuring Concept

C.A.R. Hoare The Queen's University of Belfast

This paper develops Brinch-Hansen's concept of a monitor as a method of structuring an operating system. It introduces a form of synchronization, describes a possible method of implementation in terms of semaphores and gives a suitable proof rule. Illustrative examples include a single resource scheduler, a bounded buffer, an alarm clock, a buffer pool, a disk head optimizer, and a version of the problem of readers and writers.

Key Words and Phrases: monitors, operating systems, scheduling, mutual exclusion, synchronization, system implementation languages, structured multiprogramming CR Categories: 4.31, 4.22 I. Introduction

A primary aim of an operating system is to share a computer installation among many programs making unpredictable demands upon its resources. A primary task of its designer is therefore to construct resource allocation (or scheduling) algorithms for resources of various kinds (main store, drum store, magnetic tape handlers, consoles, etc.). In order to simplify his task, he should try to construct separate schedulers for each class of resource. Each scheduler will consist of a certain amount of local administrative data, together with some procedures and functions which are called by programs wishing to acquire and release resources. Such a collection of associated data and procedures is known as a monitor; and a suitable notation can be based on the class notation of SIMULAST (bl.

tonitorname: monitor begin . . . declarations of data local to the monitor;

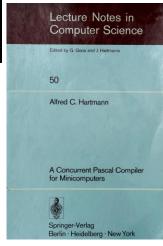
procedure procname (... formal parameters ...); begin ... procedure body ... end;

. . declarations of other procedure local to the monitor; . . initialization of local data of the monitor . . . end;

Note that the procedure bodies may have local data, in the normal way.

In order to call a procedure of a monitor, it is necessary to give the name of the monitor as well as the name of the desired procedure, separating them by a dot:

Hoare's work on refining the concept of monitor was published in 1974.



A Concurrent Pascal Compiler for minicomputers

Δ

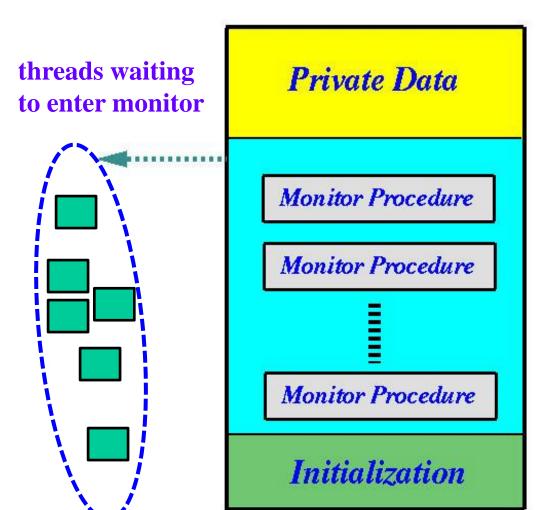
What Is a Monitor? - Basics

- Monitor is a highly structured programming language construct. It consists of
 - ***private** variables and **private** procedures that can only be used within a monitor.
 - **constructors** that initialize the monitor.
 - **A** number of (public) monitor procedures that are available to the users.
- Note that monitors have no public data.
- A monitor is a mini-OS with monitor procedures as system calls.

Monitor: Mutual Exclusion 1/2

- No more than one thread can be executing in a monitor. Thus, mutual exclusion is automatically guaranteed in a monitor.
- When a thread calls a monitor procedure and enters the monitor successfully, it is the only thread executing in the monitor.
- When a thread calls a monitor procedure and the monitor has a thread executing, the caller is blocked outside of the monitor.

Monitor: Mutual Exclusion 2/2



- If there is a thread executing in a monitor, any thread that calls a monitor procedure is blocked outside of the monitor.
- When the monitor has no executing thread, one waiting thread will be let in.

Monitor: Syntax

```
monitor Monitor-Name
   local variable declarations;
   Procedure1 (...)
   { // statements };
   Procedure2 (...)
   { // statements };
   // other procedures
      // initialization
```

- All variables are private.Why? Exercise!
- Monitor procedures are public; however, some procedures may be private so that they can only be used within a monitor.
- Initialization procedures (i.e., constructors) execute only once when the monitor is created.

Monitor: A Very Simple Example

```
monitor IncDec
                            thread Increment
                            while (1) {
   int
         count;
                                // do something
   void Increase(void)
                               .IncDec.Increase();
     count++; }
                                cout <<
                                  IncDec.GetData();
   void Decrease(void)
                                 / do something
     count--; }
                                 Is the printed value the
   int GetData(void)
                                 one just updated?
       return count;
                          initialization
      count = 0; }
                                                  9
```

Condition Variables

- Mutual exclusion is an easy task with monitors.
- While a thread is executing in a monitor, it may have to wait until an event occurs.
- Each programmer-defined event is conceptually represented by a condition variable.
- A condition variable, or a condition, has a private waiting list, and two public methods: signal and wait.
- Note that a condition variable has no value and cannot be modified.

Condition wait

- Let cv be a condition variable. The use of methods signal and wait on cv are cv.signal() and cv.wait().
- Condition wait and condition signal can only be used in a monitor.
- A thread that executes a condition wait blocks immediately and is put into the waiting list of that condition variable. The monitor becomes "empty" (i.e., no executing thread inside).
- This means that this thread is waiting for the indicated event to occur.

Condition signal

- Condition signal is used to indicate an event has occurred.
- If there are threads waiting on the signaled condition variable, one of them will be released.
- If there is no waiting thread waiting on the signaled condition variable, this signal is lost as if it never happens.
- Consider the released thread (from the signaled condition) and the thread that signals. There are two threads executing in the monitor, and mutual exclusion is violated! Something has to be done to fix this problem.

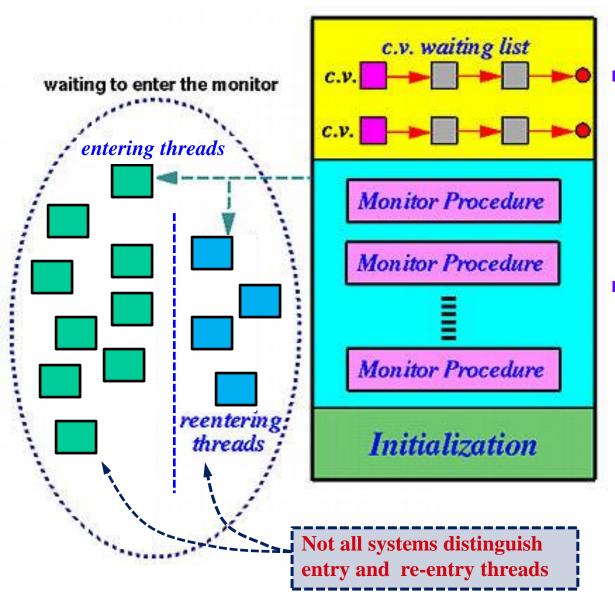
Two Types of Monitors

- After a signal, the released thread and the signaling thread may be executing in the monitor.
- There are two approaches to address this issue:
 - **Hoare Type** (proposed by C. A. R. Hoare)¹: The released thread takes the monitor and the signaling thread waits somewhere.
 - **Mesa Type** (proposed by Lampson and Redell)²: The released thread waits somewhere and the signaling thread continues to use the monitor. This is also used in Java.
 - 1. C. A. R. Hoare, Monitors: An Operating System Structuring Concept, Communications of the ACM, Vol. 17 (1974), No. 10 (October), pp. 549-557.
 - 2. Butler W. Lampson and David D. Redell, Experience with Process and Monitor in Mesa, Communications of the ACM, Vol. 23 (1980), No. 2 (February), pp. 105-117.

What Do You Mean by "Waiting Somewhere"?

- The signaling thread (Hoare type) or the released thread (Mesa type) must wait somewhere.
- You could consider there is a waiting bench for these threads to wait.
- Hence, each thread that involves in a monitor call may be in one of the four states:
 - ***Active:** The running one.
 - **Entering:** Those blocked by the monitor.
 - ***Waiting:** Those waiting on a condition variable.
 - **Inactive:** Those waiting on the waiting bench.

Monitor with Condition Variables



- Threads blocked due to signal/wait are in the re-entry list (i.e., waiting bench).
- When the monitor is free, a thread is released from either entry or reentry.

What Is the Major Difference?

```
Condition UntilHappen;

// Hoare Type
if (!event)
UntilHappen.wait();

// Mesa Type
while (!event)
UntilHappen.wait();
```

Unless stated otherwise, we only use the Hoare type monitors in this course.

With Hoare type, once a signal arrives, the signaler yields the monitor to the released thread and the condition is not changed. Thus, an if is sufficient.

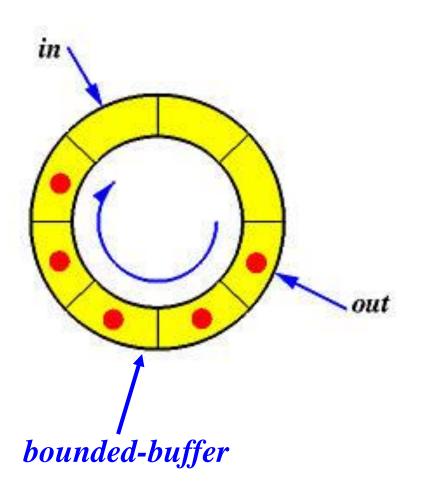
With Mesa type, the released thread may be waiting for a while before it runs. During this period, other threads may be in the monitor and change the condition. It is better to check the condition again with a while!

Examples

- ☐ Producer/Consumer
- **□** Dining Philosophers
- ☐ Alarm Clock
- ☐ Readers-Writers (Reader Priority)
- ☐ Readers-Writers (Take Turns)

We use the Hoare type monitors unless stated otherwise

Monitor: Producer/Consumer



```
monitor ProdCons
  int count, in, out;
  int Buf[SIZE];
  condition
     UntilFull,
     UntilEmpty;
  procedure PUT(int);
  procedure GET(int *);
  \{ count = 0 \}
```

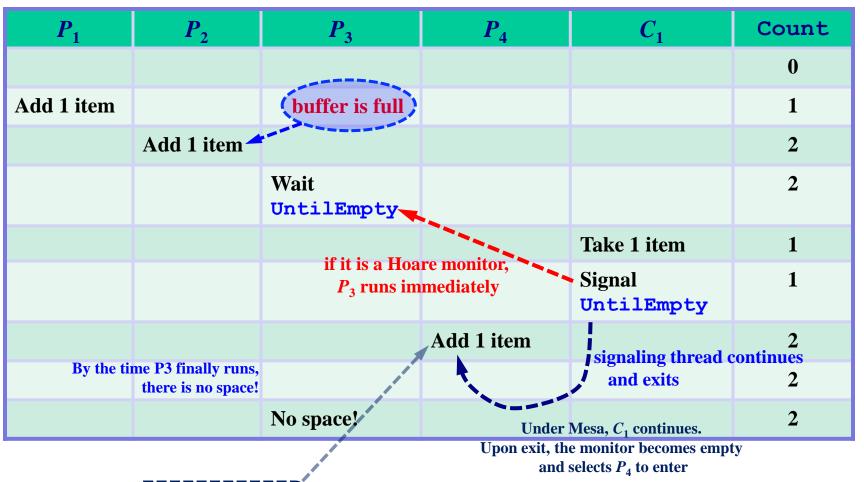
Monitor: PUT() and GET()

```
void PUT(int X)
  if (count == SIZE)
    UntilEmpty.wait();
  Buf[in] = X;
  in = (in+1) %SIZE;
  count++;
  if (count == 1)
    UntilFull.signal()
```

```
void GET(int *X)
  if (count == 0)
    UntilFull.wait();
  *X = Buf[out];
  out=(out+1)%SIZE;
  count--;
   f (count == SIZE-1)
    UntilEmpty.signal();
```

Run This Solution with Mesa?

Buffer Size = 2



monitor is empty.

allow one to enter.

Dining Philosophers, Again!

- Let us look at another solution to the dining philosophers problem.
- Recall that slides 138-139 of 08-Semaphores.pdf discussed a solution suggested by Dijkstra. This solution requires that a philosopher can eat only if he can get BOTH chopsticks at the same time.
- This solution can be implemented using a monitor easily.

Monitor Definition

```
monitor philosopher
  enum { THINKING, HUNGRY,
         EATING} state[5];
  condition self[5];
  private: CanEat(int);
  procedure GET(int); // get BOTH chopsticks
  procedure PUT(int); // release chopsticks
  { for (i=0;i<5;i++)
      state[i] = THINKING;
```

The CanEat() Procedure

```
checking whether philosopher k can eat
                                  the left and right neighbors of
                 CanEat (int k) philosopher k are not eating
                ((state[(k+4) %5] != EATING) & &
                  (state[k] == HUNGRY) &&
                  (state[(k+1)%5] != EATING)
  philosophe
                      state[k] = EATING;
  k is hungry
                      self[k].signal();
```

If the left and right neighbors of philosopher *k* are not eating and philosopher *k* is hungry, then philosopher *k* can eat. Thus, release him!

The GET() and PUT() Procedures

```
I am hungry
void GET(int i)
                               see if I can eat
   state[i] = HUNGRY
                                       If I could not eat,
   CanEat(i);
                                       block myself
   if (state[i] != EATING)
       self[i].wait();
                           void PUT(int i)
      I finished eating
                               state[i] = THINKING;
  Let my neighbors eat.
                              CanEat((i+4) % 5);
                              CanEat((i+1) % 5);
```

How about Deadlock?

- This solution does not cause deadlock, because
 - 1. The only place where eating permission is granted is in procedure CanEat(), and
 - 2. Philosopher *k* can eat only if he could get both chopsticks (i.e., no hold-and- wait and no circular waiting).

How about Bounded Waiting?

Question: The Progress condition is meet and could be proved easily. How about the Bounded Waiting condition? More precisely, is it possible that some philosophers can continue the process of thinking and eating and block some others indefinitely? Exercise.

A Simple Alarm Clock: 1/4

- A set of Sleeper threads wish to Slumber for various times and set an alarm clock to wake them when it is time to get up.
- Unfortunately, their alarm clock is a little primitive:
 - > Every hour it squirts cold water at the nearest sleeper, who immediately prods the next sleeper.
 - Each sleeper checks the time: if it is not the time for him/her to go to work, then he/she goes back to sleep.
- An external thread calls a monitor procedure every hour to initiate this waking up operation.

A Simple Alarm Clock: 2/4

Monitor Definition

A Simple Alarm Clock: 3/4

Procedure Tick

A Simple Alarm Clock: 4/4

Procedure Slumber

- The while loop controls the number of hours a Slumber can sleep.
- * The Tick () procedure updates the time and wakes up the first slumber.
- **This Slumber** wakes up the next one. The Signal is lost if no one there.
- ***** This is referred to as cascading release/signal.
- Cascading release can release all waiting threads even though the # is unknown.

The Readers-Writers Problem Reader Priority

- We still need a reading reader count reading and a waiting reader count readers.
- Two condition variables are needed: OK_to_Read and OK to Write:
 - >OK_to_Read: readers wait here if they cannot read because a writer is writing.
 - >OK_to_Write: writers wait here if they cannot write because a writer is writing, or readers are reading.

Monitor Definition

```
monitor reader-writer
  int reading = 0; // reading readers
  int readers = 0; // waiting readers
  Bool busy = FALSE; // writer is writing
  condition OK to Read, OK to Write;
  procedure read REQUEST(void);
  procedure read RELEASE(void);
  procedure write REQUEST(void);
  procedure write RELEASE(void);
```

Readers and Writers

Reader

```
while (1)
{
    // do something
    read_REQUEST();
    // reading
    read_RELEASE();
    // do something
}
```

Writer

```
while (1)
{
    // do something
    write_REQUEST();
    // writing
    write_RELEASE();
    // do something
}
```

Monitor Code for Readers

```
Reader Priority:
void read REQUEST()
                                 If there is a writer writing,
                                   this reader waits!
  if (busy) {
                       // if a writer is writing
    readers++;
                         this reader must wait
    OK to Read.wait(); // wait on OK to read
                         // released!
    readers--;
  reading++;
                         // if not busy or released
  OK to Read.signal(); // let the next reader to go
void read RELEASE()
                        // a reader has done reading
  reading--;
  if (reading == 0) // is this reader the last one?
    OK to Write.signal(); // YES, allow a writer to go
                                                       35
      if there is no reader reading, yield to a writer;
```

Monitor Code for Writers

```
Reader Priority:
void write REQUEST()
                                   As long as there is a reader reading,
  if (busy || reading != 0)
                               // if a writer is writing
                               // or readers are reading
    OK to Write.wait();
                               // this writer must wait
  busy = TRUE;
                               // otherwise, start to write;
                                   If there is no readers reading.
void write RELEASE()
                                   yield to a writer
  busy = FALSE;
                               // a writer done writing
  if (readers != 0)
                               // if some waiting readers
    OK to Read.signal();
                                    allow a reader to go
  else
                               // otherwise
    OK to Write.signal();
                                  allow a writer to go
                                                         36
```

Monitor Code: Summary

```
void read REQUEST()
                                   void write REQUEST()
                 readers only block here
  if (busy) {
                                      if (busy || reading != 0)
     readers++; 👱
                                        OK to Write.wait();
    OK to Read.wait();
                                      busy = TRUE;
     readers--;
                                      the exiting writer releases a writer
                cascading release.
                                       only if there is no waiting readers
  reading++
  OK to Read.signal()
                                      the exiting writer releases
                                     a waiting reader
void read RELEASE()
                                   void write RELEASE()
                                      busy = FALSE;
  reading--;
                                      if (readers != 0)
  if (reading == 0).
     OK to Write.signal();
                                        OK to Read.signal();
            the last reader releases
                                      else
                  a waiting writer
                                        OK to Write.signal();
                                                                   37
```

only if there is no readers waiting or reading, a writer is released

The Reader-Writer Problem: Again!

- Let us add a minor modification to the readerswriters (priority version) problem to make it a bit more realistic:
 - ➤ If a writer is waiting, the new readers should yield to a writer.
 - > Upon the exit of a reader,
 - ✓ If there are waiting writers, let one go
 - > Upon the exit of a writer,
 - ✓ If there are waiting readers, let one go
 - ✓ If there are waiting writers, let one go
 - > So, the readers and writers take turns.

Monitor Definition

```
monitor reader-writer
  int readers = 0; // waiting readers
  int reading = 0; // reading readers
  int writing = 0; // writing writers
  int writers = 0; // waiting writers
  condition OK to Read, OK to Write;
  procedure read REQUEST(void);
  procedure read RELEASE(void);
  procedure write REQUEST(void);
  procedure write RELEASE(void);
```

Monitor Code for Readers

```
void read REQUEST()
                                    a writer is writing.
                                    this reader waits!
  if (writing > 0 || writers > 0) {// if a writer writing
                          // or there are waiting writers
    readers++;
                          // this reader waits
    OK to Read.wait();
                          // this reader is released
    readers--;
  reading++;
                          // one more reading reader
  OK to Read.signal(); // allow other readers to read
void read RELEASE()
  if (--reading == 0) // if this is the last reader
    OK to Write.signal(); // let a writer go
```

Monitor Code for Writers

```
void write REQUEST()
  if (writing > 0 || readers > 0) {// if a writer writing
             // or there are waiting readers
    writers++;
    OK to Write.wait(); // this writer waits
                        // this writer is released
    writers--;
                        // this writer starts writing
  writing = 1;
void write RELEASE()
  writing = 0;
              // this writer finishes writing
  if (readers > 0) // if there are readers waiting
    OK to Read.signal(); // let a reader to go
  else
                         // otherwise
    OK to Write.signal(); // let a writer to go
                                                  41
```

Monitor Code: Summary

readers \rightarrow writer \rightarrow readers \rightarrow writer \rightarrow ... (Taking Turns)

```
void read REQUEST()
                                    void write REQUEST()
 { wait if writers are waiting or writing
                                    {wait if writer writing or readers waiting
  if (writing > 0
                                      if (writing > 0
          || writers > 0)
                                              || readers > 0)
     readers++;
                                        writers++;
     OK to Read.wait();
                                         OK to Write.wait();
                                        writers--;
     readers--;
   reading++;
                                      writing = 1;
   OK_to_Read.signal(); cascading release
                                   the exiting writer releases
                                           a waiting reader
void read RELEASE()
                                    void write RELEASE()
  the last reader releases a writer
                                                    no waiting readers
   if (--reading == 0)
                                      writing = 0; let a writer go
     OK to Write.signal();
                                      if (readers > 0)
                                         OK to Read.signal();
                                      else
   if there is no waiting writers,
                                       ►OK to Write.signal()4?
      these signals are lost!
```

The Reader-Writer Problem Exercise...

- Consider a solution that only uses a single condition, Take_Turn, rather than two:
 OK_to_Read and OK_to_Write.
- A reader waits on Take_Turn until there are no waiting writers, and then waits on Take_Turn again until there is no writer writing. In this case, this reader is safe to read.
- A writer waits on Take_Turn if there are readers reading or a writer writing.
- A reader or writer signals Take_Turn when they finish reading or writing.

 43

Monitor Definition

```
monitor reader-writer
  int reading = 0; // reading readers
  int writing = 0; // writing writers
  int writers = 0; // waiting writers
  condition Take Turn;
  procedure read REQUEST(void);
  procedure read RELEASE(void);
  procedure write REQUEST(void);
  procedure write RELEASE(void);
```

Monitor Code for Readers

```
void read REQUEST()
  if (writers > 0) // if there are writers waiting,
    Take Turn.wait(); // wait until released
 if (writing > 0) // if there is a writer writing,
    Take Turn.wait(); \forall // wait, of course

    ★ because no writer is writing,

 reading++;
                            a reader can read
                          no writers waiting
void read RELEASE()
  reading--;
                      |// a reader has done reading
  Take Turn.signal(); \// let one reader/writer to go
                no writer writing
```

Monitor Code for Writers

```
void write REQUEST()
                       this means reading or writing being non-zero
                       // one more write request
  writers++;
  if (reading || writing) // if there are readers reading
    Take Turn.wait(); // or a writer writing, wait
                       // if no readers reading and no
  writing++;
                       // writer writing, then go!
void write RELEASE()
  writing--;
                       // reduce writing count
                       // reduce writer count
  writers--;
  Take Turn.signal(); // let some one to proceed
```

The Reader-Writer Problem Question...

- Is this solution correct?
- If you think that this solution is correct, please prove that this solution does implement the requirements correctly (*i.e.*, taking turns correctly).
- If you think that this solution is incorrect, then
 - Explain what the problems are and use execution sequences to show that this solution fails to implement taking turns correctly.
 - >And, modify this solution to make it working!

Three Extensions empty(), Priority wait(), broadcast (i.e., signal_all)

The Reader-Writer: Again and Again!

- Some monitor implementations may have additional features. For example, in addition to wait() and signal(), a condition variable may have one more function empty(), which returns TRUE if the condition variable has no waiting threads.
- C. A. R. Hoare proposed the following version:
 - > A new reader should not be permitted to start if there is a waiting writer;
 - ➤ At the end of a write operation, waiting readers are given preference over waiting writers.
- We have seen this previously.

Monitor Definition

```
monitor reader-writer
  int reading = 0;  // reading readers
 Bool writing = FALSE; // writing?
  condition OKtoRead, // readers wait here
             OKtoWrite; // writers wait here
  procedure read REQUEST(void);
  procedure read RELEASE(void);
  procedure write REQUEST(void);
  procedure write RELEASE(void);
```

Monitor Code for Readers

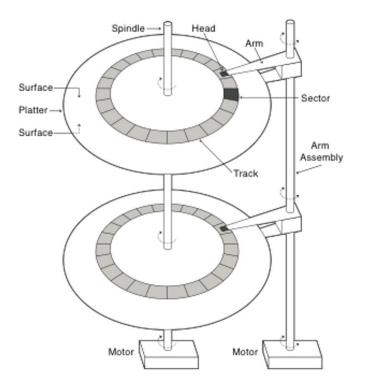
```
--- not all systems have empty ()
void read REQUEST()
  if (writing || !empty(OKtoWrite);) // if writing or
    OKtoRead.wait(); // writer queue not empty, wait!
                       // then this reader can read
  reading++;
                       // allow a waiting reader to go
  OKtoRead.signal();
                       // cascading release here!
void read RELEASE()
                       // a reader has done reading
  reading--;
  if (reading == 0)
    OKtoWrite.signal();
                       because of empty (), a waiting
```

writers count is eliminated

Monitor Code for Writers

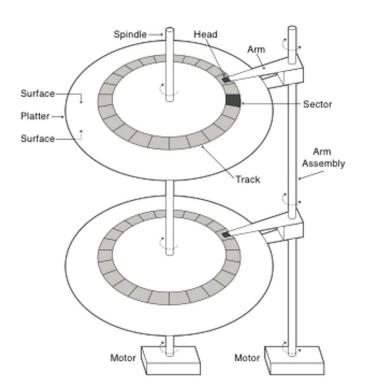
```
void write REQUEST()
  if (writing || reading > 0) // if someone writing
    OKtoWrite.wait(); // or readers waiting, wait here
  writing = TRUE;  // after released, write!
                                --- not all systems have empty ()
void write RELEASE()
 writing = FALSE;  // finished writing
  if (! empty(OKtoRead)) // if there are waiting readers
    OKtoRead.signal(); // let one go (cascading!)
  else
    OKtoWrite.signal(); // or let a waiting writer go
                         // if there is no waiting
                            writer, this signal is lost.
because of empty (), a waiting
                                                      52
  readers count is eliminated
```

Disk Organization: 1/2



- The information stored on the disk is organized into a series of cylinders, each of which consists of several tracks, and each track consists of several sectors or blocks.
- When a read/write request comes, the disk head assembly is moved to the correct cylinder, the head corresponding to the required track is selected, and the information may be read or written when the required sector passes the head as the disk rotates.
- The time to move the head assembly to a track is referred to as the *seek* time, which is mechanical and time consuming.
- Because there are many cylinders to be accessed, how can we minimize the total seek time?

Disk Organization: 2/2



- Instead accessing the cylinders based on the incoming order, it would be more efficient and probably cause less mechanical strain on the device if the requests could be handled by making continuous sweeps across the disk surfaces, first in one direction and the other.
- Some systems have an extended wait() method.
- The wait (Priority) has an integer argument Priority.
- When a condition variable is signaled, the thread with the highest priority is released. A smaller value of Priority usually means a higher priority. WARNING: check the system for the details.

Monitor Definition

```
monitor Disk_Scheduler
  int Head = 0;  // initially at cylinder 0
  int Dir = IN;  // thus, IN sweep first
  int MAX = ...; // highest cylinder number
 Bool Busy = FALSE; // initially the disk is free
  condition IN Sweep, // IN sweep queue
            OUT Sweep; // OUT sweep queue
  procedure Request(int Cy); // request to use cylinder Cy
  procedure Release(void); // release cylinder Cy
```

The Use of Disk_Scheduler

```
Request(Cylinder); // request to move the
// disk assembly to
// cylinder Signal all
activate the desired track
read or write the needed sector/block

Release(); // release the cylinder
// after the access
```

Monitor Code for Request

```
void Request(int Cy) // request to access Cylinder Cy
                                     // if disk is busy, wait!
   if (Busy) {
      if ((Head < Cy) ||
                                     // if beyond the current pos
          (Head == Cy && Dir == IN)) // or at the current pos
                                             & direction is IN
                                     // wait on IN sweep CV
         IN Sweep.Wait(Cy);
      else
                                     // Otherwise, wait for OUT
         OUT_Sweep.Wait(MAX - Cy); // sweep using a reversed
                                     // priority order
                                     // Once get out, can access
   Busy = TRUE;
   Head = Cy;
                                     // the disk head is at Head
```

When the disk head moves inward, cylinders are accessed in an ascending order.

Thus, a smaller cylinder number means a higher priority.

When the disk head moves outward, cylinders are accessed in a descending order.

Thus, a smaller cylinder number means a lower priority, and MAX — Cy is used.

Monitor Code for Release

```
void Release(void)
                               // request to release
                               // done using cylinder Cy
  Busy = FALSE;
   if (Dir == IN) {
                               // IN sweep?
     if (!IN Sweep.empty())
                               // if there are waiting in IN
         IN Sweep.signal();
                                     release one of them
     else {
                               // no one waiting in IN
                                     switch direction to OUT
         Dir = OUT;
                                     release one of them
         OUT Sweep.signal();
                                   symmetric for IN and OUT
                               // current OUT sweep
   else {
     if (!OUT Sweep.empty())
                                  if there are waiting in OUT
         OUT Sweep.signal();
                                     release one of them
                               // no one waiting in OUT
     else {
                                     switch direction to IN
         Dir = IN;
         IN Sweep.signal();
                                     release one of them
                                                          58
```

The Broadcast/Signal_all Method

- In some systems, each condition variable may have a broadcast or signal_all method.
- Let cv be a condition variable. cv.broadcast() or cv.signal_all() releases all waiting threads on condition cv with a single signal call.
- When cv.broadcast() (or cv.signal_all()) is called, all waiting threads on condition variable cv are released and changed to inactive/re-entry.
- This will save the use of a cascading release.

A Simple Alarm Clock Procedures Tick & Slumber

```
void Tick(void) // called by an external timer
                 // every hour
  Now = Now + 1; // This is internal hour count
  Wake.broadcast(); // wake up all sleepers
void Slumber(int n) // n = sleeping hours
  int AlarmCall;
                     // time to wake up
                    cascading release is not needed
  AlarmCall = Now + n; // update my alarm clock
  while (Now < AlarmCall) { // as long as I can sleep
    Wake.Wait(); // sleep
    Wake.signal(); // no more needed
                                                60
```

Readers-Writers (Reader Priority)

```
void read REQUEST()
  if
      (busy)
    readers++:
    OK to Read.wait();
    readers--;
              cascading release
              no more needed.
  reading++
  -OK-to-Read.signal()-,-
                 cascading release
void read_RELEASE()
  reading--;
  if (reading == 0)
    OK to Write.signal();
            the last reader releases
                 a waiting writer
```

```
void write REQUEST()
  if (busy || reading != 0)
     OK to Write.wait();
  busy = TRUE;
  the exiting writer releases a writer
   only if there is no waiting readers
  the exiting writer releases
  a waiting reader
void write RELEASE()
  busv = FALSE;
                        release ALL
  if (readers != 0) waiting readers
    OK to Read.broadcast(
  else
     OK to Write.signal();
                                61
```

Readers-Writers: Taking Turns

```
void write REQUEST()
void read REQUEST()
  if (writing > 0
                                  if (writing > 0
         || writers >
                                         || readers > 0) {
    readers++; 🗡
                                    writers++;
   OK to Read.wait();
                                    OK to Write.wait();
    readers--;
                                    writers--;
               cascading release
               no more needed
  reading++;
                                  writing = 1;
  OK-to-Read-signal-()->
cascading release
  is no more needed
void read RELEASE()
                                void write RELEASE()
  if (--reading == 0)
                                  writing = 0;
                                                     release ALL
                                  if (readers > 0) \( \text{waiting readers} \)
    OK to Write.signal();
                                    OK to Read.broadcast();
                                  else
                                    OK to Write.signal();
```

An Important Note

cv.broadcast() is not the same as

```
while (!cv.empty())
    cv.signal();
```

- In each iteration of the while loop, a waiting thread is released, and this released thread gets the monitor (Hoare type) while the signaling thread becomes inactive/re-entry.
- When the released thread exits, the monitor mechanism picks a thread to enter. If the picked one is the signaling thread, the situation is fine.
- Otherwise, the newcomer could cause problems₆₃

Example: Barrier 1/8

- Monitor Barrier has an initial value n > 0 and a Barrier_wait() method.
- A thread that calls $Barrier_wait()$ blocks if the number of blocked threads is less than n-1.
- When the n-th thread calls Barrier_wait(), it releases all waiting n-1 threads.
- In other words, barrier Barrier blocks the first n-1 threads that calls Barrier_wait() and the n-th one releases all waiting threads. Thus, all n threads would act in a single group.

Example: Barrier 2/8

```
monitor Barrier
   int n;
                                // the maximum threads to be queued
   int count = 0;
                                // counting the current # queued
   condition barrier;
                                // queue threads here
   procedure Barrier wait() // method for threads to call
              cascading release
                                // one more thread entering barrier
      count++;
                                // is this the last one?
      if (count \( \frac{1}{2} \) n)
         barrier.wait();
                                     no! queue this thread
         barrier.signal();
                                    when released, releases the next
      else {
                                // this is the n-th thread ...
        count = 0;
                                // reset the counter
         barrier.signal();
                                // releases a waiting one and cascades
                    must be before the signal()
```

Example: Barrier 3/8

What if the order of count=0 and
barrier.signal() is switched as follows?
barrier.signal();
count = 0;

Does this solution work? In other words, after switching the two statements in the else part, can the n-th thread correctly release the waiting n-1 threads?

Example: Barrier 4/8

- NO, this does not work properly.
- Suppose n = 3. We may have the following scenario:
 - $\succ T_1$ calls Barrier_wait() and waits (count = 1)
 - $\succ T_2$ calls Barrier_wait() and waits (count = 2)
 - $\succ T_3$ calls Barrier_wait() and executes the else:
 - $\checkmark T_3$ calls barrier. signal () (currently count = 3) and assume that T_1 is released. T_3 yields the monitor.
 - $\checkmark T_1$ is active and calls barrier. signal (). This releases T_2 .
 - $\checkmark T_2$ calls barrier.signal() and exits.
 - ✓ Who will enter? If it is the signaling thread T_3 , it is OK. Otherwise, a new thread enters, sees count < 3 being FALSE and go to else and start another signal (). This is a wrong program logic.

Example: Barrier 5/8

```
monitor Barrier
                              // the maximum threads to be queued
   int n;
   int count = 0;
                              // counting the current # queued
   condition barrier;
                              // queue threads here
   procedure Barrier wait() ___// method for threads to call
              all threads waiting here are released
      count++; and cascading release is not needed; count++;
                            // is this the last one?
      if (count < n) {--
       - barrier.wait();
                               // no! queue this thread
        -barrier.signal()-;-
      else {
                              // this is the n-th thread ...
        count = 0;
                             // reset the counter
        barrier.broadcast(); // releases a waiting one and cascades
```

Example: Barrier 6/8

```
monitor Barrier
                          // the maximum threads to be queued
  int n:
  int count = 0, i;
                         // counting the current # queued
  condition barrier;
                          // queue threads here
  procedure Barrier wait() // method for threads to call
     count++;
                        // one more thread entering barrier
                          // is this the last one?
     if (count < n)
       else {
       for (i = 1; i \leq n-1; i++) // releases the threads using for
         barrier.signal(); //
                                one at a time
                    if this is the n-th thread, there are n-1 threads waiting
```

NOTE: This is a Hoare type monitor. This is wrong!₆₉

Example: Barrier 7/8

T_1	T_2	$T_{^{\circ}3}$	barrier	count	i
			Ø	0	
count++			Ø	1	
wait()			T_1	1	
	count++		T_1	2	
	wait()		T_1, T_2	2	
T_1 is released $Count++$			T_1, T_2	3	
		count = 0	T_1, T_2	0	
T_1 comes back fast signal ()			T_2	0	1
continue & exit			T_2	0	1
come back fast			T_2	0	1
count++			T_2	1	1
wait() T_1 is released again!			T_1, T_2	1	1
		signal()	T_1, T_2	1	2
continue			T_2	1	2

Example: Barrier 8/8

- If you use a Hoare type monitor, try to use cascading release if it is possible.
- If broadcast() is available, this is a good alternative of cascading release.
- Do not use a for or a while loop to release the blocked threads on a condition.
- QUESTION: Is the for/while version working correctly under a MESA monitor?
- QUESTION: Are cascading and broadcast releases work correctly under a MESA monitor?

Semaphore and Monitor Equivalence

Semaphore and Monitor Equivalence

- In terms of expressive power, semaphores and monitors are equivalent.
- A semaphore can easily be implemented with a monitor.
- Conversely, a monitor and its condition variables can also be implemented with multiple semaphores, although this is a bit tedious.
- Therefore, semaphores and monitors are equivalent because one may be implemented by the other.

Semaphore Implementation Using a Monitor

Semaphores in Hoare Type: 1/2

- On the right is a possible implementation of a semaphore using a Hoare type monitor.
- count is the semaphore counter.
- c is a condition variable for blocking a thread that must wait.
- Does it work?

```
monitor Hoare Sem
              count = 0;
   int
   condition
              // wait
   P(void)
      count--;
      if (count < 0)
         c.wait();
   V(void)
               // signal
      count++;
      c.signal();
```

Semaphores in Hoare Type: 2/2

- Why is there no if?
- Recall that abs (count)
 is the number of waiting
 threads on that semaphore.
- It count++ is not positive, there are waiting threads and c.signal() will release one.
- If count++ is positive, no waiting threads on c and c.signal() is lost.

```
monitor Hoare Sem
               count = 0;
   int
   condition
               // wait
   P (void)
      count--;
      if (count < 0)
          c.wait();
   V(void)
               // signal
      count++;
      c.signal();
```

Semaphores in Mesa Type: 1/2

- Now let us try the same implementation using the Mesa type monitor.
- count is the semaphore counter.
- c is a condition variable for blocking a thread that has to wait.
- Does it work?

```
monitor Mesa Sem
              count = 0;
   int
   condition
              // wait
   P(void)
      count--;
      while (count < 0)
         c.wait();
   V(void)
              // signal
      count++;
      c.signal();
```

Semaphores in Mesa Type: 2/2

T_1	T_2	T_3	count	Comment
			0	
count	ID	T. A.	-1	T_1 calls P ()
while ()		esa type, T_2 continu Jpon exit, T_3 (rather	1	
c.wait	than T_1) gets the	control of monitor.	-1	T_1 blocks
	count++		0	T_2 calls V ()
	c.signal		0	T_1 released
waiting to	T ₂ exit	count	-1	T_3 calls P ()
re-enter		while ()	-1	
		c.wait	-1	T_3 blocks
while ()			-1	T ₁ re-enters
c.wait	a de la companya de		- 1	T_1 blocks
	T_1 and	T_3 are both blocked	-1	T_1 and T_3 both blocked

A correct signal fails to release a waiting thread properly.

Hoare Monitor Implementation Using Semaphores

Simulating a Hoare Monitor Using Semaphores: 1/4

- Because the boundary of a monitor must be mutually exclusive, a mutex lock is needed.
- Let this mutex lock be Mutex.
- Because the signaling thread must yield the monitor and wait, a semaphore Suspended, initialized to 0, is needed for this purpose.
- Threads waiting on Suspended are inactive ones.
- Suspended_count counts the number of threads suspended on semaphore Suspended.

Simulating a Hoare Monitor Using Semaphores: 2/4

Each monitor procedure has the following form:

- If there are inactive/re-entry threads, let one go!
- The baton, Mutex, is given to the released thread.
- Therefore, this is in favor of the inactive/re-entry threads.

Simulating a Hoare Monitor Using Semaphores: 3/4

• For condition cv, a semaphore cv_Sem and a counter cv Count are needed, initialized to 0.

```
cv.wait():
                             // one more waiting
   cv Count++;
  if (Suspended count > 0)
                             // some inactive waiting
      Suspended.signal();
                             // let one go
  else
                             // otherwise, no inactive
                                   release the monitor
    Mutex.signal();
   cv Sem.wait();
                             // condition wait
   cv Count--;
                             // cv signal received
                             // cv count decreases
```

- Before a thread waits on a cv, an inactive/re-entry thread is released if there is one.
- Otherwise, it releases the Mutex.
- Then, wait on the cv.

Simulating a Hoare Monitor Using Semaphores: 4/4

Here is cv.signal():

- Before the signaling thread yields the monitor to the released one,
 the signaling thread signals cv_Sem to release a waiting thread on that cv.
- If there is no inactive thread, this signal is considered never happened.

Simulating a Mesa Monitor Using Semaphores

- We have seen the simulation of a Hoare type monitor using semaphores.
- It is certainly possible to simulate the Mesa type monitor using semaphores.
- Please think this over and do it as an exercise.
- The implementation discussed earlier gives inactive/re-entry threads higher priority to enter a monitor. What if there is no such requirement? In other words, an inactive/re-entry thread is pushed outside of the monitor and waits alongside other entering threads?

Comparisons

Hoare Type vs. Mesa Type

- When a signal occurs, a Hoare type monitor uses two context switches, one switching the signaling thread out and the other switching the released in. However, a Mesa type monitor uses one.
- Thread scheduling must be very reliable with Hoare type monitors to ensure once the signaling thread is switched out the next one to execute in the monitor must be the released thread.
- With Mesa type monitors, a condition may be evaluated multiple times. However, incorrect signals will do less harm because every thread checks its own condition.

Semaphores vs. Monitors

Semaphores	Monitors
Can be used anywhere, but should not be in a monitor	Can only be accessed with monitor procedure calls
No connection between the semaphore and the data this semaphore protects	Data and access procedures are in the same place (i.e., a monitor)
Semaphores are low level assembly language-like instructions	Monitors are well-structured higher-level construct
Not easy to use and prone to bugs	Easy of use and good protection of vital data

Semaphores vs. Conditions

Semaphores	Condition Variables
Can be used anywhere, but not in a monitor	Can only be used in monitors
wait() does not always block its caller	wait() always blocks its caller
signal () either releases a thread, or increases the semaphore counter	signal () either releases a thread, or the signal is lost as if it never occurs
If signal () releases a thread, the caller and the released both continue	If signal () releases a thread, either the caller or the released continues, but not both

Remarks

Reminder 1

It is suggested that a signal () should be the last executed statement before exiting a monitor procedure. We have seen this in many examples.

Why?

- ➤ Whether this monitor is a Hoare type or Mesa type, the difference does not matter.
- ➤ If this is a Hoare type, the signaling thread will not do anything even though it is inactive.
- If this is a Mesa type, the signaling thread will exit immediately.
- **➤**Therefore, this is a safer approach.

Reminder 2

- Why is using semaphores in a monitor a bad idea?
 - If a thread T is executing in monitor A, monitor A is "occupied" by thread T.
 - ➤ What if thread *T* calls a semaphore wait in monitor **A** and is blocked?
 - Estimate Because thread T is executing in monitor A, if it waits on a semaphore in monitor A, monitor A cannot be used as monitor A is occupied by the thread T.
 - Therefore, using semaphores in a monitor is not a good idea.

 91

Reminder 3

- Why is calling a monitor procedure in another monitor from this monitor a bad idea?
 - If a thread T is executing in monitor A, monitor A is "occupied" by thread T.
 - ▶ If thread T successfully enters another monitor B, monitor A is not empty and cannot be used by any thread. This is inefficient!
 - ➤ Worse, what if thread *T* waits on a condition variable in monitor **B**? Then, monitor **A** could not be used for a long time until thread *T* returns from monitor **B** and exits.

Monitors with ThreadMentor

ThreadMentor does not support empty(), Priority wait() and broadcast() or signal_all()

Monitor: Definition

```
class MyMon::public Monitor
 public:
   MyMon(); // constructor
   MonitorProcedure-1();
   MonitorProcedure-2();
    // other procedures
 private:
    // variables used in
    // this monitor
};
```

- A monitor must be a derived class of class Monitor.
- The initialization part should be in constructors.
- Make monitor procedures public.
- Local variables should be private/protected.

Monitor: Monitor Procedures

```
int MyMon::MonProc(...)
{
    MonitorBegin();
    // other statements
    // of this procedure
    MonitorEnd();
}
```

MonitorBegin () locks the monitor and MonitorEnd() unlocks it. Thus, mutual exclusion is guaranteed.

- Monitor procedures are C/C++ functions.
- Before you do anything, call MonitorBegin().
- Before exit, call MonitorEnd().
- The following is wrong:

```
int MyMon::MonProc()
{
    MonitorBegin();
    // other stuffs
    return 0;
    MonitorEnd();
```

Monitor: A Simple Example

```
Class Count
      ::public Monitor.
  public:
    int
         Inc()
         Dec();.
    int
    Count();
  private:
    int Counter;
Count::Count(void)
{ Counter = 0; }
```

```
int Count::Inc()
   MonitorBegin();
      Counter++;
   MonitorEnd();
   return Counter;
int Count::Dec()
   MonitorBegin();
      Counter--;
   MonitorEnd();
   return Counter;
```

Monitor: Condition Variables

```
Condition Event;
Event.Wait();
Event.Signal();
```

- Condition is a class and has two methods, Wait() and Signal().
- Waiting on a condition variable means waiting for that event to occur.
- Signaling a condition variable means that the event has occurred.

Philosopher Monitor Definition

condition variable pointers, one for each philosopher

```
class Mon::public Monitor
                               Mon:: Mon ()
                get and put
                                f'int i;
                chopsticks
   public:
                                  for (i=0; i < 5; i++){
      Mon();...
                                    State[i] = THINKING;
      GET(int); PUT(int)
   private:
      Condition Self[5];
           State[5];
      int
      int CanEat(int);
};
                                  state of each philosopher
 are both chopsticks available?
```

Philosopher Monitor Implementation

```
void Mon::GET(int k)
int Mon::CanEat(int k)
 if ((state[(k+4)%5] !=
                              MonitorBegin();
      EATING
                              state[k] = HUNGRY;
   &&(state[k] ==
                              CanEat(k);
                              if (state[k] != EATING)
      HUNGRY
   && (state[(k+1)%5]!=
                                 self[k].Wait();
                              MonitorEnd();
      EATING))
    state[k] = EATING;
    self[k].Signal();
                                    if I cannot eat, wait
   check to see if I can eat
```

Specifying a Monitor Type

Replace HOARE with MESA if you wish to use a Mesa type monitor.

- A monitor type must be specified in your monitor constructor.
- Use HOARE or MESA for Hoare type and Mesa type monitors.

The End