# Part III Synchronization Deadlocks and Livelocks

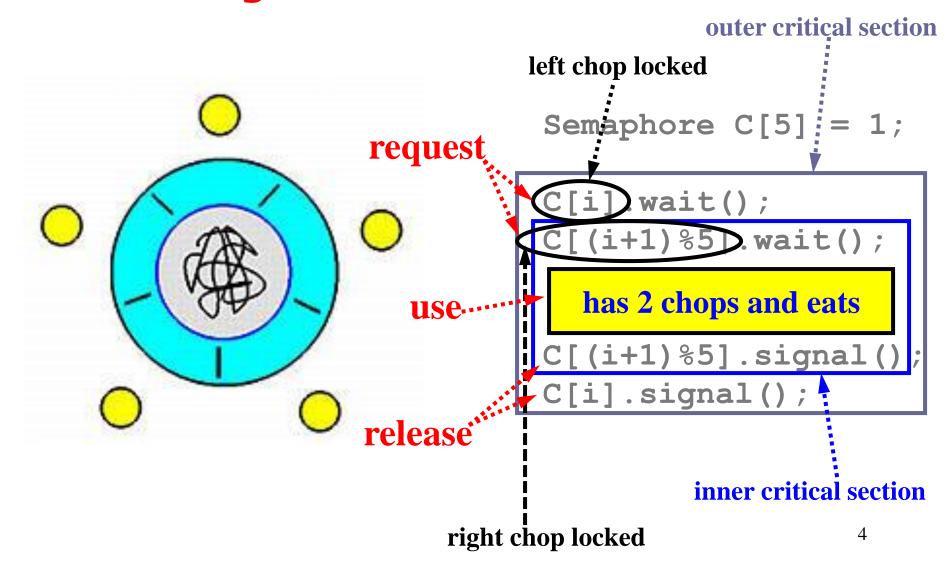
You think you know when you learn, are more sure when you can write, even more when you can teach, but certain when you can program.

### **Deadlock Basics**

#### System Model: 1/2

- System resources are used in the following way:
  - \*Request: If a process makes a request (i.e., semaphore wait or monitor acquire) to use a system resource which cannot be granted immediately, then the requesting process blocks until it can acquire the resource successfully.
  - **Use:** The process operates on the resource (i.e., in critical section).
  - **Release:** The process releases the resource (i.e., semaphore signal or monitor release).

#### System Model: 2/2



#### **Deadlock: Definition**

- A set of processes is in a **deadlock** state when every process in the set is waiting for an event that can only be caused by another process in the same set.
- The key here is that processes are all in the waiting state.

#### **Deadlock Necessary Conditions**

- If a deadlock occurs, then each of the following four conditions must hold.
  - **Mutual Exclusion**: At least one resource must be held in a non-sharable way.
  - **\*Hold-and-Wait:** A process must be holding a resource and waiting for another.
  - **No Preemption:** Resource cannot be preempted.
  - **Circular Waiting:**  $P_1$  waits for  $P_2$ ,  $P_2$  waits for  $P_3$ , ...,  $P_{n-1}$  waits for  $P_n$ , and  $P_n$  waits for  $P_1$ .

#### **Deadlock Necessary Conditions**

- Note that the conditions are necessary.
- This means if a deadlock occurs ALL conditions are met.
- Because  $p \Rightarrow q$  is equivalent to  $\neg q \Rightarrow \neg p$ , where  $\neg q$  means not all conditions are met and  $\neg p$  means no deadlock, as long as one of the four conditions fails there will be no deadlock.

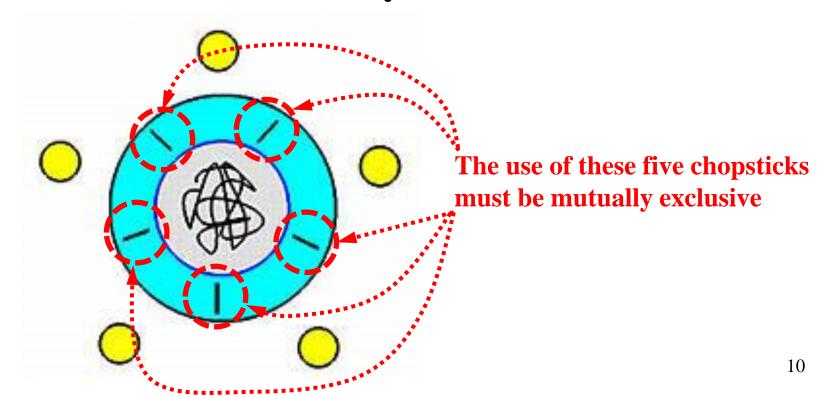
### Deadlock Prevention Make sure deadlock will never happen

#### **Deadlock Prevention: 1/13**

- Deadlock Prevention means making sure deadlocks never occur.
- To this end, if we can make sure at least one of the four conditions fails, there will be no deadlock.

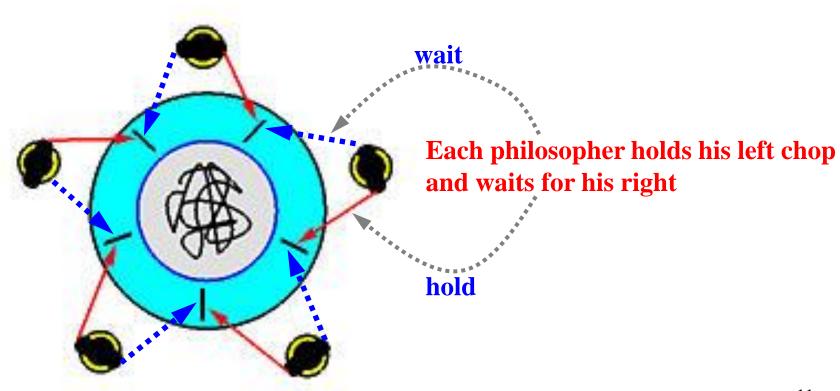
### Deadlock Prevention: 2/13 Mutual Exclusion

 Mutual Exclusion: Some sharable resources must be accessed exclusively, which means we cannot deny mutual exclusion



### Deadlock Prevention: 3/13 Hold-and-Wait

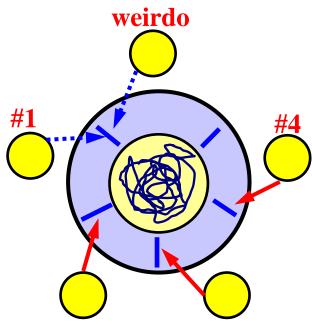
 Hold-and-Wait: A process holds some resources and requests for other resources.



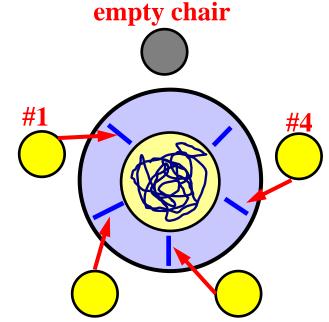
### Deadlock Prevention: 4/13 Hold-and-Wait

- Solution: Make sure no process can hold some resources and request for other resources.
- Two strategies are possible (the monitor solution to the philosophers problem):
  - **A** process must acquire *all* resources before it runs.
  - **\*** When a process requests for resources, it must hold none (i.e., returning resources before requesting for more).
- Resource utilization may be low, since many resources will be held and unused for a long time.
- Starvation is possible. A process that needs some popular resources may have to wait indefinitely.

### Deadlock Prevention: 5/13 Hold-and-Wait



If weirdo is faster than #1, #1 cannot eat and the weirdo or #4 can eat but not both. If weirdo is slower than #1, #4 can eat Since there is no hold-and-wait, there is no deadlock.



In this case, #4 has no right neighbor and can take his right chop.
Since there is no hold-and-wait, there is no deadlock.

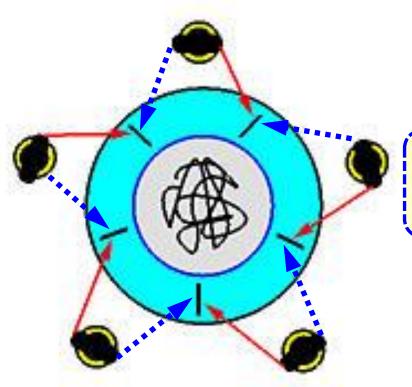
The solution with THINKING-HUNGRY-EATING states forces a philosopher to have both chops before eating. Hence, no hold-and-wait.

### Deadlock Prevention: 6/13 No Preemption

- This means resources being held by a process cannot be taken away (i.e., no preemption).
- To negate this no preemption condition, a process may deallocate all resources it holds so that the other processes can use.
- This is sometimes not doable: while philosopher *i* is eating, his neighbors cannot take *i*'s chops away forcing *i* to stop eating.
- Some resources cannot be reproduced cheaply (e.g., printer). However, the CPU can be preempted. Remember interrupts?

# Deadlock Prevention: 7/13 Circular Waiting

• Circular Waiting:  $P_1$  waits for  $P_2$ ,  $P_2$  waits for  $P_3$ , ...,  $P_{n-1}$  waits for  $P_n$ , and  $P_n$  waits for  $P_1$ .



The weirdo, 4-chair, and monitor solutions all avoid circular waiting and there is no deadlock.

How to break circular waiting in general?
Hierarchical Ordering of resources

# Deadlock Prevention: 8/13 Circular Waiting

Hierarchical Ordering: 1/6

- Resources are ordered in a hierarchical way.
- A process can only request resources <u>higher</u> than those it owns.
- To request lower order resources, a process must release all resources higher than or equal to the requesting ones.

# Deadlock Prevention: 9/13 Circular Waiting

Hierarchical Ordering: 2/6

- Suppose a system orders its resources into 9 levels,  $L_1, L_2, ..., L_8, L_9$ .
- Suppose the highest level of resources a process  $P_i$  has is  $L_5$ .
- Then,  $P_i$  can only request resources at  $L_6$ ,  $L_7$ ,  $L_8$  and  $L_9$ .
- If  $P_i$  wishes to use a resource at  $L_3$ ,  $P_i$  must release all of its resources at  $L_3$ ,  $L_4$  and  $L_5$ .

# Deadlock Prevention: 10/13 Circular Waiting

Hierarchical Ordering: 3/6

- With Hierarchical Ordering, no circular waiting can occur.
  - Suppose that processes  $P_1, P_2, ..., P_n$  involve in a resource waiting chain:  $P_i$   $(1 \le i < n)$  waits for some higher-level resources being held by  $P_{i+1}$ .
  - $\triangleright$  The resource level of  $P_n$  is the highest of all  $P_i$ 's.
    - 1. If  $P_n$  does not need any resource, it may run and release some resources for other processes to use.
    - 2. If  $P_n$  requests further resources at a higher level, because  $P_n$  is the last in the chain  $P_n$  can have these resources and run. Thus, we are back to (1) and Deadlock-free again!

## Deadlock Prevention: 11/13 Circular Waiting

Hierarchical Ordering: 4/6

- Let us do a proof formally.
- Let  $L(P_i)$  be the highest-level of resources  $P_i$  has.
  - Example 2 Because  $P_i$  requests for resources at a higher-level being held by  $P_{i+1}$ , we have  $L(P_i) < L(P_{i+1})$ .
  - > Therefore, we have

$$L(P_1) < L(P_2) < ... < L(P_{n-1}) < L(P_n)$$

 $\triangleright$  If  $P_n$  waits for a resource being held by  $P_1$ , we have

$$L(P_n) < L(P_1)$$

- $\triangleright$  We have  $L(P_1) < L(P_n) < L(P_1)$ , which is impossible.
- Therefore, circular wait is impossible.

# Deadlock Prevention: 12/13 Circular Waiting

Hierarchical Ordering: 5/6

- Let us revisit the dining philosophers problem.
- Chopsticks are ordered as  $C_0 < C_1 < C_2 < C_3 < C_4$ .
- Philosopher i requests chopstick  $C_i$  followed by  $C_{i+1}$  (i.e., lower to higher).
- Philosopher 4 cannot request  $C_4$  followed by  $C_0$ , because this violates the hierarchical ordering. To ensure hierarchical ordering, philosopher 4 must request  $C_0$  followed by  $C_4$  (lower to higher).
- Isn't this the weirdo (or lefty-righty) version?
- Therefore, the weirdo version is deadlock-free.

# Deadlock Prevention: 13/13 Circular Waiting

Hierarchical Ordering: 6/6

Users supervisor mode file system communication I/O & Device virtual mem process hardware

- The hierarchical ordering was used in the THE operating system developed by Dijkstra.
- The lowest layer is process management.
- Each layer *only uses the operations provided* by *lower layers* and does not have to know their implementation.
- Each layer hides the existence of certain data structures, operations and hardware from higher-level layers.
- A major difficulty in the hierarchical ordering is that it is difficult to arrange resources into a linear ordering.

Edsger W. Dijkstra, The Structure of the "THE"-Multiprogramming System, Communications of the ACM, Vol. 11 (1968), No. 5 (May), pp. 341-346.

**THE** = **Technische Hogeschool Eindhoven** 

21

### Deadlock Avoidance Deadlocks can be avoided but not totally prevented

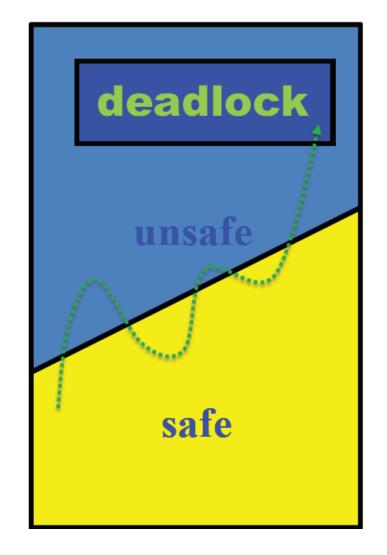
### Safe Sequences and Safe State

### Safe Sequences

- Each process provides the maximum number of resources of each type it needs.
- There are algorithms that can ensure the system will never enter a deadlock state. This is deadlock avoidance.
- A sequence of processes  $\langle P_1, P_2, ..., P_n \rangle$  is a **safe sequence** if for each process  $P_i$  in the sequence, its resource requests can be satisfied by the remaining resources and the sum of all resources that are being held by  $P_1, P_2, ..., P_{i-1}$ . This means we can suspend  $P_i$  and run  $P_1, P_2, ..., P_{i-1}$  until they complete. Then,  $P_i$  will have all resources to run.

#### **Safe State**

- A state is **safe** if the system can allocate resources to each process (up to its maximum, of course) in some order and still avoid a deadlock.
- Thus, a state is **safe** if there is a **safe sequence**. Otherwise, if no safe sequence exists, the system state is **unsafe**.
- An unsafe state is not necessarily a deadlock state.
   On the other hand, a deadlock state is an unsafe state.



#### Example: 1/2

• A system has 12 tapes and three processes A, B, C. At time  $t_0$ , we have:

3 free tapes

	Max needs	Current holding		Will need	
A	10		5		5
B	4		2		2
C	9		2	į	7

- Then,  $\langle B, A, C \rangle$  is a safe sequence (safe state).
- The system has 12-(5+2+2)=3 free tapes.
- Because B needs 2 tapes, it can take 2, run, and return 4. After B completes, the system has (3-2)+4=5 tapes. A now can take all 5 tapes and run. Finally, A returns 10 tapes for C to take 7 of them.

#### Example: 2/2

• A system has 12 tapes and three processes A, B, C. At time  $t_1$ , C has one more tape:

	Max needs	<b>Current holding</b>	Will need
$\boldsymbol{A}$	10	5	5
B	4	2	2
C	9	3	6

- The system has 12-(5+2+3)=2 free tapes.
- At this point, only *B* can take these 2 and run. It returns 4, making 4 free tapes available.
- But, none of A and C can run, and a deadlock occurs.
- The problem is due to granting C one more tape.

# Deadlock Avoidance and Banker's Algorithm

#### **Deadlock Avoidance**

- A deadlock avoidance algorithm ensures that the system is always in a safe state. Therefore, no deadlock can occur.
- Resource requests are granted only if in doing so the system is still in a safe state.
- Consequently, resource utilization may be lower than those systems without using a deadlock avoidance algorithm.

#### **Basic Idea: 1/2**

- Suppose you are a banker with a limited asset.
- If a customer A wishes to borrow X amount.
- If you have no less than X, A can have it.
- If you have less than X, A's request cannot be satisfied, and must wait until other customers pay their loans off so that the returned amount plus the amount in hand can satisfy A's request.

#### Basic Idea: 2/2

- Suppose you have \$1,000 and your brothers Adam, Bill and Chuck borrowed \$400, \$300 and \$100, respectively.
- Thus, you have only \$200 in hand.
- What if Adam asks for \$600 more?
- You don't have that amount for him. But, you think: if Bill and Chuck will return their \$400 and \$300 (assuming everyone is honest), respectively, you will have enough for Adam.
- In this way, you just ask Adam to wait until Bill and Chuck return the money.

#### **Assumptions**

- The system has m resource types.
- Each resource type has several units.
- There are *n* processes, each of which declares its maximum needs although these resources may not all be used at the same time.
- The deadlock avoidance system requires several matrices to keep track the states of resource allocation.

#### **Data Structures: 1/2**

- The following arrays are used:
  - ightharpoonup Available[1..m]: one entry for each resource. Available[i]=k means resource type i has k units available.
  - ightharpoonup Max[1..n,1..m]: maximum demand of each process. Max[i,j]=k means process i needs k units of resource j.
  - ightharpoonup Allocation[1..n,1..m]: resources allocated to each process. Allocation[i,j]=k means process i is currently allocated k units of resource j.
  - $\triangleright$  Need[1..n,1..m]: the remaining resource need of each process. Need[i,j]=k means process i needs k more units of resource j.
  - $\triangleright$  Thus, Max = Allocation + Need.

#### **Data Structures: 2/2**

- We will use A[i,\*] to indicate the i-th row of matrix A.
- Given two arrays A[1..m] and B[1..m],  $A \le B$  if  $A[i] \le B[i]$  for all i.
- Given A[1..n, 1..m] and B[1..n, 1..m], •  $A[i,*] \le B[i,*]$  if  $A[i,j] \le B[i,j]$  for all j.

# Finding a Safety Sequence The Safety Algorithm

- 1. Let Work[1..m] and Finish[1..n] be two working arrays.
- 2. Work := Available and Finish[i]=FALSE for all i
- 3. Find an *i* such that the following both hold
  - ightharpoonup Finish[i] = FALSE // process i is not yet done
- 4. Work = Work + Allocation[i,\*] // run it and reclaim

  Finish[i] = TRUE // process i completes

  go to Step 3
- 5. If Finish[i] = TRUE for all i, the system is in a safe state.

add this process i to the safe sequence

### **Banker's Algorithm: 1/3**

- In addition to *Available*, *Max*, *Allocation* and *Need*, we need one more array *Request*[1..n,1..m] to ensure a resource request can be granted.
- Request[i,j] = k: process i requests k units of resource type j.
- Process *i* calls Request() with an argument x[1..m], where x[j] = k means process *i* requests k units of resource type j.
- After process i finishes using some resources, it calls **Release()** with argument x[1..m], where x[j] = k means process i releases k units of resource type j.

# Banker's Algorithm: 2/3

- Process i calls Request(x) with its resource request in argument x.
- The **Resource-Request** algorithm uses x to update the Request[1..n,1..m] matrix.
- Then, Resource-Request runs the checking procedure by calling the Safety Algorithm to determine whether the system is in a safe state if the request is allocated.
- If it is safe, the process gets the resources. Otherwise, it is blocked until the process can get the needed resources.

# Banker's Algorithm: 3/3 Resource-Request

- 1. Let Request[1..n, 1..m] be the request matrix. Request[i,j]=k means process i requests k units of resource j.
- 2. If  $Request[i,*] \le Need[i,*]$ , go to Step 3. Otherwise, it is an error.
- 3. If  $Request[i,*] \le Available$ , go to Step 4. Otherwise, process i waits.
- 4. Do the following:

```
Available = Available - Request[i,*]
Allocation[i,*] = Allocation[i,*] + Request[i,*]
Need[i,*] = Need[i,*] - Request[i,*]
```

- 5. Call the Safety Algorithm.
- 6. If the request is safe, the requested resources are granted. Otherwise, process i waits and the resource-allocation tables are restored back to the original.

#### **Example 1: 1/4**

	All	ocati	on		Max	Λ	leed=	Max-	Alloc	$\boldsymbol{A}$	vailal	ble
	X	Y	Z	X	Y	Z	X	Y	Z	X	Y	Z
A	0	1	0	7	5	3	7	4	3	3	3	2
В	2	0	0	3	2	2	1	2	2			
C	3	0	2	9	0	2	6	0	0			
D	2	1	1	2	2	2	0	1	1			
E	0	0	2	4	3	3	4	3	1			

- The table is always scanned top-down.
- Because B's  $Need = [1,2,2] \le Available = [3,2,2], B$  runs.
- After *B* returns its *Allocation*. Therefore, *Allocation* = [2,0,0] + [3,3,2] = [5,3,2].
- Safe Sequence: B

#### **Example 1: 2/4**

	All	ocati	on		Max	Λ	leed=	Max-	Alloc	$\boldsymbol{A}$	vailal	ble
	X	Y	Z	X	Y	Z	X	Y	Z	X	Y	Z
A	0	1	0	7	5	3	7	4	3	5	3	2
В												
C	3	0	2	9	0	2	6	0	0			
D	2	1	1	2	2	2	0	1	1 🔻			
E	0	0	2	4	3	3	4	3	1			

- The table is always scanned top-down.
- *D* runs next. After this, *Available* = [5,3,2] + [2,1,1] = [7,4,3].
- Safe Sequence: B, D

#### **Example 1: 3/4**

	All	ocatio	on		Max	Λ	leed=	Max-	Alloc	$\boldsymbol{A}$	vailal	ble
	X	Y	Z	X	Y	Z	X	Y	Z	X	Y	Z
A	0	1	0	7	5	3	7	4	3	7	4	3
В												
C	3	0	2	9	0	2	6	0	0			
D									Ť			
E	0	0	2	4	3	3	4	3	1			

- The table is always scanned top-down.
- A runs next.
- Then, Available = [7,4,3] + [0,1,0] = [7,5,3].
- Safe Sequence: B, D, A

#### **Example 1: 4/4**

	All	ocatio	on		Max	Λ	leed=	Max-	Alloc	$\boldsymbol{A}$	vailal	ble
	X	Y	Z	X	Y	Z	X	Y	Z	X	Y	Z
A										7	5	3
В												
C	3	0	2	9	0	2	6	0	0			
D									Ť			
E	0	0	2	4	3	3	4	3	1			

- The table is always scanned top-down.
- Because *C*'s  $[6,0,0] \le Available = [7,5,3]$ , *C* runs.
- After this, Available = [7,5,3] + [3,0,2] = [10,5,7] and E runs.
- Safe Sequence: B, D, A, C, E.
- There are other safe sequences:  $\langle D, E, B, A, C \rangle$ ,  $\langle D, B, A, E, C \rangle$ , ...

#### Example 2: <mark>1/3</mark>

- Now suppose process B asks for 1 X and 2 Zs. More precisely,  $Request_B = [1,0,2]$ . Is the system still in a safe state if this request is granted?
- Since  $Request_B = [1,0,2] \le Available = [3,3,2]$ , this request may be granted if the system is safe.
- If this request is granted, we have the following:

	All	ocatio	on		Max	Λ	leed=	Max-	Alloc	A	vailabl	e
	X	Y	Z	X	Y	Z	X	Y	Z	X	Y	Z
A	0	1	0	7	5	3	7	4	3	2	3	0
B	3	0	2.	3	2	2	0	2	0.	-444	•••••	***
C	3	<b>/</b> 0	2	9	0	2	6	<b>/</b> 0	0		1	
D	2,	1	1	2	2	2	0 /	1	1		,	
E	0	0	2	4	3	3	4	3	1			

$$[3,0,2]=[2,0,0]+[1,0,2]$$

$$[0,2,0]=[1,2,2]-[1,0,2]$$

$$[3,0,2]=[2,0,0]+[1,0,2]$$
  $[0,2,0]=[1,2,2]-[1,0,2]$   $[2,3,0]=[3,3,2]-[1,0,2]$ 

#### **Example 2: 2/3**

	All	ocati	on		Max	Λ	leed=	Max-	Alloc	$\boldsymbol{A}$	vailal	ble
	X	Y	Z	X	Y	Z	X	Y	Z	X	Y	Z
A	0	1	0	7	5	3	7	4	3	2	3	0
B	3	0	2	3	2	2	0	2	0.	4 8 8 8	*****	
C	3	0	2	9	0	2	6	0	0			
D	2	1	1	2	2	2	0	1	1			
E	0	0	2	4	3	3	4	3	1			

- Is the system in a safe state after this allocation?
- Yes, because the safety algorithm will provide a safe sequence  $\langle B,D,E,A,C \rangle$ . Verify it yourself.
- Hence, B's request of [1,0,2] can safely be made.

#### **Example 2: 3/3**

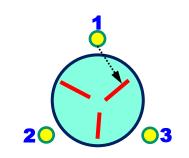
	All	ocati	on		Max	Λ	leed=	Max-	Alloc	$\boldsymbol{A}$	vailal	ble
	X	Y	Z	X	Y	Z	X	Y	Z	X	Y	Z
A	0	1	0	7	5	3	7	4	3	2	3	0
В	3	0	2	3	2	2	0	2	0			
C	3	0	2	9	0	2	6	0	0			
D	2	1	1	2	2	2	0	1	1			
E	0	0	2	4	3	3	4	3	1			

- Then, E's request  $Request_E = [3,3,0]$  cannot be granted because  $Request_E = [3,3,0] \le [2,3,0]$  is false.
- A's request  $Request_A = [0,2,0]$  cannot be granted because the system will be unsafe.
- If  $Request_A = [0,2,0]$  is granted, Available = [2,1,0] = [2,3,0] [0,2,0].
- None of the five processes can finish and the system is unsafe.

## **Example 3: 1/12**

- Let us look at how Banker's algorithm handles the dining philosophers problem.
- We use three philosophers  $P_1$ ,  $P_2$  and  $P_3$  and three chopsticks  $C_1$ ,  $C_2$  and  $C_3$ .
- Philosopher  $P_i$  uses chopsticks  $C_i$  and  $C_{i\%3+1}$ .
- What if  $P_1$  requests  $C_1$  (left),  $P_2$  requests  $C_2$  (left),  $P_3$  requests  $C_3$  (left),  $P_1$  requests  $C_2$  (right),  $P_2$  requests  $C_3$  (right),  $P_3$  requests  $C_1$  (right).
- This request sequence can cause a deadlock.
- Can Banker's algorithm overcome this problem?

# **Example 3: 2/12**



#### • Initially we have the following:

_	Al	locati	on		Max		Need	=Max	-Alloc	A	vailal	ble	R	eques	<u>t</u>
	$C_1$	$C_2$	$C_3$	$C_1$	$C_2$	$C_3$	$C_1$	$C_2$	<i>C</i> <sub>3</sub>	$C_1$	$C_2$	<i>C</i> <sub>3</sub>	$C_1$	$C_2$	<i>C</i> <sub>3</sub>
<b>P</b> <sub>1</sub>	0	0	0	1	1	0	1	1	0	1	1	1			
$P_2$	0	0	0	0	1	1	0	1	1						
<b>P</b> <sub>3</sub>	0	0	0	1	0	1	1	0	1						

#### • $P_1$ requests to use $C_1$ , Request = [1,0,0].

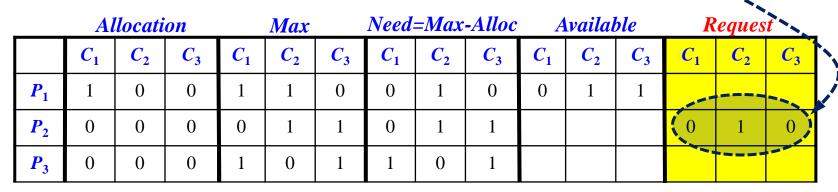
	$\boldsymbol{A}$	llocati	on		Max	_	Need	=Max	-Alloc	A	vaila	ble	R	Reques	st	_
	<i>C</i> <sub>1</sub>	$C_2$	$C_3$	$C_1$	$C_2$	$C_3$	<i>C</i> <sub>1</sub>	$C_2$	<i>C</i> <sub>3</sub>	$C_1$	<i>C</i> <sub>2</sub>	$C_3$	$C_1$	_C <sub>2</sub> _	$C_3$	4.
$P_1$	0	0	0	1	1	0	1	1	0	1	1	1	(1	0	0	
$P_2$	0	0	0	0	1	1	0	1	1							
<b>P</b> <sub>3</sub>	0	0	0	1	0	1	1	0	1							

# **Example 3: 3/12**

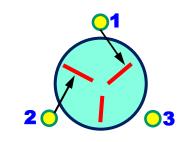
Assume this request is granted:

	A	locati	on		Max		Need	<u>=Max</u>	-Alloc	A	vailal	ble	R	eques	t
	$C_1$	$C_2$	$C_3$	<i>C</i> <sub>1</sub>	$C_2$	$C_3$	<i>C</i> <sub>1</sub>	$C_2$	<i>C</i> <sub>3</sub>	$C_1$	$C_2$	$C_3$	<i>C</i> <sub>1</sub>	$C_2$	<i>C</i> <sub>3</sub>
<b>P</b> <sub>1</sub>	1	0	0	1	1	0	0	1	0	0	1	1	1	0	0
$P_2$	0	0	0	0	1	1	0	1	1						
$P_3$	0	0	0	1	0	1	1	0	1						

- Obviously, the system is safe.
- $P_2$  requests to use  $C_2$ , Request = [0,1,0].



# **Example 3: 4/12**



Assume this request is granted:

	Al	locati	on		Max		Need	=Max	-Alloc	A	<u>vailal</u>	ble	R	eques	t
	$C_1$	$C_2$	C	<i>C</i> <sub>1</sub>	$C_2$	$C_3$	$C_1$	$C_2$	$C_3$	$C_1$	$C_2$	C <sub>3</sub>	<i>C</i> <sub>1</sub>	$C_2$	<i>C</i> <sub>3</sub>
<b>P</b> <sub>1</sub>	1	0	0	1	1	0	0	1	0	0	0	1		-	
<b>P</b> <sub>2</sub>	0	1	0	0	1	1	0	0	1				0	1	0
<b>P</b> <sub>3</sub>	0	0	0	1	0	1	1	0	1						

- Obviously, the system is safe, again.
- $P_3$  requests to use  $C_3$ , Request = [0,0,1].

	Al	locati	on	_	Max		Need	<u>=Max</u>	-Alloc	A	vaila	ble	R	eques	st
	$C_1$	$C_2$	$C_3$	$C_1$	$C_2$	$C_3$	<i>C</i> <sub>1</sub>	$C_2$	$C_3$	$C_1$	$C_2$	<i>C</i> <sub>3</sub>	<i>C</i> <sub>1</sub>	$C_2$	Ca
<b>P</b> <sub>1</sub>	1	0	0	1	1	0	0	1	0	0	0	1			
$P_2$	0	1	0	0	1	1	0	0	1						
<b>P</b> <sub>3</sub>	0	0	0	1	0	1	1	0	1				(0	0	1)

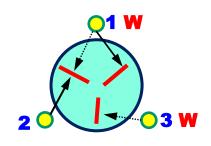
# **Example 3: 5/12**

Assume this request is granted:

	Al	locati	on		Max		Need	=Max	-Alloc	A	vailal	ble	R	eques	t
	$C_1$	$C_2$	$C_3$	$C_1$	$C_2$	$C_3$	$C_1$	_C <sub>2</sub> _	C <sub>5</sub>	<u>-C</u> 1-	- C <sub>2</sub> -	$C_3$	$C_1$	$C_2$	<i>C</i> <sub>3</sub>
<b>P</b> <sub>1</sub>	1	0	0	1-	1	0	0	1	0	0	0	0		1	
$P_2$	0	1	0 /	0	1	1	0	0	1						
<b>P</b> <sub>3</sub>	0	0	1	1	0	1	1	0	0				0	0	1

- Is the system safe?
- No! There is no safe sequence, because none of the three Need array rows is ≤ Available.
- As a result,  $P_3$  must wait, and a possible deadlock is avoided.

# **Example 3: 6/12**



 $P_3$ 's request cannot be granted:

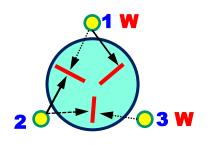
		A	locati	on		Max	_	Need	=Max	-Alloc	A	vaila	ble	R	Reques	st	_
Ī		$C_1$	$C_2$	$C_3$	$C_1$	$C_2$	$C_3$	<i>C</i> <sub>1</sub>	$C_2$	<i>C</i> <sub>3</sub>	$C_1$	$C_2$	$C_3$	<i>C</i> <sub>1</sub>	$C_2$	<i>C</i> <sub>3</sub>	
Ī	<b>P</b> <sub>1</sub>	1	0	0	1	1	0	0	1	0	0	0	1				
Ī	$P_2$	0	1	0	0	1	1	0	0	1							
Ī	<b>P</b> <sub>3</sub>	0	0	0	1	0	1	1	0	1				0	0	1	wai

•  $P_1$  requests to use  $C_2$ , Request = [0,1,0].-----

		A	locati	on	-	Max	_	Need	=Max	-Alloc	A	vailal	ble	R	Reques	st	
I		$C_1$	$C_2$	$C_3$	$C_1$	$C_2$	$C_3$	<b>C</b> <sub>1</sub>	$C_2$	$C_3$	$C_1$	$C_2$	<i>C</i> <sub>3</sub>	$C_1$	<u>C</u> 2	$C_3$	
ſ	<b>P</b> <sub>1</sub>	1	0	0	1	1	0	0	1	0	0	0	1	(0	1	0	
Ī	<b>P</b> <sub>2</sub>	0	1	0	0	1	1	0	0	1							
Ī	<b>P</b> <sub>3</sub>	0	0	0	1	0	1	1	0	1				0	0	1	wai

■ Because  $P_1$ 's *Request* is not  $\leq A$  *vailable*,  $P_1$  waits!

## **Example 3: 7/12**



Now we have the following:

	A	locati	on		Max		Need	=Max	-Alloc	A	vaila	ble	R	eques	st .	_
	$C_1$	$C_2$	$C_3$	$C_1$	$C_2$	$C_3$	$C_1$	$C_2$	$C_3$	$C_1$	$C_2$	<i>C</i> <sub>3</sub>	$C_1$	$C_2$	$C_3$	
<b>P</b> <sub>1</sub>	1	0	0	1	1	0	0	1	0	0	0	1	0	1	0	wait
<b>P</b> <sub>2</sub>	0	1	0	0	1	1	0	0	1							
<b>P</b> <sub>3</sub>	0	0	0	1	0	1	1	0	1				0	0	1	wait

•  $P_2$  requests  $C_3$ ? Thus, Request = [0,0,1].

		Al	locati	on		Max		Need	=Max	-Alloc	A	vaila	ble	R	Reques	st	_
	C	71	$C_2$	$C_3$	$C_1$	$C_2$	$C_3$	$C_1$	$C_2$	<i>C</i> <sub>3</sub>	$C_1$	$C_2$	$C_3$	$C_1$	$C_2$	<i>C</i> <sub>3</sub>	
P	1	1	0	0	1	1	0	0	1	0	0	0	1	0	1	0	wait
P	2 (	0	1	0	0	1	1	0	0	1				0	0	1	41
P	3 (	0	0	0	1	0	1	1	0	1				0	0	1	wait

Can this request be satisfied?

## **Example 3: 8/12**

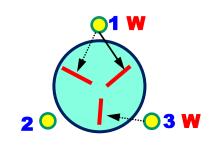
• Assume  $P_2$ 's Request = [0,0,1] is granted: 26

		$\boldsymbol{A}$	llocati	on		Max	_	Need	=Max	-Alloc	A	vaila	ble	R	eques	<u>t</u>	_
		$C_1$	C <sub>2</sub>	$C_3$	$C_1$	$C_2$	<u>C</u> <sub>2</sub>	_C <sub>F</sub> .	-C <sub>2</sub> -	C <sub>3</sub> -	- <b>C</b> T -	_C <sub>2</sub>	<i>C</i> <sub>3</sub>	$C_1$	$C_2$	$C_3$	
Ī	<b>P</b> <sub>1</sub>	1	0	0	مام	1	0	0	1	0	0	0	0	0	1	0	wait
Ī	<b>P</b> <sub>2</sub>	0	1	1	0	1	1	0	0	0				0	0	1	
	<b>P</b> <sub>3</sub>	0	0	0	1	0	1	1	0	1				0	0	1	wait

- Because  $P_2$ 's  $Need \le Available$ ,  $P_2$  can run. Note that  $P_1$  and  $P_3$  are waiting, and the system is safe.
- After  $P_2$  finishes using the requested resources,  $P_2$  releases  $C_2$  and  $C_3$ :

_	t	eques	R	ble	vailal	A	<u>-Alloc</u>	<u>=Max</u>	Need:		Max		on	locati	A	
	$C_3$	$C_2$	$C_1$	$C_3$	$C_2$	$C_1$	$C_3$	$C_2$	$C_1$	$C_3$	$C_2$	$C_1$	$C_3$	$C_2$	$C_1$	
wait	0	1	0	1	1	0	0	1	0	0	1	1	0	0	1	<b>P</b> <sub>1</sub>
							1	1	0	1	1	0	0	0	0	$P_2$
wait	1	0	0				1	0	1	1	0	1	0	0	0	<b>P</b> <sub>3</sub>

# **Example 3: 9/12**



Now we have the following:

_	$\boldsymbol{A}$	llocati	on		Max		Need	=Max	-Alloc	A	vaila	ble	R	eques	st.	_
	$C_1$	$C_2$	<i>C</i> <sub>3</sub>	$C_1$	$C_2$	$C_3$	$C_1$	$C_2$	<i>C</i> <sub>3</sub>	$C_1$	<b>C</b> <sub>2</sub>	<i>C</i> <sub>3</sub>	$C_1$	$C_2$	$C_3$	
<b>P</b> <sub>1</sub>	1	0	0	1	1	0	0	1	0	0	1	1	0	1	0	wait
$P_2$	0	0	0	0	1	1	0	1	1							
<b>P</b> <sub>3</sub>	0	0	0	1	0	1	1	0	1				0	0	1	wait

- $P_1$  and  $P_3$  are waiting. Can one of them run?
- Because  $P_1$ 's  $Need = [0,1,0] \le Available = [0,1,1]$ ,  $P_1$  can run.
- Note that  $P_3$  cannot run, because its *Need* is not  $\leq Available$ .

# **Example 3: 10/12**

• After  $P_1$  finishes, we have the following:

	$\boldsymbol{A}$	llocati	on	_	Max		Need	=Max	-Alloc	A	vailal	ble	R	eques	t
	$C_1$	$C_2$	<i>C</i> <sub>3</sub>	$C_1$	$C_2$	$C_3$	$C_1$	$C_2$	<i>C</i> <sub>3</sub>	$C_1$	$C_2$	<i>C</i> <sub>3</sub>	$C_1$	$C_2$	$C_3$
$P_1$	0	0	0	1	1	0	1	1	0	1	1	1			
$P_2$	0	0	0	0	1	1	0	1	1						
$P_3$	0	0	0	1	0	1	1	0	1				0	0	1

- wait
- Now  $P_3$  can run, because its  $Need = [0,0,1] \le Available = [1,1,1]$ .
- Please verify the remaining activities.
- Banker's algorithm avoids possible deadlocks.

# **Example 3: 11/12**

- Please verify the following:
  - ➤ We have seen the version that requires each philosopher to pick up **BOTH** chopsticks at the same time. Thus, requests by P<sub>1</sub>, P<sub>2</sub> and P<sub>3</sub> are [1,1,0], [0,1,1] and [1,0,1], respectively. Use Banker's algorithm to show that it is deadlock-free.
  - ➤ Do the same with the weirdo (or lefty-righty) version.

# **Example 3: 12/12**

- After a process returns its resources, a waiting process may be released if its Request ≤ Available and the system is safe.
- When searching for a waiting process to be resumed, one may implement some search order such as
  - **Priority**
  - **>FIFO**
  - >Aging: priority gets higher as waiting longer
  - Some form of "fitting": first-fit, best-fit, worst-fit
- Banker's algorithm can easily be implemented with a monitor.
  57

# Deadlock Detection Deadlocks are not avoided but they will be detected

#### **Deadlock Detection: 1/2**

- If deadlock prevention or avoidance are not used, then a deadlock situation may occur. Thus, we need:
  - An algorithm that can examine the system state to determine if a deadlock has occurred. This is a *deadlock detection* algorithm.
  - An algorithm that can help recover from a deadlock. This is a *recovery* algorithm.

#### **Deadlock Detection: 2/2**

- A deadlock detection algorithm does not have to know the maximum need Max of a process.
- Thus, the current need Need is not available.
- A deadlock detection algorithm only uses Available, Allocation and Request.
- A deadlock detection algorithm works like Banker's algorithm without Need.
  - Whenever  $Request \leq Available$ , the requested resources are allocated.
  - Those processes that cannot get the requested resources may involve in a deadlock.

# **Deadlock Detection Algorithm**

- 1. Let Work[1..m] and Finish[1..n] be two working arrays.
- 2. Work := Available and Finish[i]=FALSE for all i
- 3. Find an *i* such that both
  - ightharpoonup Finish[i] = FALSE // process i is not yet done
- 4. Work = Work + Allocation[i,\*] // run it and reclaim

  Finish[i] = TRUE // process i completes

  go to Step 3
- 5. If Finish[i] = TRUE for all i, the system is in a safe state. If Finish[i] = FALSE, then process  $P_i$  is deadlocked.

## **Example 1: 1/5**

	All	ocati	on	R	eque	s <b>t</b>	Av	railab	le
	X	Y	Z	X	Y	Z	X	Y	Z
A	0	1	0	0	0	0	0	0	0
B	2	0	0	2	0	2			
C	3	0	3	0	0	0			
D	2	1	1	1	0	0			
E	0	0	2	0	0	2			

- **Suppose maximum available resource is [7,2,6] and the current state of resource allocation is shown above.**
- Is the system deadlocked?

# **Example 1: 2/5**

	All	ocati	on	R	eque.	st	Av	railab	le
	X	Y	Z	X	Y	Z	X	Y	Z
A	0	1	0	0	0	0	0	0	0
B	2	0	0	2	0	2			
<b>C</b>	3	0	3	0	0	0			
D	2	1	1	1	0	0			
E	0	0	2	0	0	2			

- Because A's  $Request \le Available$ , run A first.
- Then, A returns [0,1,0] and the new A vailable = [0,0,0] + [0,1,0] = [0,1,0].

#### **Example 1: 3/5**

	All	ocati	on	R	eque	st .	Av	railab	ole
	X	Y	Z	X	Y	Z	X	Y	Z
A						-	0	1	0
B	2	0	0	2	0	2			
<b>C</b>	3	0	3	0	0	0			
D	2	1	1	1	0	0			
E	0	0	2	0	0	2 *			

- Because C's  $Request \leq Available$ , run C.
- Then, C returns [3,0,3] and the new A vailable = [0,1,0] + [3,0,3] = [3,1,3].

#### **Example 1: 4/5**

	All	ocati	on	Request			Available		
	X	Y	Z	X	Y	Z	X	Y	Z
A							3	1	3
B	2	0	0	2	0	2			
C									
D	2	1	1	1	0	0			
E	0	0	2	0	0	2 *			

- Because B's  $Request \leq Available$ , run B.
- Then, B returns [2,0,0] and the new A vailable = [3,1,3] + [2,0,0] = [5,1,3].

## **Example 1: 5/5**

	All	ocati	on	Request			Available		
	X	Y	Z	X	Y	Z	X	Y	Z
A							5	1	3
B									
<b>C</b>									
D	2	1	1	1	0	0			
E	0	0	2	0	0	2 *			

- Then, we can run either D or E.
- As a result, this is not a deadlock state.

# **Example 2: 1/3**

	All	locati	on	R	eque	est	Available		
	X	Y	Z	X	Y	Z	X	Y	Z
A	0	1	0	0	0	0	0	0	0
B	2	0	0	2	0	2			
<b>C</b>	3	0	3	0	0	1			
D	2	1	1	1	0	0			
E	0	0	2	0	0	2			

- Suppose C requests for one more resource Z.
- Is this a deadlock state?

#### **Example 2: 2/3**

	All	locati	on	Request			Available		
	X	Y	Z	X	Y	Z	X	Y	Z
A	0	1	0	0	0	0	0	0	0
B	2	0	0	2	0	2			
<b>C</b>	3	0	3	0	0	1			
D	2	1	1	1	0	0			
E	0	0	2	0	0	2			

- Because A's  $Request \leq Available$ , A runs.
- Then, A returns its Allocation, and the new Available = [0,0,0] + [0,1,0] = [0,1,0].

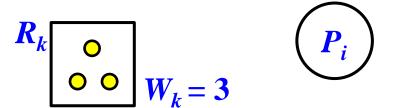
#### **Example 2: 3/3**

	Allocation			Request			Available		
	X	Y	Z	X	Y	Z	X	Y	Z
A							0	1	0
B	2	0	0	2	0	2			
<b>C</b>	3	0	3	0	0	1			
D	2	1	1	1	0	0			
E	0	0	2	0	0	2			

- None of the four remaining processes can run.
- Therefore, this is a deadlock state.

## Resource Allocation Graph: 1/9

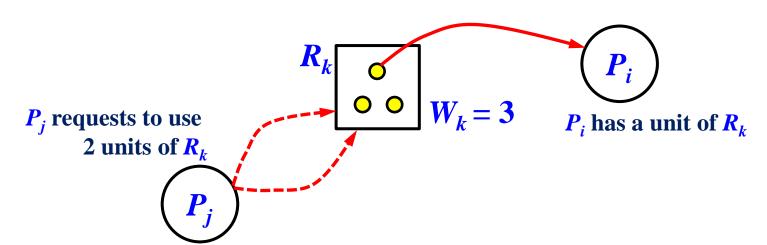
- Suppose we have m resource types,  $R_1, R_2, ..., R_m$  and resource type  $R_i$  has  $W_i$  units.
- Each resource type is represented by a *rectangular* node in which each "dot" indicates a unit. Therefore, resource type  $R_i$  has  $W_i$  dots.
- Suppose there are n processes  $P_1, P_2, ..., P_n$ , each process is represented by a *circular* node.





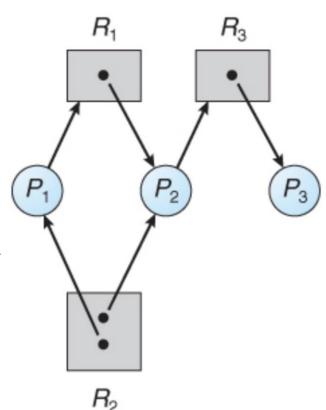
## Resource Allocation Graph: 2/9

- If process  $P_i$  makes a request to use a unit of resource  $R_j$ , draw a directed edge from  $P_i$  to  $R_j$ .
- If process  $P_i$  receives a resource of type  $R_j$ , draw a directed edge from a **dot** of  $R_i$  to  $P_i$ .
- This is a Resource Allocation Graph, RAG.



#### Resource Allocation Graph: 3/9

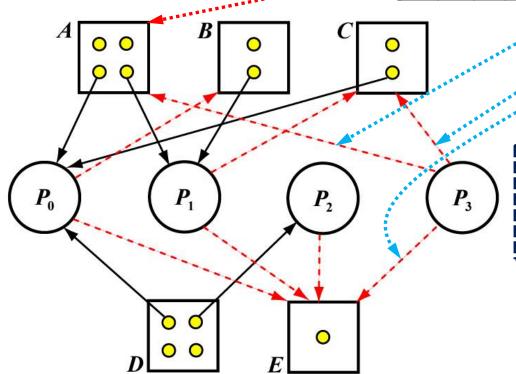
- The right diagram has 3 processes  $P_1$ ,  $P_2$  and  $P_3$  and 3 types of resources  $R_1$ ,  $R_2$  and  $R_3$ .
- Process  $P_1$  has a unit of  $R_2$  and requests a unit of  $R_1$ .
- Process  $P_2$  has a unit of  $R_1$  and a unit of  $R_2$ , and requests a unit of  $R_3$ .
- Process  $P_3$  has a unit of  $R_3$ .



## Resource Allocation Graph: 4/9

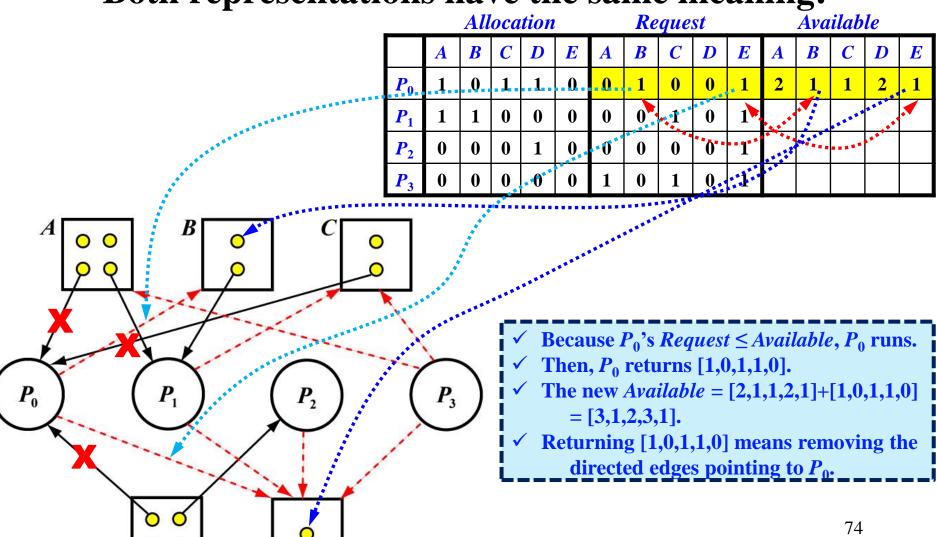
One can convert a table to a RAG and vice versa.

		_	All	loca	tion		Request					Available					
		A	B	C	<del>-D</del> -	E	<b>-</b> A -	B	c	D.	E	A	B	C	D	E	
	$P_0$	1	0	1	1	0	0	1	0	0	1	.2	1	1	2	1	
	$P_1$	1	1	0	0	0	Q	0	1	0	1						
	$P_2$	0	. 0	.0.	1	0	0	0	0	0	1						
*****	$P_3$	0	0	0	0	0	. 1	0	. 1	0	. 1						

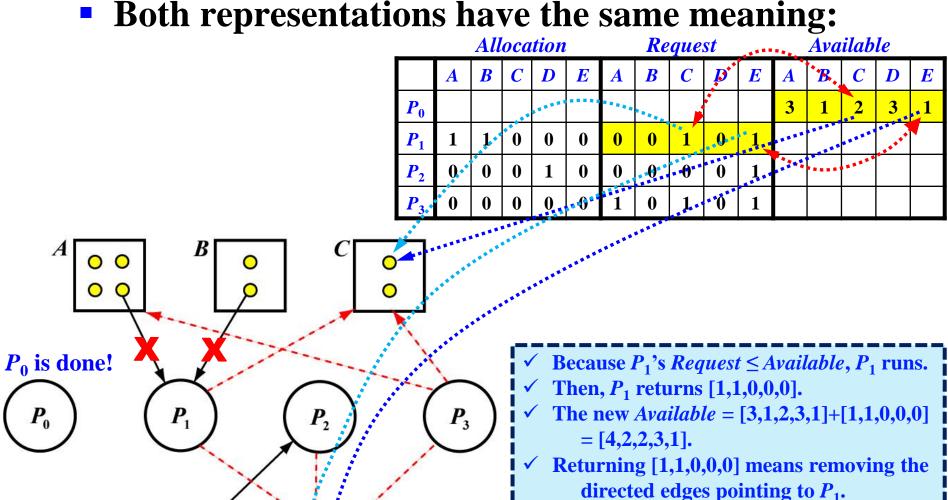


- ✓ The total number of units of each resource type can easily be computed.
- **✓ Dashed arrows are requests**
- ✓ Solid arrows are allocations.

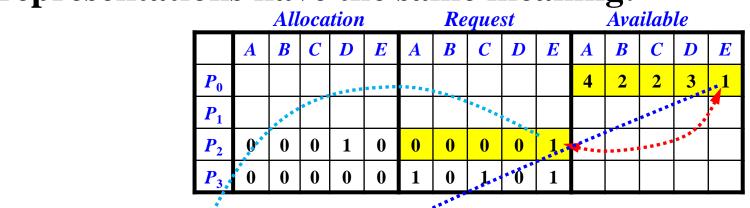
## Resource Allocation Graph: 5/9

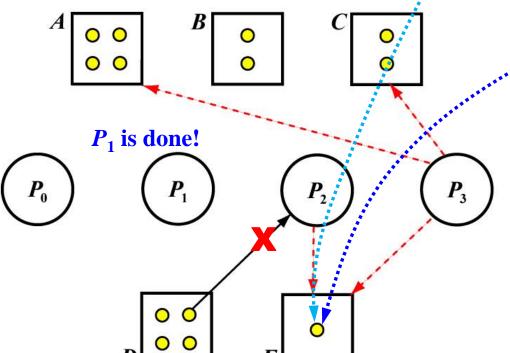


## Resource Allocation Graph: 6/9



## Resource Allocation Graph: 7/9

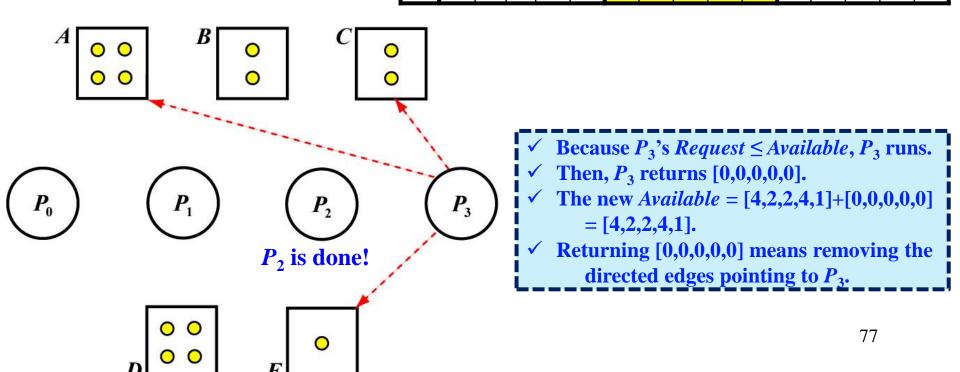




- ✓ Because  $P_2$ 's Request  $\leq$  Available,  $P_2$  runs.
- ✓ Then,  $P_2$  returns [0,0,0,1,0].
- ✓ The new Available = [4,2,2,3,1]+[0,0,0,1,0]= [4,2,2,4,1].
- ✓ Returning [0,0,0,1,0] means removing the directed edges pointing to  $P_2$ .

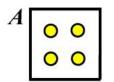
## Resource Allocation Graph: 8/9

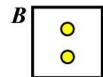
	_	All	oca	tion	!	Request					Available					
	A	B	<b>C</b>	D	E	A	B	C	D	E	A	B	C	D	E	
$P_0$											4	2	2	4	1	
$P_1$																
$P_2$																
$P_3$	0	0	0	0	0	1	0	1	0	1						

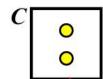


## Resource Allocation Graph: 9/9

	_	All	oca	tion	,	Request					Available				
	A	B	<b>C</b>	D	E	A	B	<b>C</b>	D	E	A	B	C	D	E
$P_0$											4	2	2	4	1
$P_1$															
$P_2$															
$P_3$	0	0	0	0	0	1	0	1	0	1					









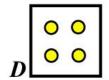


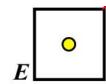






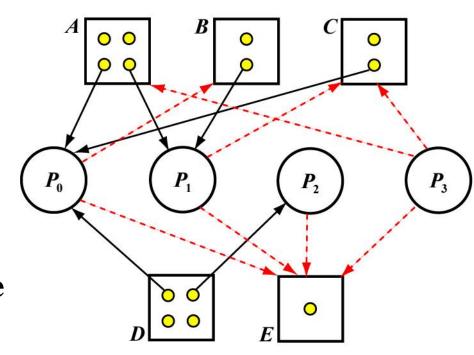
- ✓ Because  $P_3$ 's Request  $\leq$  Available,  $P_3$  runs.
- Then,  $P_3$  returns [0,0,0,0,0]. The new Available = [4,2,2,4,1] + [0,0,0,0,0]
- Returning [0,0,0,0,0] means removing the directed edges pointing to  $P_3$ .





#### **RAG** Reduction

- Consider a process P<sub>i</sub>, which is neither blocked nor an isolated node.
- A RAG is reduced by a process  $P_i$  by removing all edges to and from  $P_i$ .
- This reduction can be done if P<sub>i</sub> can have all its requests satisfied (i.e., Request ≤ Available).
- A RAG is completely reducible if all edges of the graph can be removed by a sequence of reductions.



- $\triangleright$   $P_0$  is a valid reduction, because  $P_0$ 's needs (i.e., 1 unit of B and 1 unit of E) are available.
- $\triangleright$   $P_1, P_2$  and  $P_3$  can also be good candidates.

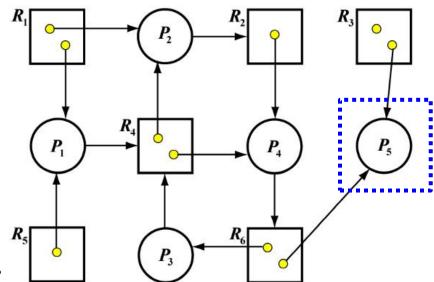
#### The Deadlock Theorem: 1/8

- The Deadlock Theorem: A state of resource allocation is a deadlock state if and only if its corresponding RAG is not completely reducible.
- The necessary conditions for having a deadlock implies that hold-and-wait and circular waiting must both hold.
- Therefore, this incompletely reducible RAG must contain a cycle (i.e., circular waiting).
- If a RAG does not contain a cycle, the state of resource allocation is deadlock-free.
- It is easy to determine whether a RAG contains a cycle (e.g., depth-first search).

  80

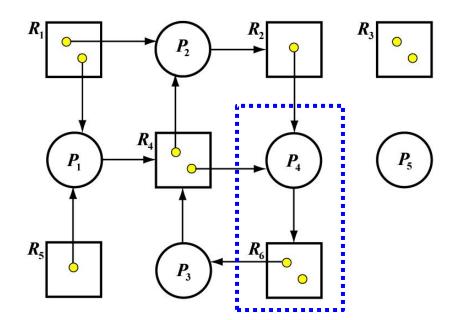
#### The Deadlock Theorem: 2/8

- Consider the right RAG.
- This RAG obviously can be reduced by process P<sub>5</sub>, because P<sub>5</sub> has everything it needs.
- Therefore, all edges to and from process P<sub>5</sub> are removed.



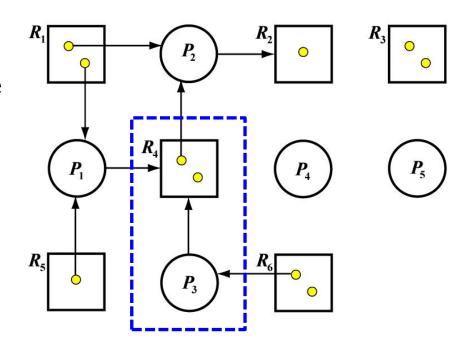
#### The Deadlock Theorem: 3/8

- After removing all edges from and to P<sub>5</sub>, we have the right RAG.
- Now  $P_4$  is a candidate, because it needs a unit of resource  $R_6$ , which is available.
- Then, all edges to and from process  $P_4$  are removed.



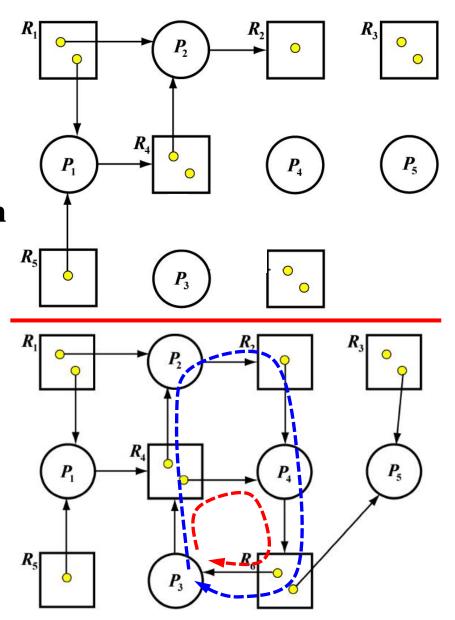
#### The Deadlock Theorem: 4/8

- After removing all edges from and to  $P_4$ , we have the right **RAG**.
- Now  $P_3$  is a candidate, because it needs a unit of resource  $R_4$ , which is available.
- Then, all edges to and from process  $P_3$  are removed.



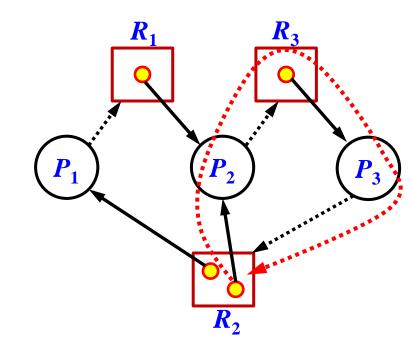
#### The Deadlock Theorem: 5/8

- Now,  $P_2$  can take the only unit of  $R_2$  and run.
- Then, process P<sub>1</sub> runs, and the reduction completes.
- Note that process  $P_1$  can run because its only need is a unit of  $R_4$ , and this needed unit is available.
- This RAG is deadlock-free, but it contains cycles!



#### The Deadlock Theorem: 6/8

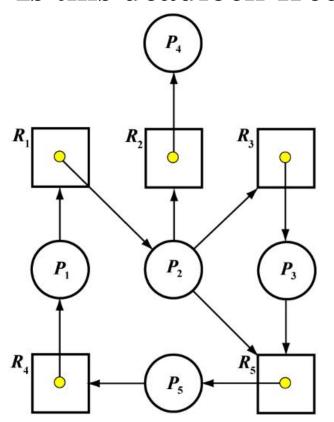
- $P_3$ 's request (1 unit of  $R_2$ ) cannot be met.
- $P_2$ 's request (1 unit of  $R_3$ ) cannot be met.
- $P_1$ 's request (1 unit of  $R_1$ ) cannot be met.
- Because this RAG cannot be completely reduced, a deadlock exists.
- There is a cycle.



The existence of a cycle does not always imply a deadlock state as discussed on this and the previous slides.

#### The Deadlock Theorem: 7/8

**Exercise:** Convert the following **RAG** to their table representations and do a step-by-step execution of the deadlock detection and **RAG** reduction. Is this deadlock-free?



#### The Deadlock Theorem: 8/8

- What we can conclude is: if a RAG has no directed cycle, the corresponding allocation state is deadlock-free.
- If a RAG has a directed cycle, we have two cases:
  - ➤ If there is only one unit per resource type, there is a deadlock. This is a perfect circular waiting.
  - ➤ Otherwise, a deadlock is only a possibility. Use **RAG** reduction or deadlock detection to determine whether the system is in a deadlock state.

## The Use of a Detection Algorithm

- The deadlock detection algorithm is not run when a request is made.
- It is run periodically when it is needed.
  - ➤ If deadlocks occur frequently, the detection algorithm should be invoked frequently.
  - **▶** Once per hour or whenever CPU utilization becomes low (*i.e.*, below 40%). Low CPU utilization usually means more processes are waiting.

#### **How to Recover**

- When a detection algorithm determines a deadlock has occurred, the algorithm may inform the system administrator to deal with it.
   Or, allow the system to recover from a deadlock.
- There are two options.
  - > Process Termination
  - **▶** Resource Preemption
- These two options are not mutually exclusive (*i.e.*, can have both if needed).

#### **Process Termination**

- Abort all deadlocked processes
- Abort one process at a time until the deadlock cycle is eliminated
- Problems:
  - Aborting a process may not be easy. What if a process is updating or printing a large file? The system must find some way to maintain the states of the files and printer before they can be reused.
  - Termination may be determined by the priority/importance of a process.

## **Resource Preemption**

- Selecting a victim: which resources and which processes are to be preempted?
- Rollback: If we preempt a resource from a process, what should be done with that process?
  - > Total Rollback: abort the process and restart it
  - ➤ Partial Rollback: rollback the process only as far as necessary to break the deadlock.
- Starvation: We cannot always pick the same process as a victim. Some limit must be set.

## Livelock

### Livelock: 1/3

- Livelock: If two or more processes continually repeat the same interaction in response to changes in the other processes without doing any useful work.
- These processes are not in the waiting state, and they are running concurrently.
- This is different from a deadlock because in a deadlock all processes are in the waiting state.

## Livelock: 2/3

- There are two locks, a process locks its own lock first.
- Then, checks whether the other lock is locked.
  - > If the other lock is not locked, this process locks it.
  - > Otherwise, this process releases its lock and locks its own lock again later

## Livelock: 3/3

```
MutexLock Mutex1, Mutex2;

Process 1

Mutex1.Lock();
while (Mutex2.isLocked()) {
   Mutex1.Unlock();
   // wait for a while
   Mutex1.Lock();
}
Mutex2.Lock();

Mutex2.Lock();

// wait for a while
   Mutex1.Lock();

Mutex2.Lock();

Mutex2.Lock();

Mutex2.Lock();

Mutex1.Lock();

Mutex1.Lock();
```

- What if both processes execute in a fully synchronized way?
- We have a livelock, each process reacts to the other's activity.
- Nothing meaningful is achieved.
- To avoid this type of livelock, order the locking sequence in a hierarchical way (i.e., both lock Mutex1 first followed by Mutex2.
- Can we combine both locks into only one lock?

# The End